

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI

Source Systems, Data Ingestion, and Pipelines

Week 1



DeepLearning.AI

Source Systems, Data Ingestion, and Pipelines

Welcome

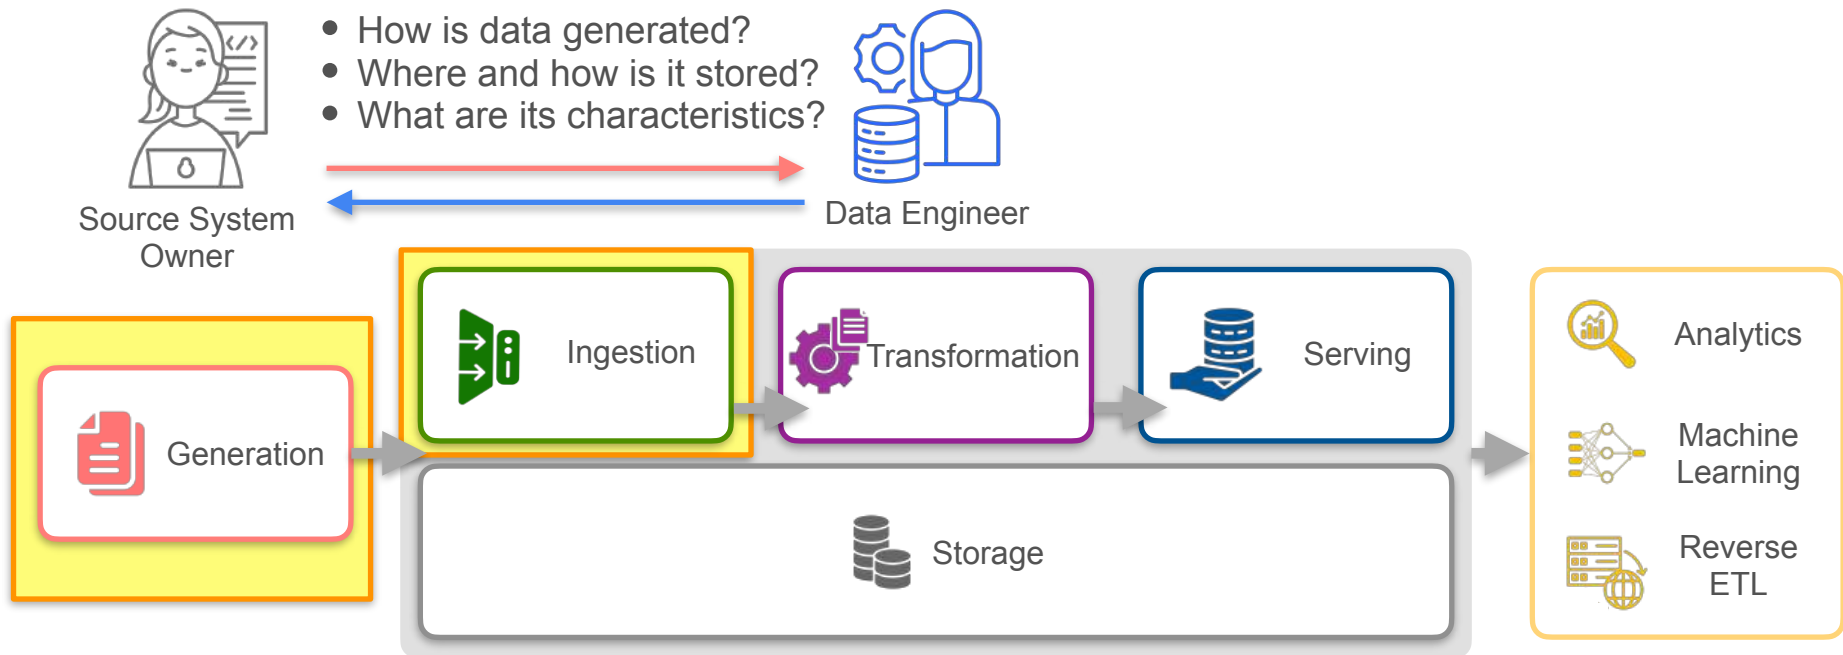


DeepLearning.AI

Working with Source Systems

Course 2 Overview

Data Engineering Lifecycle



Course Plan

Week 1

Common source systems

- Databases, object storage, and streaming sources
- Working with source systems on AWS

Week 2

Setting up ingestion from source systems

Week 3

DataOps undercurrent

- Automating some of your pipeline tasks
- Monitoring data quality

Week 4

Orchestration, monitoring, and automating data pipelines

- Setting up directed acyclic graphs
- Working with infrastructure as code



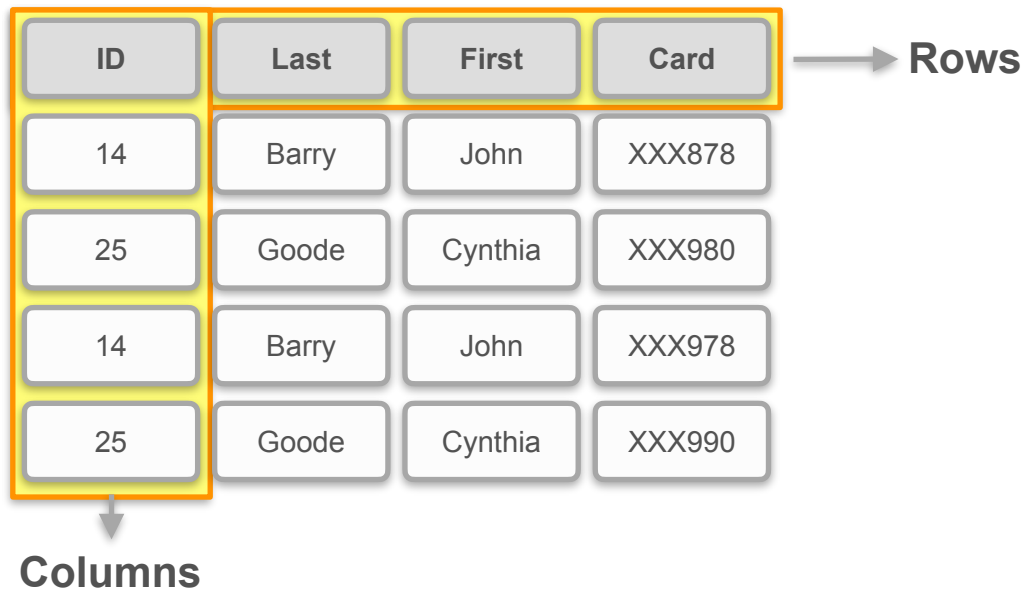
DeepLearning.AI

Introduction to Source Systems

Different Types of Source Systems

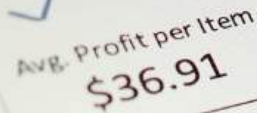
Structured Data

Data organized as tables of rows and columns



The diagram illustrates a table of structured data. The table has four columns: ID, Last, First, and Card. The first column (ID) is highlighted with a yellow background and a yellow border. The first row (headers) is highlighted with a yellow background and an orange border. An arrow points from the text 'Rows' to the first row, and another arrow points from the text 'Columns' to the first column.

ID	Last	First	Card
14	Barry	John	XXX878
25	Goode	Cynthia	XXX980
14	Barry	John	XXX978
25	Goode	Cynthia	XXX990



ROI
419%

Video by Adobe Stock (paid license)



```
import csv
with open('eggs.csv', newline='') as csvfile:
    spamreader = csv.reader(csvfile, delimiter=' ', quotechar='|')
    for row in spamreader:
        print(', '.join(row))
```

Structured Data

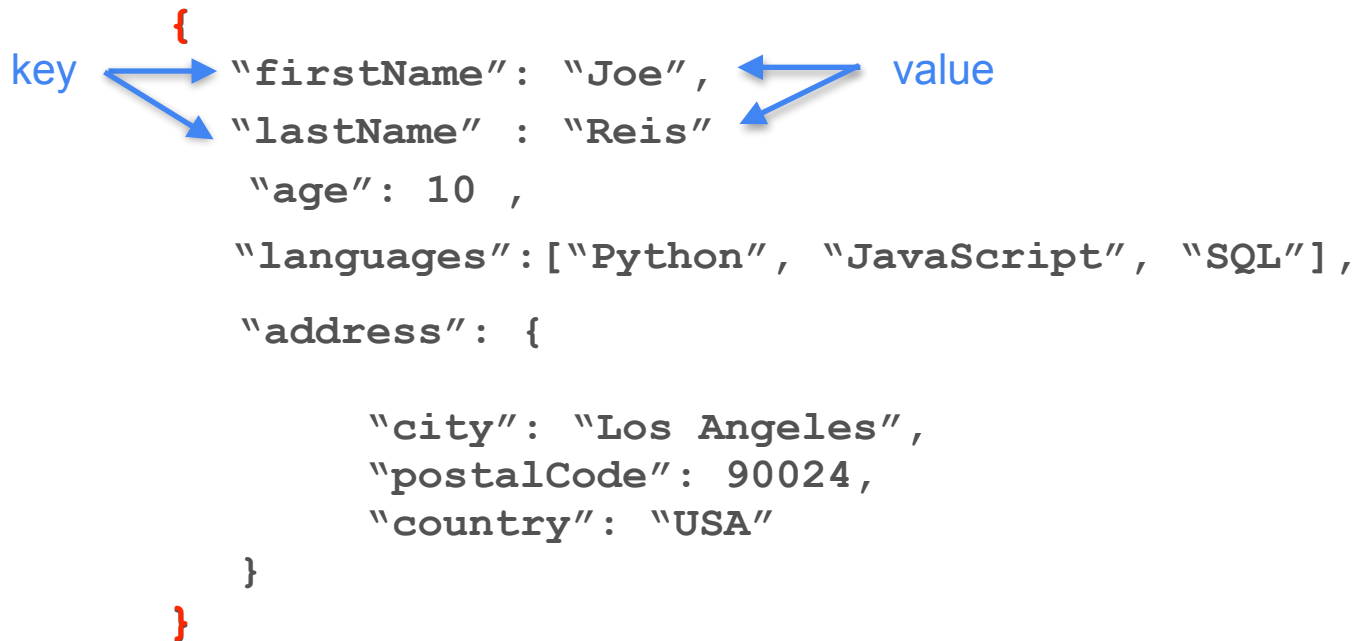
Data organized as tables of rows and columns

Semi-Structured Data

Data that is not in tabular form but still has some structure

JavaScript Object Notation (JSON)

A series of key-value pairs



```
{  
  "firstName": "Joe",  
  "lastName" : "Reis"  
  "age": 10 ,  
  "languages": ["Python", "JavaScript", "SQL"],  
  "address": {  
    "city": "Los Angeles",  
    "postalCode": 90024,  
    "country": "USA"  
  }  
}
```

Structured Data

Data organized as tables of rows and columns

Semi-Structured Data

Data that is not in tabular form but still has some structure

JavaScript Object Notation (JSON)

A series of key-value pairs

```
{  
  "firstName": "Joe",  
  "lastName": "Reis",  
  "age": 10,  
  "languages": ["Python", "JavaScript", "SQL"],  
  "address": {  
    "city": "Los Angeles",  
    "postalCode": 90024,  
    "country": "USA"  
  }  
}
```

The diagram illustrates the JSON structure with key-value pairs and nested objects. The main object is a JavaScript Object Notation (JSON) object. It contains several key-value pairs: "firstName": "Joe", "lastName": "Reis", "age": 10, "languages": ["Python", "JavaScript", "SQL"], and "address": { ... }. The "lastName" value "Reis" is highlighted with an orange box. The "age" value 10 is highlighted with an orange circle. The "languages" array ["Python", "JavaScript", "SQL"] is highlighted with an orange box. The "address" object is a nested JSON object, highlighted with an orange box. The "address" object contains three key-value pairs: "city": "Los Angeles", "postalCode": 90024, and "country": "USA". The "city" value "Los Angeles" is highlighted with an orange box. The "postalCode" value 90024 is highlighted with an orange box. The "country" value "USA" is highlighted with an orange box. Blue arrows point from the word "key" to the keys "firstName", "lastName", "age", "languages", and "address". Blue arrows point from the word "value" to the values "Joe", "Reis", 10, ["Python", "JavaScript", "SQL"], and the "address" object. Blue arrows point from the word "keys" to the keys "city", "postalCode", and "country". Blue arrows point from the word "values" to the values "Los Angeles", 90024, and "USA".

**Nested
JSON
format**

Structured Data

Data organized as tables of rows and columns

Semi-Structured Data

Data that is not in tabular form but still has some structure

Unstructured Data

Data that does not have any predefined structure

Text



Video



Audio



Images



- dimensions
- pixel colors



Databases

Structured data
Semi-structured data



Files



Streaming Systems

Semi-structured data



Databases

Store data in an organized way

Structured data

Semi-structured data

Create

Read

Update

Deleate



Database
Storage



Person/
Application





Databases

Store data in an organized way

Structured data

Semi-structured data

Create
Read
Uppdate
Deleate

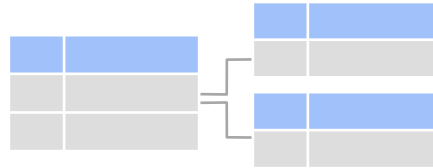


Database
Storage



Person/
Application

Relational databases



Tables with rows and columns

Non-relational (NoSQL) databases



Non-tabular data



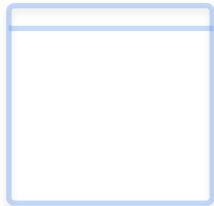
Files

Sequence of bytes
representing information

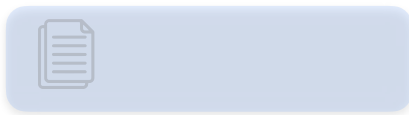


	A	B	C
1			
2			
3			

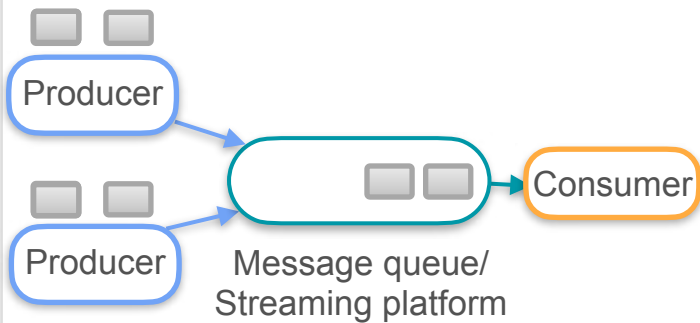
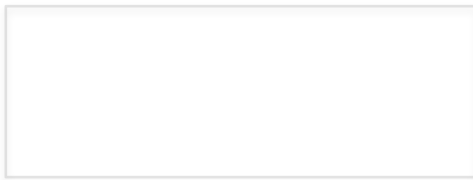
```
{  
  "firstName": "Joe",  
  "lastName" : "Reis",  
  "languages":["R", "SQL"],  
}
```

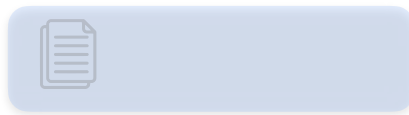
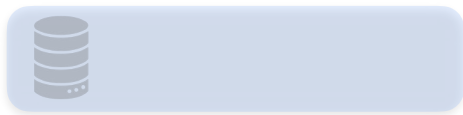


Amazon S3



	A	B	C
1			
2			
3			



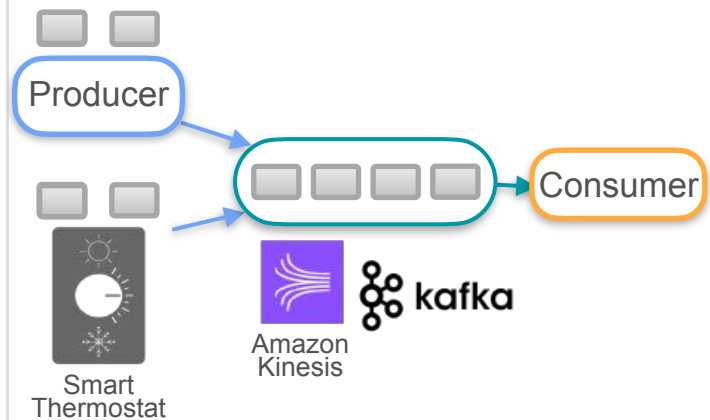
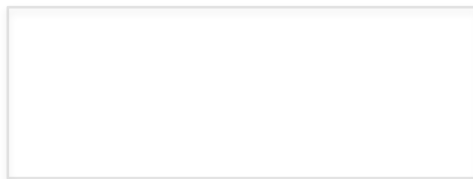


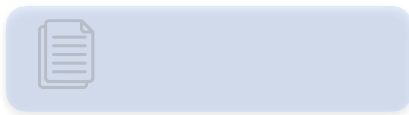
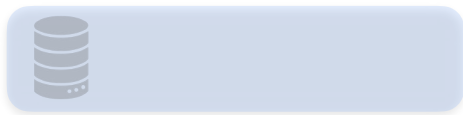
Continuous flow of data

Semi-structured data



	A	B	C
1			
2			
3			



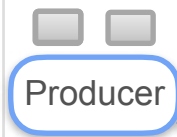
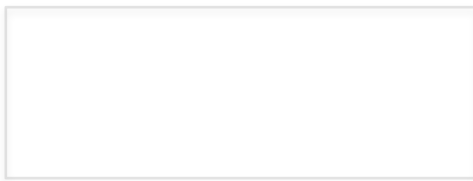


Continuous flow of data

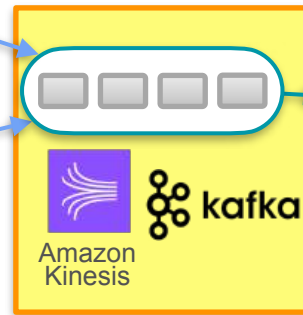
Semi-structured data



	A	B	C
1			
2			
3			



Source System



Your ingestion pipeline starts here



Databases

Store data in an organized way



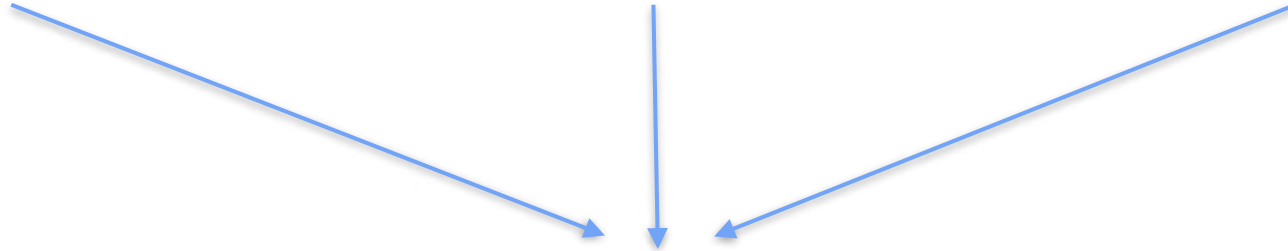
Files

Sequence of bytes
representing information



Streaming Systems

Continuous flow of data



Ingest

- Structured
- Semi-structured
- Unstructured

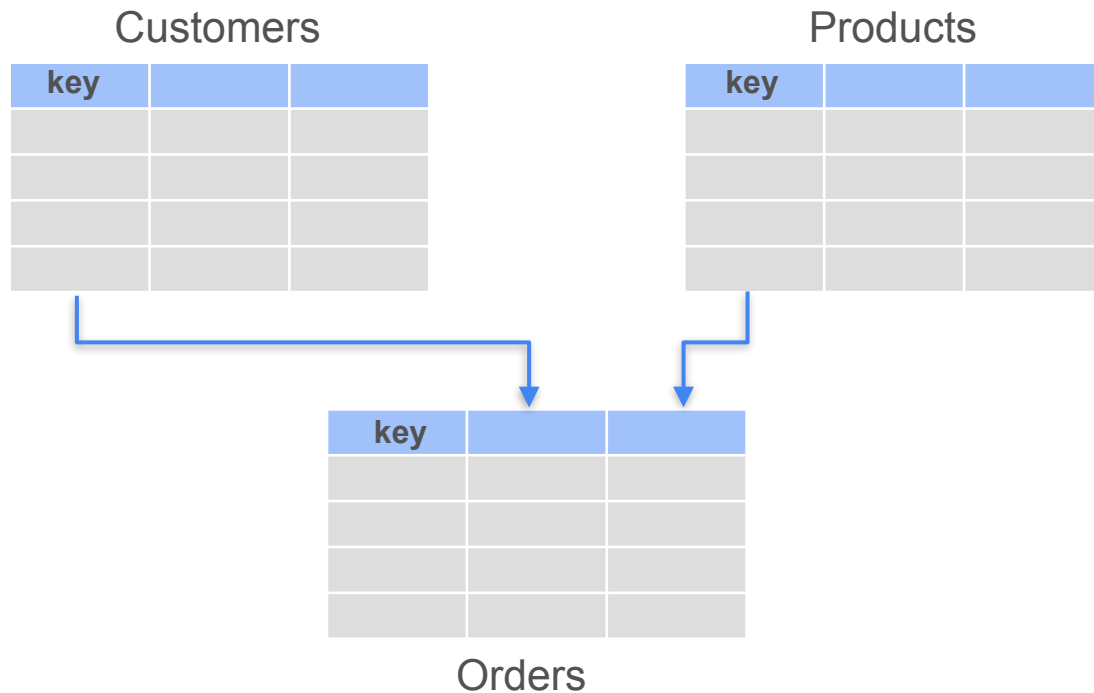


DeepLearning.AI

Introduction to Source Systems

Relational Databases

Relational Databases



- Reduce redundancy
- Make data easier to manage

Relational Databases

One big table for everything!

name	address	phone	date_time	amount	brand	SKU	description
Jane Doe	74th Street	12345678	12/08/2024	700	ABC	B32	Blender
Jane Doe	74th Street	12345678	12/08/2024	99	XYZ	i56	Iron
Jane Doe	74th Street	12345678	12/08/2024	100	GHJ	k70	Kettle



Jane Doe



Relational Databases

One big table for everything!

name	address	phone	date_time	amount	brand	SKU	description
Jane Doe	74th Street	12345678	12/08/2024	700	ABC	B32	Blender
Jane Doe	74th Street	12345678	12/08/2024	99	XYZ	i56	Iron
Jane Doe	74th Street	12345678	12/08/2024	100	GHJ	k70	Kettle
Mary Ann	19th Avenue	98765432	13/08/2024	899	STU	w40	Washer
John Ken	1st Link	36891623	14/08/2024	899	STU	w40	Washer
Ivy Tan	67th Street	98639513	15/08/2024	899	STU	w40	Washer



Relational Databases

One big table for everything!

Inconsistency

name	address	phone	date_time	amount	brand	SKU	description
Jane Doe	11th Avenue	12345678	12/08/2024	700	ABC	B32	Blender
Jane Doe	11th Avenue	12345678	12/08/2024	99	XYZ	i56	Iron
Jane Doe	74th Street	12345678	12/08/2024	100	GHJ	k70	Kettle
Mary Ann	19th Avenue	98765432	13/08/2024	899	STU	w31	Washer
John Ken	1st Link	36891623	14/08/2024	899	STU	w31	Washer
Ivy Tan	67th Street	98639513	15/08/2024	899	STU	w40	Washer



Jane Doe
now lives on 11th Avenue



SKU
now w31

Inconsistency

Relational Databases

Single customer

Customers

id	first_name	last_name	age	address
1	Jane	Doe	24	11th Ave.
2	Mary	Ann	65	19th Ave.
3	John	Ken	27	1st Link
4	Ivy	Tan	18	67th St.

Products

Single product

id	brand	SKU	description
1	ABC	b32	Blender
2	XYZ	i56	Iron
3	GHJ	k70	Kettle
4	STU	w31	Washer

Orders

id	customer_id	product_id	date_time	purchase_amount

Database schema

Relational Databases

Keys

Primary key:
uniquely
identifies each
row in a table

Customers

id	first_name	last_name	age	address
1	Jane	Doe	24	11th Ave.
2	Mary	Ann	65	19th Ave.
3	John	Ken	27	1st Link
4	Ivy	Tan	18	67th St.

Products

id	brand	SKU	description
1	ABC	b32	Blender
2	XYZ	i56	Iron
3	GHJ	k70	Kettle
4	STU	w31	Washer

Orders

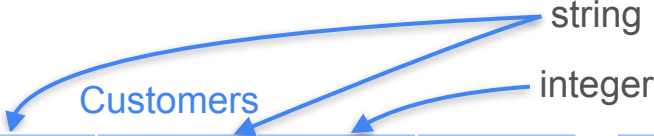
id	customer_id	product_id	date_time	purchase_amount
1	1	1	12/08/2024	700
2	1	2	12/08/2024	99
3	1	3	12/08/2024	100
4	2	4	13/08/2024	899
5	3	4	14/08/2024	899

Foreign key:

references the primary key of the customer table

Database schema

Relational Databases



id	first_name	last_name	age	address
1	Jane	Doe	24	11th Ave.
2	Mary	Ann	65	19th Ave.
3	John	Ken	27	1st Link
4	Ivy	Tan	18	67th St.

id	brand	SKU	description
1	ABC	b32	Blender
2	XYZ	i56	Iron
3	GHJ	k70	Kettle
4	STU	w31	Washer

Orders

id	customer_id	product_id	date_time	purchase_amount
1	1	1	12/08/2024	700
2	1	2	12/08/2024	99
3	1	3	12/08/2024	100
4	2	4	13/08/2024	899
5	3	4	14/08/2024	899

Database schema

Each row in a table has to follow the same column structure:
same sequence of columns and data types

Relational Databases

Customers

id	first_name	last_name	age	address
1	Jane	Doe	24	11th Ave.
2	Mary	Ann	65	19th Ave.
3	John	Ken	27	1st Link
4	Ivy	Tan	18	67th St.

Products

id	brand	SKU	description
1	ABC	b32	Blender
2	XYZ	i56	Iron
3	GHJ	k70	Kettle
4	STU	w31	Washer

Orders

id	customer_id	product_id	date_time	purchase_amount
1	1	1	12/08/2024	700
2	1	2	12/08/2024	99
3	1	3	12/08/2024	100
4	2	4	13/08/2024	899
5	3	4	14/08/2024	899
6	1	4	15/08/2024	899

Relational Databases

One big table for everything!

name	address	phone	date_time	amount	brand	SKU	description
Jane Doe	74th Street	12345678	12/08/2024	700	ABC	B32	Blender
Jane Doe	74th Street	12345678	12/08/2024	99	XYZ	i56	Iron
Jane Doe	74th Street	12345678	12/08/2024	100	GHJ	k70	Kettle
Mary Ann	19th Avenue	98765432	13/08/2024	899	STU	w40	Washer
John Ken	1st Link	36891623	14/08/2024	899	STU	w40	Washer
Ivy Tan	67th Street	98639513	15/08/2024	899	STU	w40	Washer

One Big Table (OBT) approach: use cases that need faster processing

Relational Databases

Relational Database Management System (RDBMS)

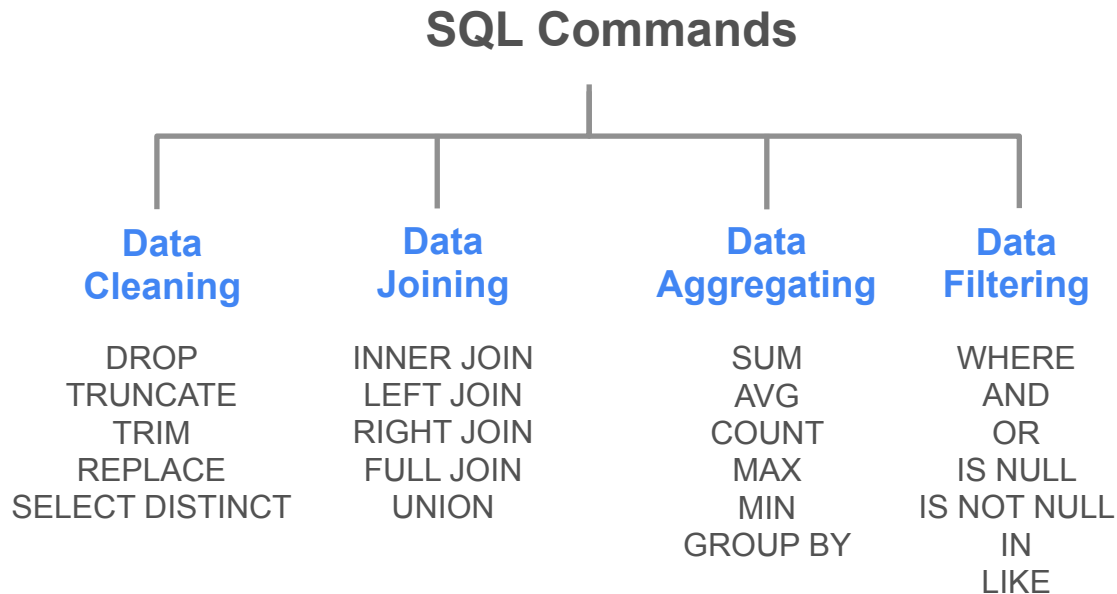


Software layer that sits on top of a relational database to manage and interact with the data.



Structured Query Language (SQL)

Relational Databases





DeepLearning.AI

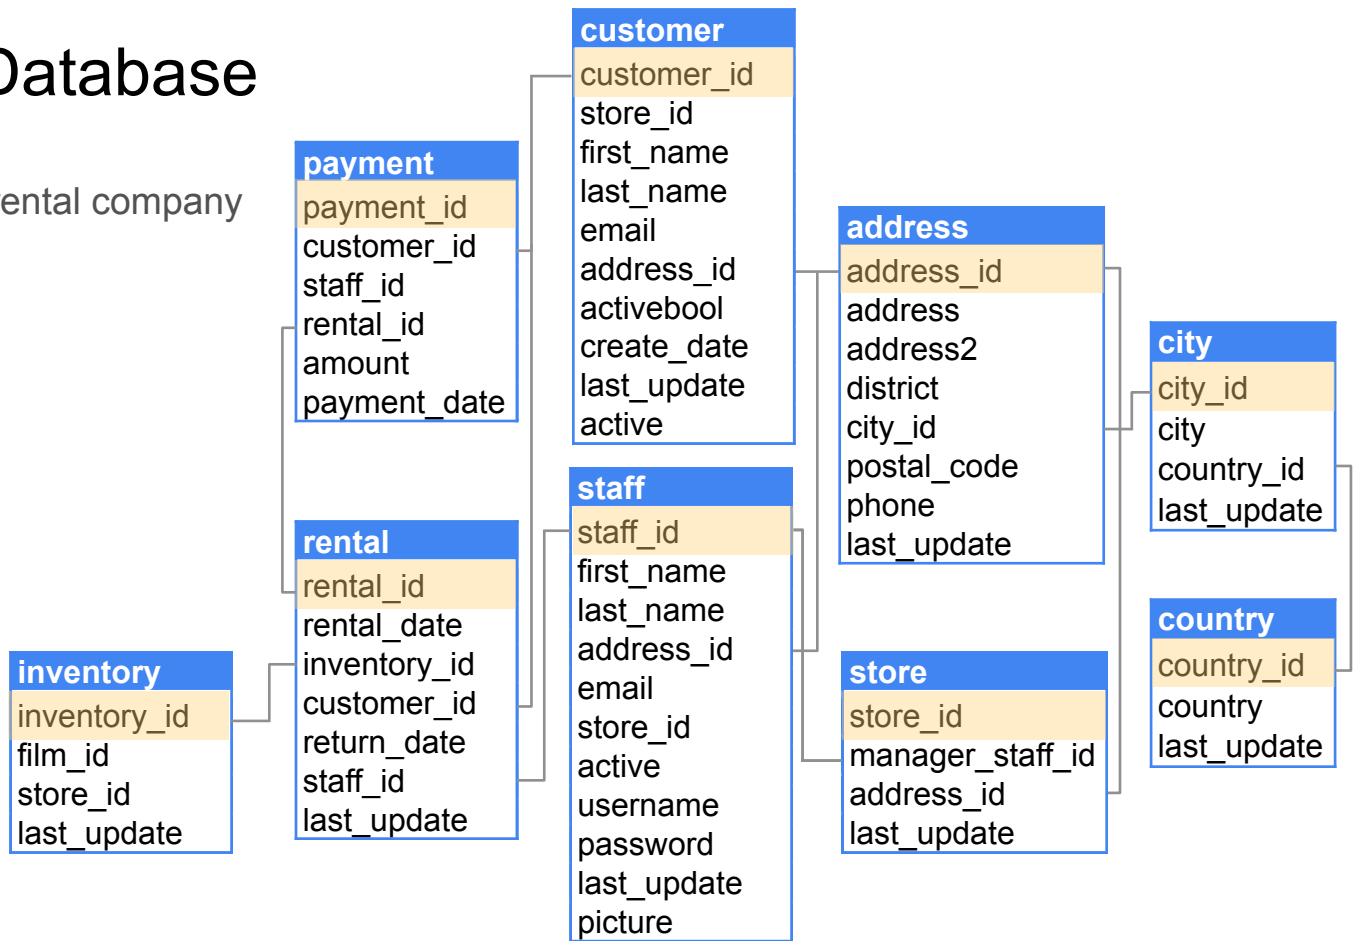
Introduction to Source Systems

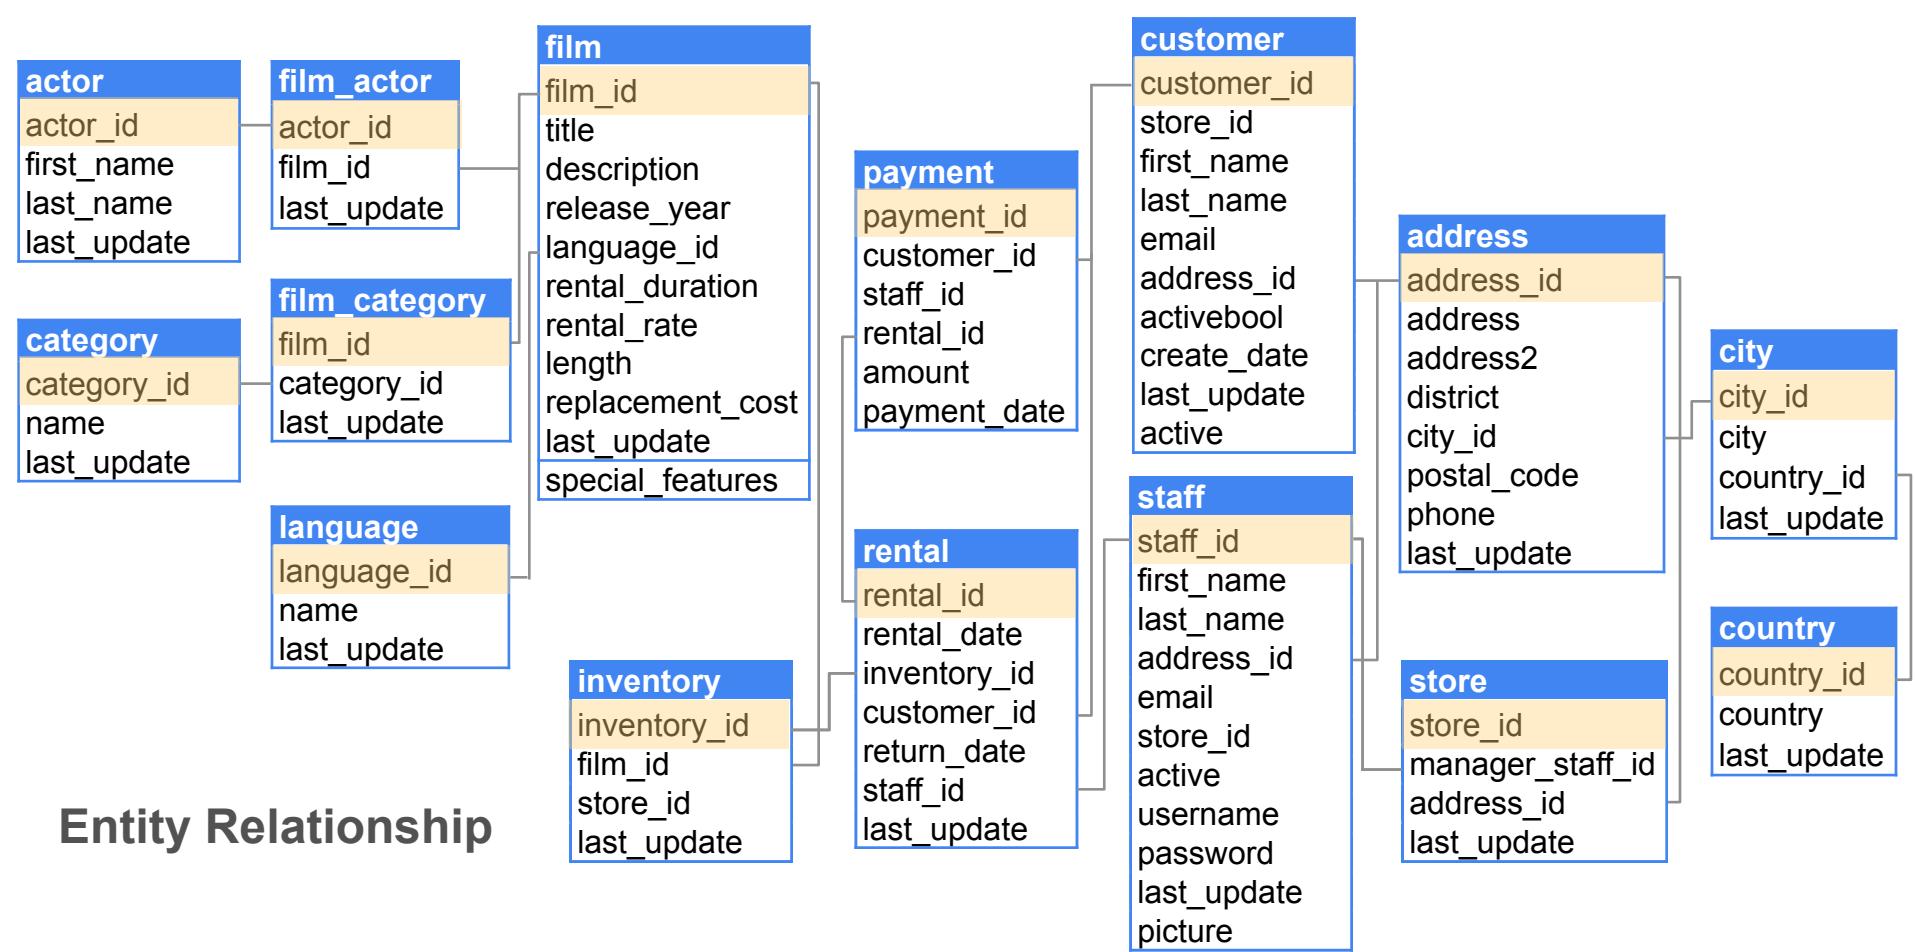
SQL Queries

The Relational Database

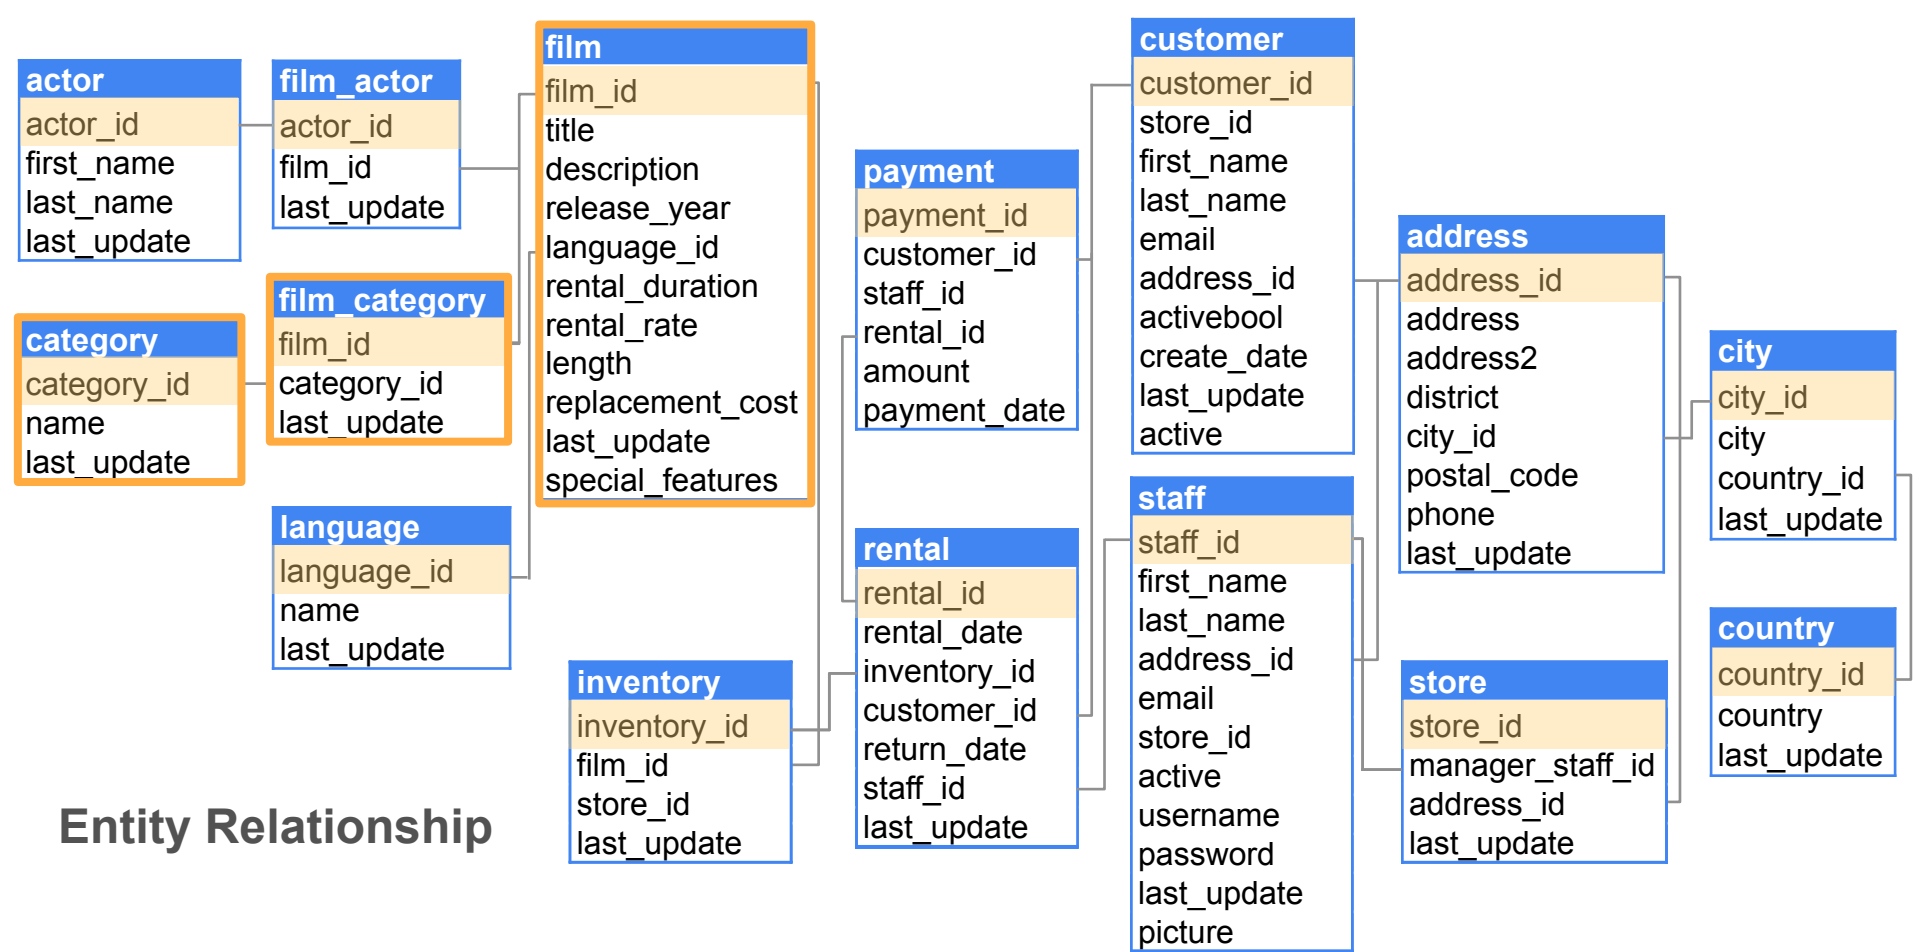
- Database for a fictitious DVD rental company called **Rentio**
- Database schema

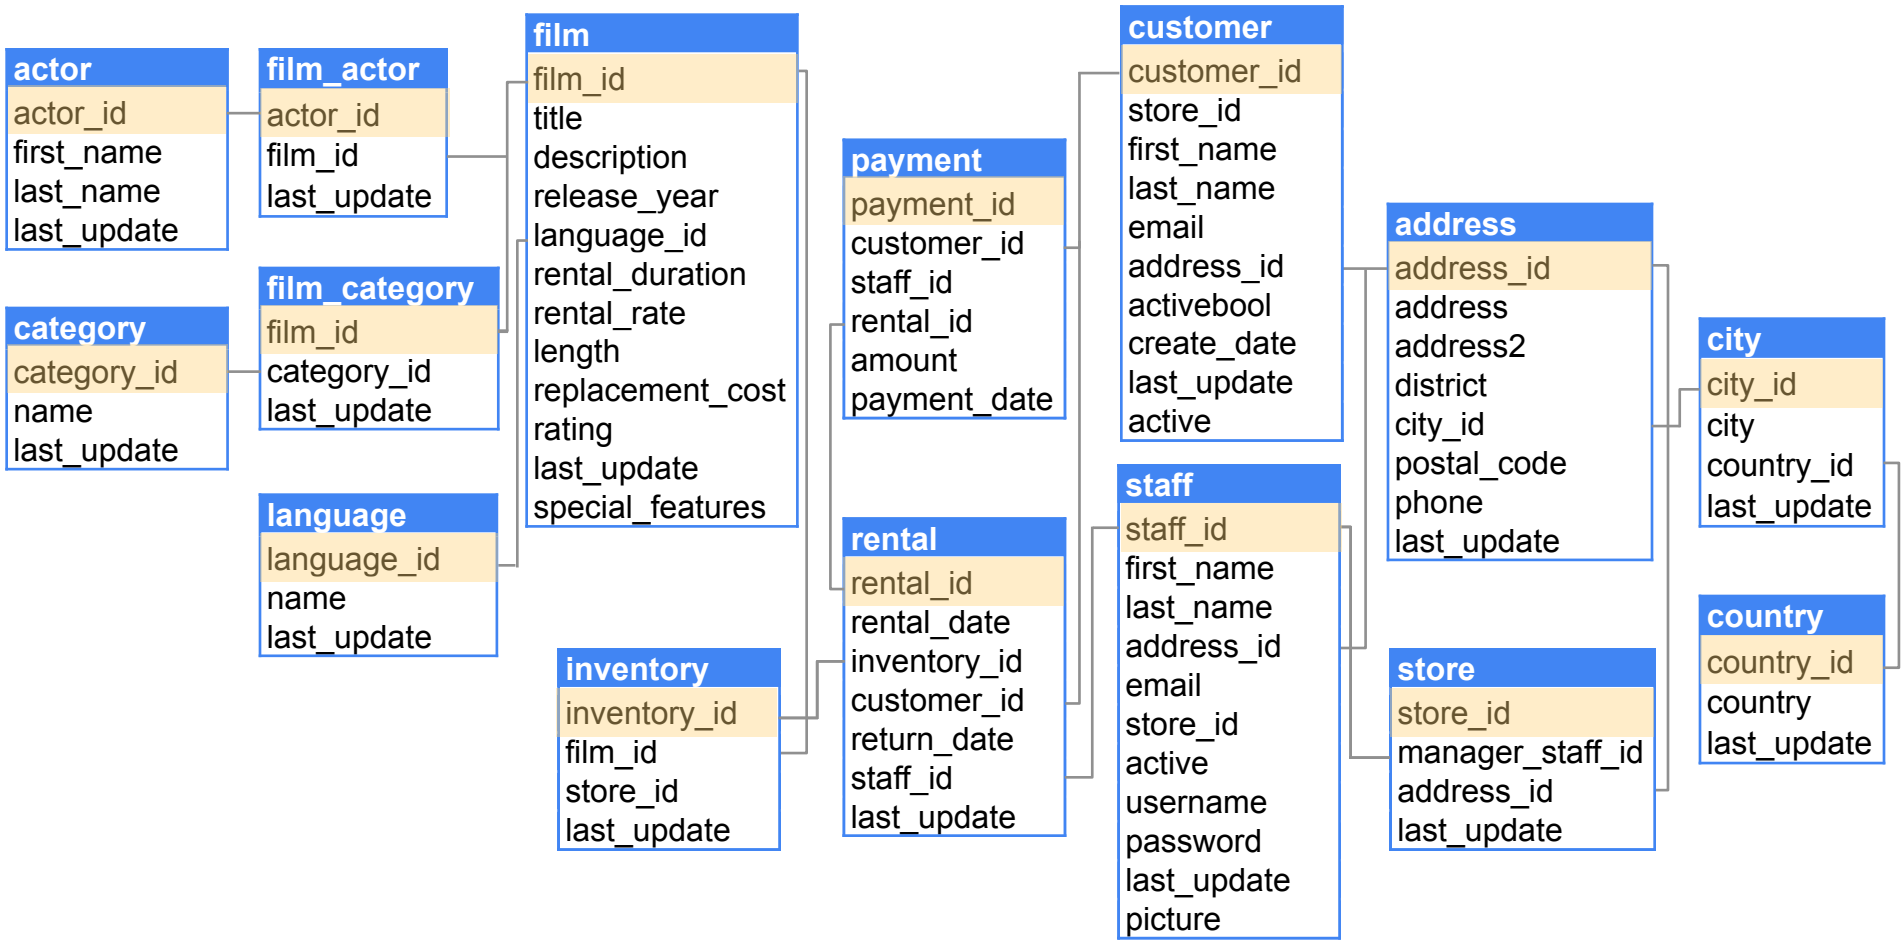
SQL queries
→
Answer business questions

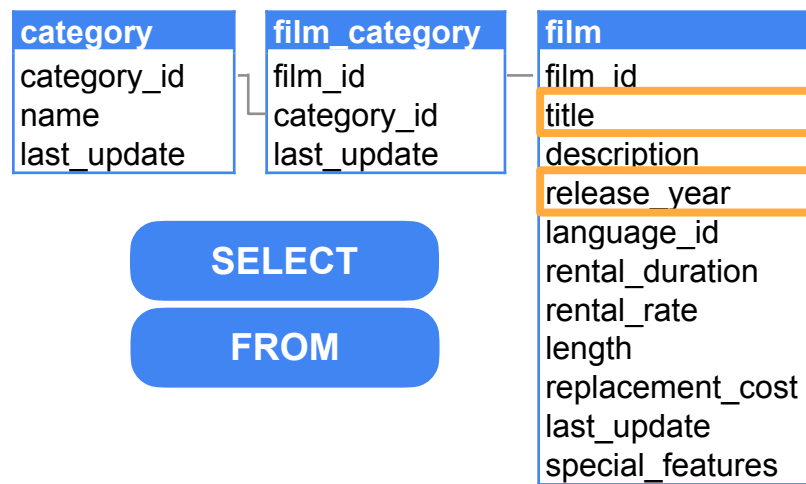




Entity Relationship



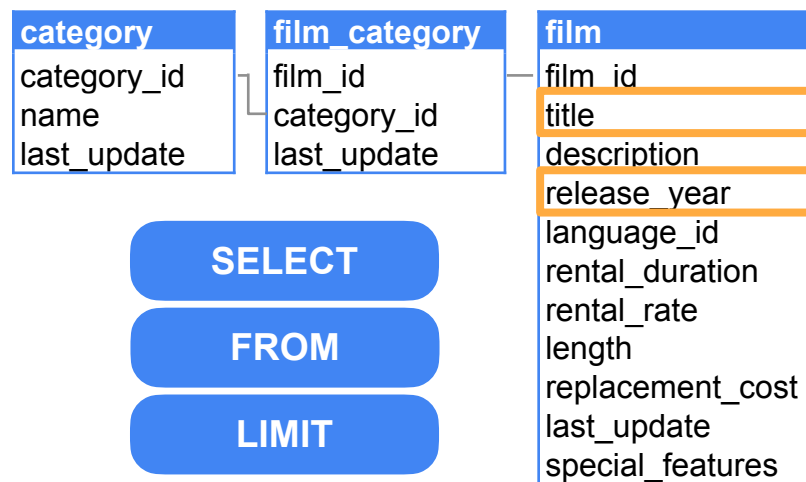




```
In [1]: %load_ext sql
        %sql mysql+pymysql://root:adminpwr@localhost:3306/sakila
```

```
In [ ]: %%sql
        |
```

```
In [ ]: 
```

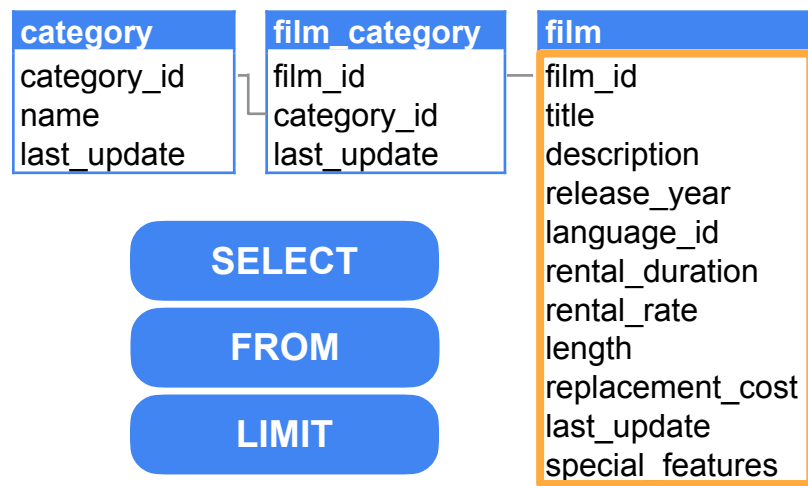
```
In [1]: %load_ext sql
        %sql mysql+pymysql://root:adminpwr@localhost:3306/sakila

In [2]: %%sql
        SELECT title, release_year
        FROM film

* mysql+pymysql://root:***@localhost:3306/sakila
1000 rows affected.

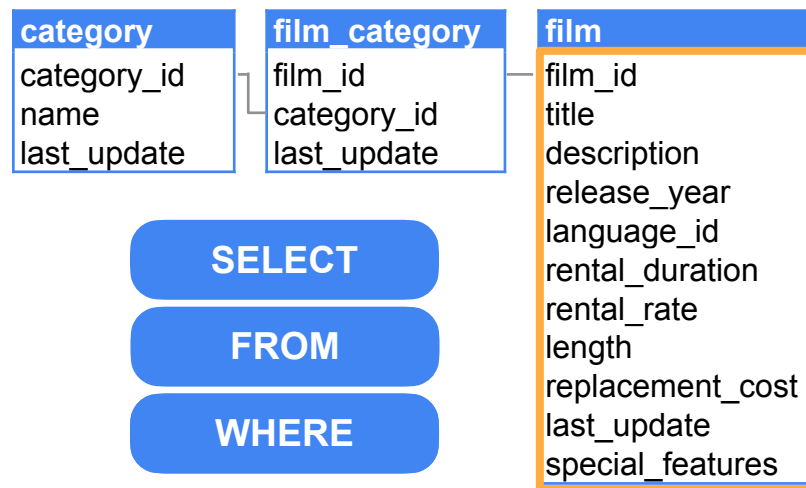
Out[2]:
```

	title	release_year
	ACADEMY DINOSAUR	2006
	ACE GOLDFINGER	2006
	ADAPTATION HOLES	2006
	AFFAIR PREJUDICE	2006
	AFRICAN EGG	2006
	AGENT TRUMAN	2006



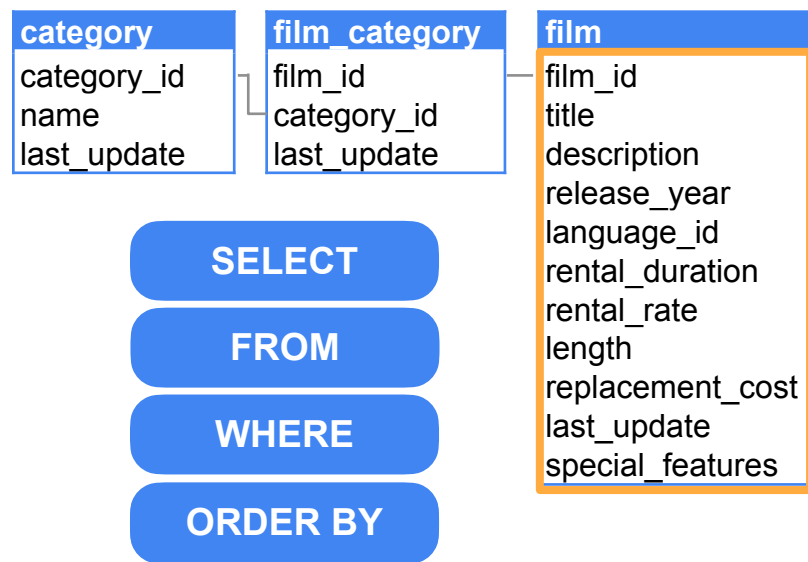
ADAPTATION HOLES	2006
AFFAIR PREJUDICE	2006
AFRICAN EGG	2006
AGENT TRUMAN	2006
AIRPLANE SIERRA	2006
AIRPORT POLLOCK	2006
ALABAMA DEVIL	2006
ALADDIN CALENDAR	2006

```
In [ ]: %%sql
```



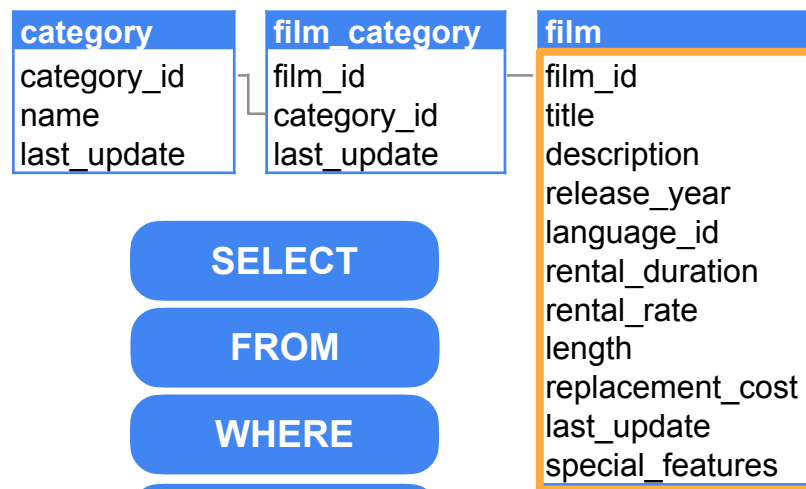
Exploring the films that are less than 60 minutes long.





```
In [4]: %%sql
        SELECT *
        FROM film
        WHERE length < 60
```

original_language_id	rental_duration	rental_rate	length	replacement_cost	rating
None	3	4.99	48	12.99	G



SELECT

FROM

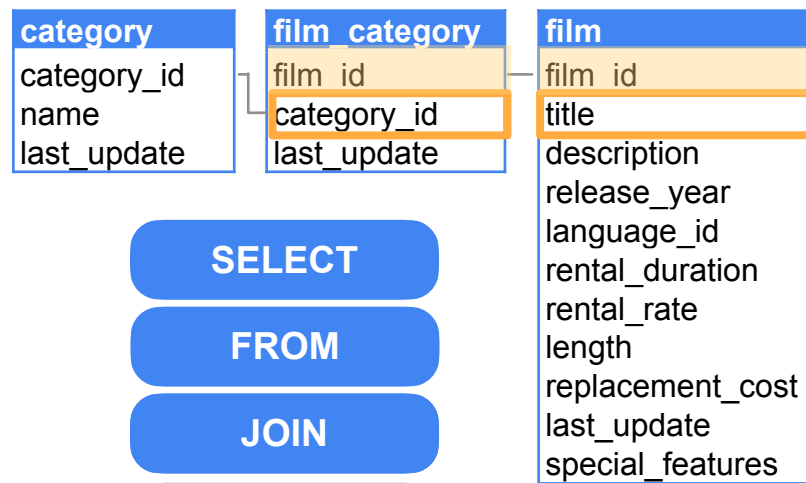
WHERE

ORDER BY

LIMIT

```
In [5]: %%sql
        SELECT *
        FROM film
        WHERE length < 60
        ORDER BY length
```

original_language_id	rental_duration	rental_rate	length	replacement_cost	rating
3	4.99	48	12.99	G	Trailers, Deleted Scer



SELECT

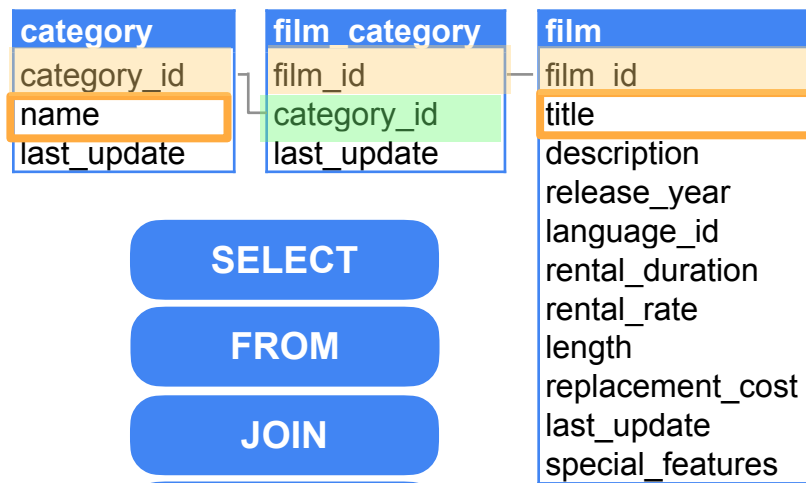
FROM

JOIN

ORDER BY

LIMIT

```
In [ ]: %%sql
        SELECT *
        FROM film
        WHERE length < 60
```



SELECT

FROM

JOIN

WHERE

ORDER BY

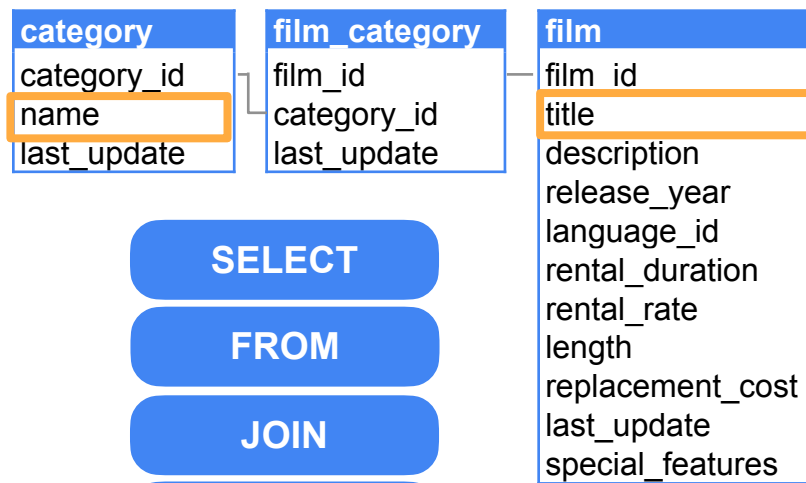
LIMIT

Monkey in Ancient India

```

In [5]: %%sql
        SELECT *
        FROM film
        JOIN film_category
        ON film.film_id = film_category.film_id
        WHERE length < 60
  
```

rating	special_features	last_update	film_id_1	category_id	last_update_1
G	Trailers, Deleted Scenes	2006-02-15 05:03:42	2	11	2006-02-15 05:07:09



SELECT

FROM

JOIN

WHERE

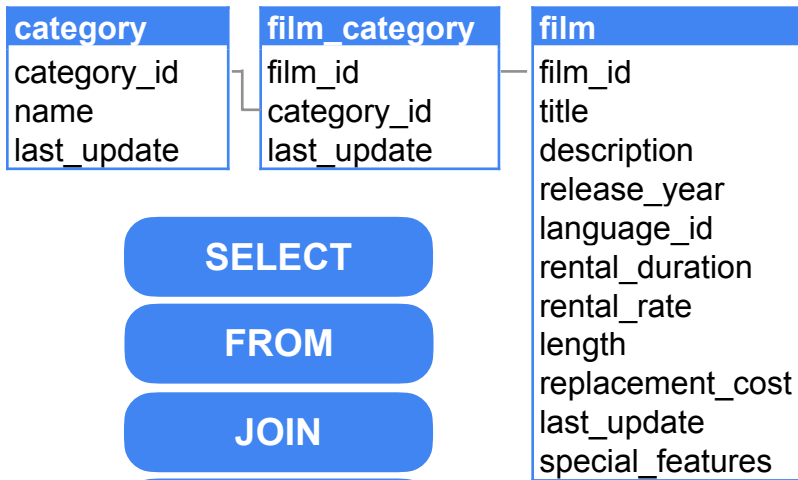
ORDER BY

LIMIT

Monkey in Ancient India

```
In [6]: %%sql
        SELECT *
        FROM film
        JOIN film_category
        ON film.film_id = film_category.film_id
        JOIN category
        ON film_category.category_id = category.category_id
        WHERE length < 60
```

date	film_id_1	category_id	last_update_1	category_id_1	name	last_update_2
2-15	2	11	2006-02-15	11	Horror	2006-02-15



SELECT

FROM

JOIN

WHERE

ORDER BY

LIMIT

INNER JOIN

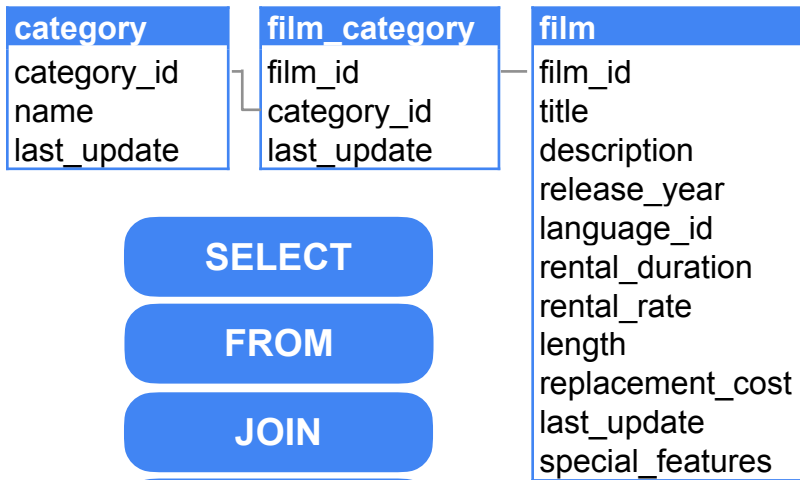
JOIN: combine the records from both tables that have a matching column value specified in the ON statement.

film has a row with film_id = 123

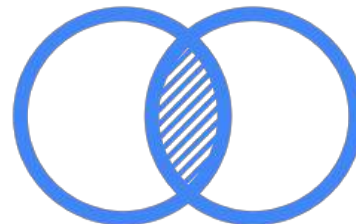
film_category does not have a row with film_id= 123



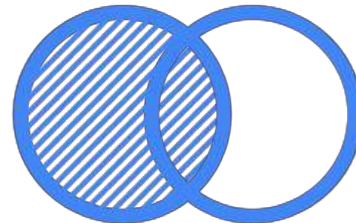
The row with film_id = 123 will not be in the join results



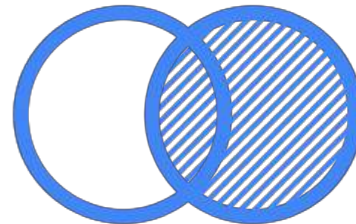
INNER JOIN



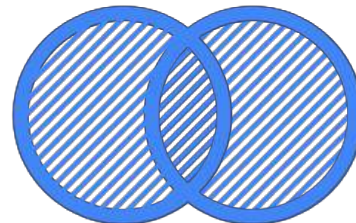
LEFT JOIN

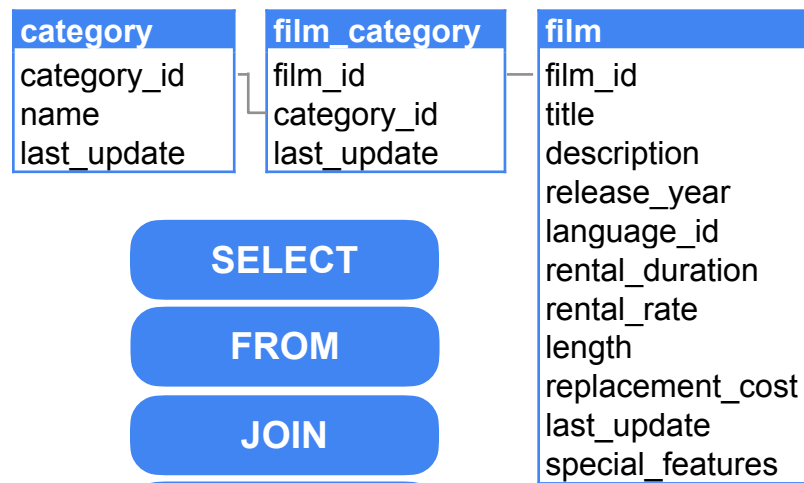


RIGHT JOIN



FULL JOIN



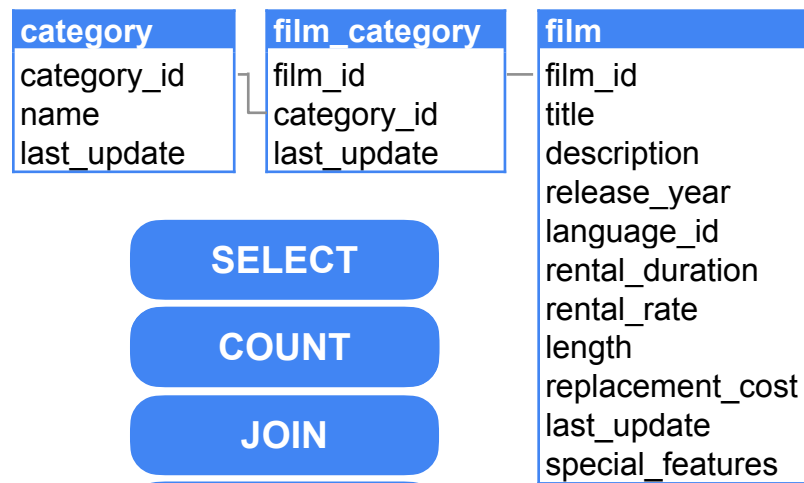


Monkey in Ancient India

```
In [7]: %%sql
        SELECT film.title, category.name
        FROM film
        JOIN film_category
        ON film.film_id = film_category.film_id
        JOIN category
        ON film_category.category_id = category.category_id
        WHERE length<60
```

Out[7]:

title	name
ACE GOLDFINGER	Horror
ADAPTATION HOLES	Documentary
AIRPORT POLLOCK	Horror
ALIEN CENTER	Foreign
ALTER VICTORY	Animation



SELECT

COUNT

JOIN

WHERE

GROUP BY

ORDER BY

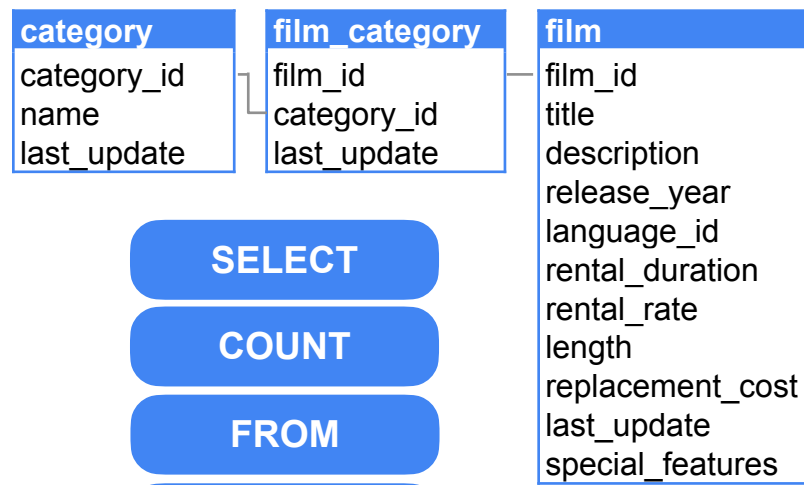
LIMIT

```
Monkey in Ancient India

In [7]: %%sql
        SELECT film.title, category.name
        FROM film
        JOIN film_category
        ON film.film_id = film_category.film_id
        JOIN category
        ON film_category.category_id = category.category_id
        WHERE length < 60

Out[7]:
```

title	name
ACE GOLDFINGER	Horror
ADAPTATION HOLES	Documentary
AIRPORT POLLOCK	Horror
ALIEN CENTER	Foreign
ALTER VICTORY	Animation



SELECT

COUNT

FROM

JOIN

WHERE

GROUP BY

ORDER BY

LIMIT

```
Monkey in Ancient India

In [7]: %%sql
        SELECT film.title, category.name
        FROM film
        JOIN film_category
        ON film.film_id = film_category.film_id
        JOIN category
        ON film_category.category_id = category.category_id
        WHERE length < 60
```

DAWN POND	Games
DEEP CRUSADE	Documentary
DESTINY SATURDAY	New
DIVORCE SHINING	Sports
DOCTOR GRAIL	Children
DOORS PRESIDENT	Animation

Common SQL Commands

SELECT

COUNT

FROM

JOIN

WHERE

GROUP BY

ORDER BY

LIMIT

Data Manipulation Operations

CREATE

INSERT
INTO

UPDATE

DELETE



DeepLearning.AI

Introduction to Source Systems

NoSQL Databases

NoSQL Databases

NoSQL

NoSQL Databases

No SQL

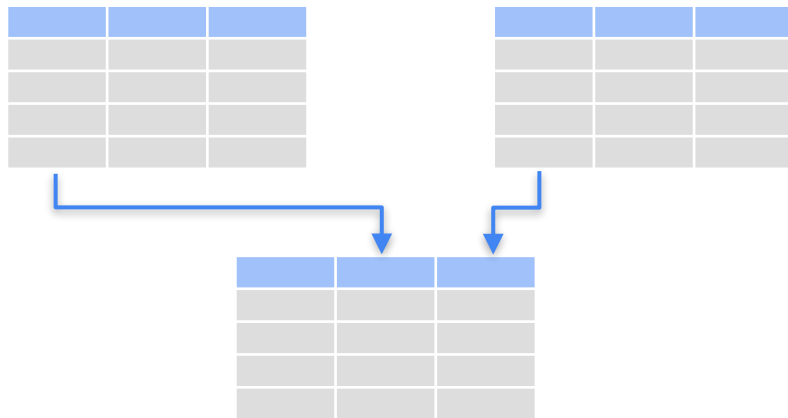
NoSQL Databases

Not Only SQL

Non-Relational Databases

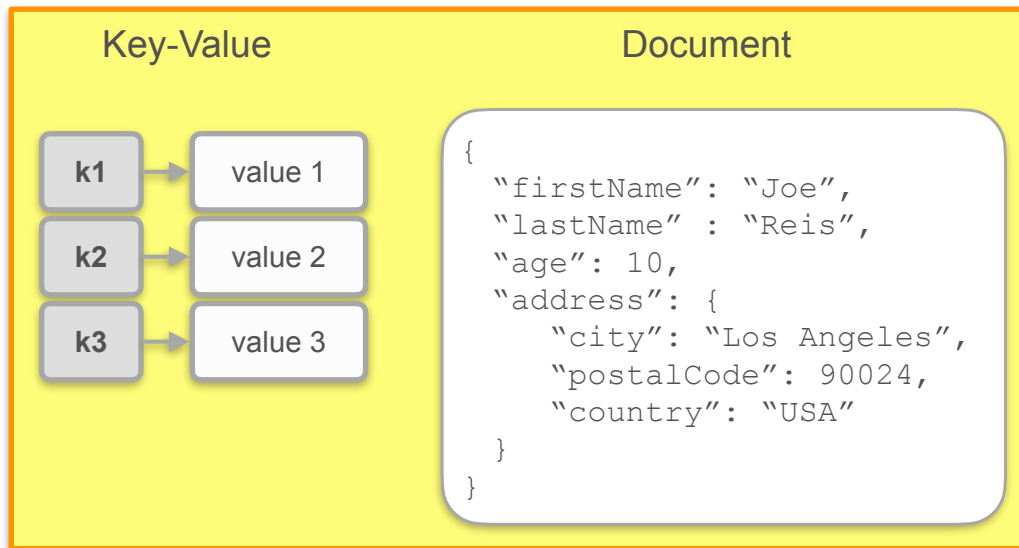
It can still support SQL or SQL-like query languages.

Relational Databases



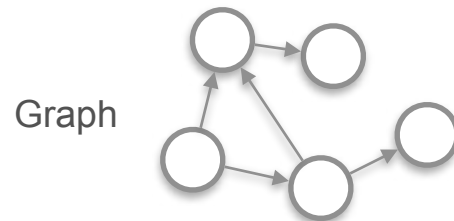
NoSQL Databases

Non-tabular structures

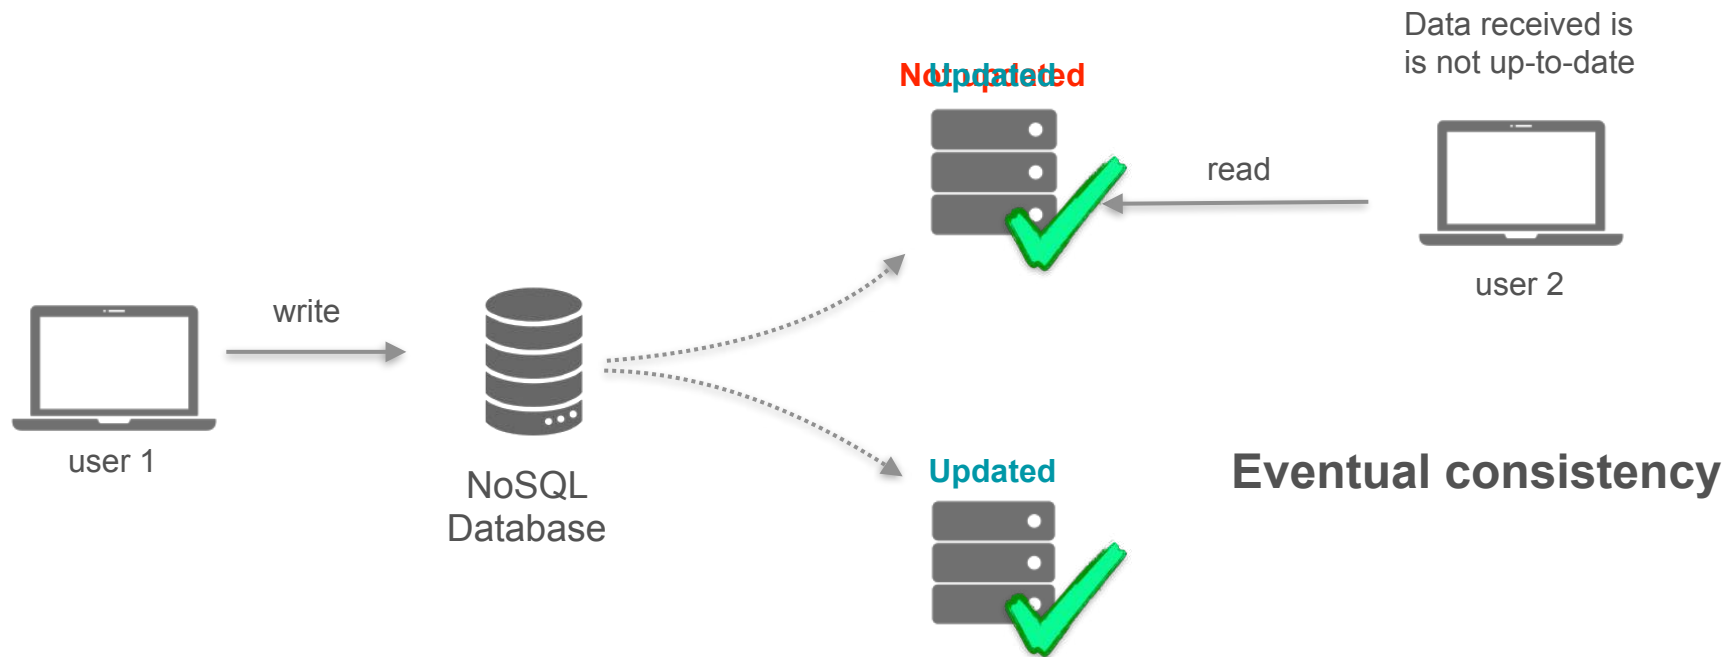


- No predefined schemas
- More flexibility when storing your data

Wide-Column



Horizontal Scaling



Consistency

NoSQL Databases	Relational Databases
<p data-bbox="355 452 768 495">Eventual Consistency</p> <ul data-bbox="258 554 842 710" style="list-style-type: none"><li data-bbox="258 554 602 598">• Speed is prioritized<li data-bbox="258 626 842 710">• System availability and scalability are important	<p data-bbox="1120 452 1491 495">Strong Consistency</p> <ul data-bbox="1031 576 1572 667" style="list-style-type: none"><li data-bbox="1031 576 1572 667">• Read data only when all nodes have been updated

NoSQL Databases

Not all NoSQL databases guarantee:

ACID compliance

Atomicity

Consistency

Isolation

Durability



Specialized Query Language

Example of NoSQL Data

```
{  
  "id": 1,  
  "key": "Blender",  
  "qty": 6,  
  "sku": "b32"  
}
```

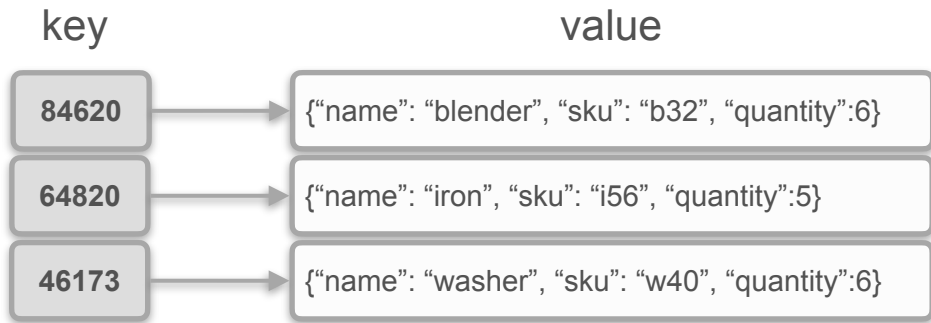
Query

```
db.products.find({qty: {$gt: 4}})
```

[Ref: AWS docs](#)

Types of No-SQL Databases

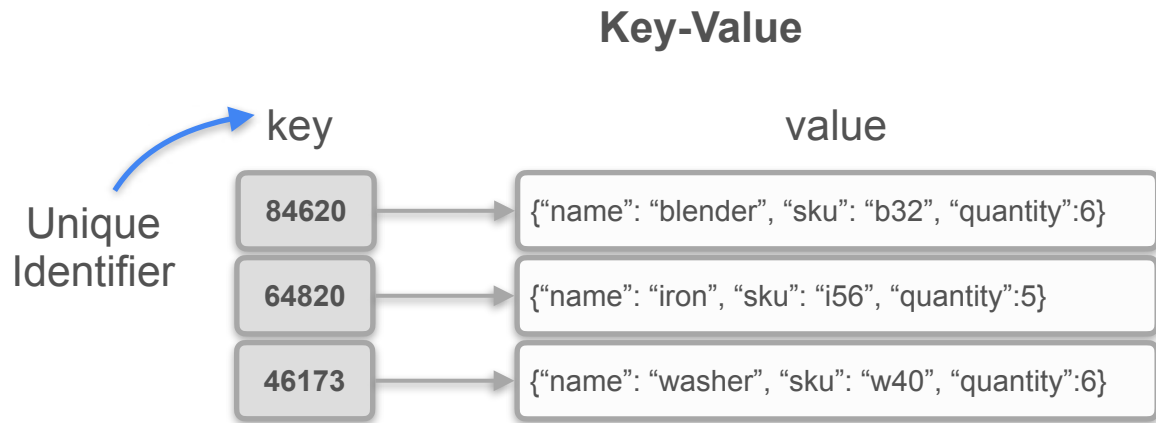
Key-Value



Document

```
{  
  "firstName": "Joe",  
  "lastName" : "Reis",  
  "age": 10,  
  "address": {  
    "city": "Los Angeles",  
    "postalCode": 90024,  
    "country": "USA"  
  }  
}
```


Key-Value Database



Fast lookup: such as caching user session data

- viewing different products
- adding items to the shopping cart
- checking out



Document Database

Collection *(Like a table)*

```
{
  "users" : [
    {
      "id": 1234,
      "name": {
        "first": "Joe",
        "last": "Reis"
      },
      "favorite_bands" : ["AC/DC", "Slayer", "WuTang Clan", "Action Bronson" ]
    },
    {
      "id": 1235,
      "name": {
        "first": "Matt",
        "last": "Housley"
      },
      "favorite_bands" : ["Dave Matthews Band", "Creed", "Nickelback"]
    }
  ]
}
```

**Single users
Documents**
(Like a row)

Document Database

```
{
  "users" : [
    {
      "id": 1234,
      "name": {
        "first": "Joe",
        "last": "Reis"
      },
      "favorite_bands" : ["AC/DC", "Slayer", "WuTang Clan", "Action Bronson" ]
    },
    {
      "id": 1235,
      "name": {
        "first": "Matt",
        "last": "Housley"
      },
      "favorite_bands" : ["Dave Matthews Band", "Creed", "Nickelback"]
    }
  ]
}
```

- Easy to retrieve all the information about a user (locality)
- Document stores don't support joins
- Flexible schema

user_id	band_id	band_id	band_name
1234	1	1	AC/DC
1234	2	2	Slayer
1234	5	3	Creed
1234	6	4	Nickelback
1235	7	5	Wutan Clan
1235	3	6	Action Bronson
1235	4	7	Dave Matthews Band


user_id	first_name	last_name
1234	Joe	Reis
1235	Matt	Housely

Fixed schema

Document Database

Use cases

- Content management
- Catalogs
- Sensor readings

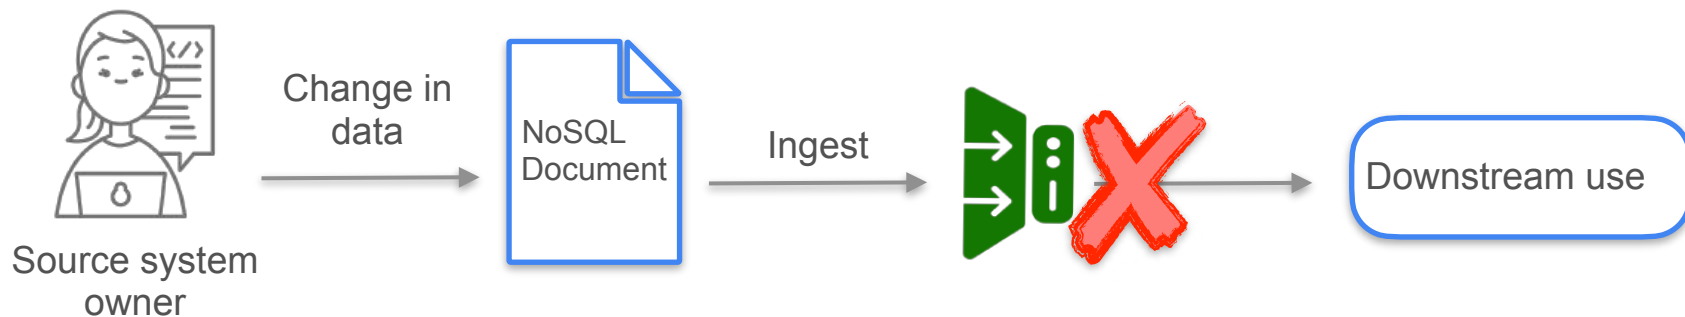


```
{
  "iot" : [
    {
      "id": 24,
      "interaction": "some_interaction",
      "device": "my_device",
      "sensor_reading": 34
    }
  ]
}
```

Flexible Schema

Document Database

Document databases become absolute nightmares to manage and query.





DeepLearning.AI

Introduction to Source Systems

Database ACID Compliance

OLTP Systems

OLTP



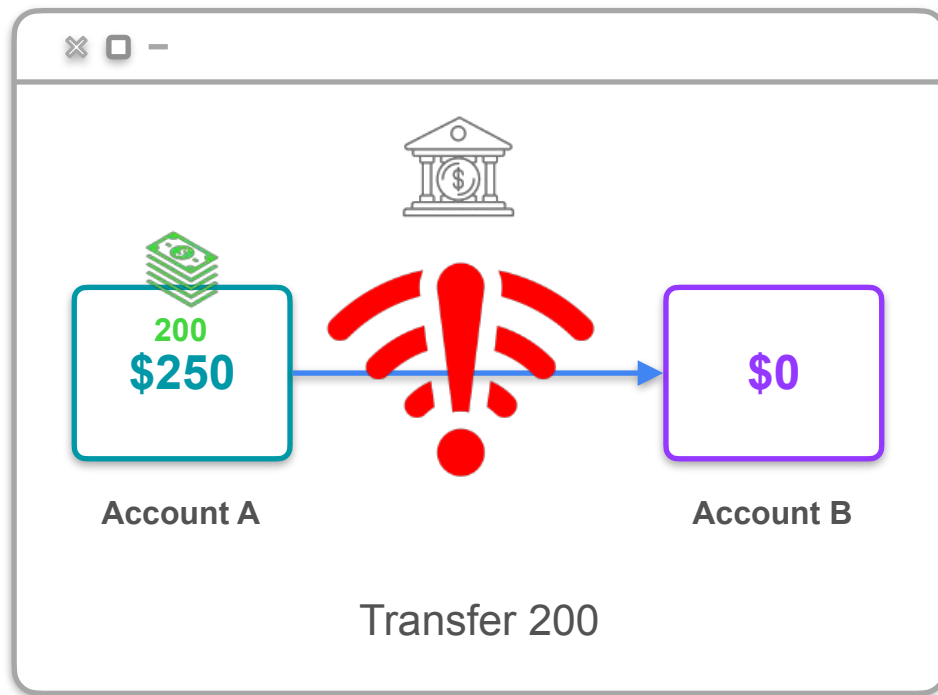
Online Transaction Processing

Support very high transaction rates (bank account balances, online orders)

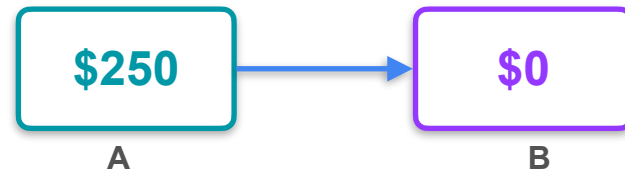
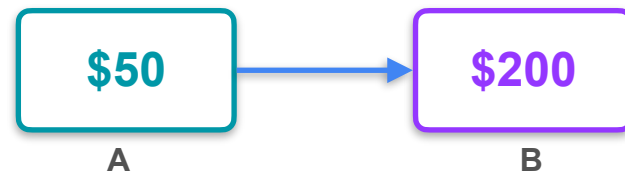
ACID Compliance

Relational Databases	NoSQL Databases
ACID compliant	Not ACID compliant by default
A tomicity	
C onsistency	
I solation	
D urability	
They help ensure transactions are processed reliably and accurately in an OLTP system.	

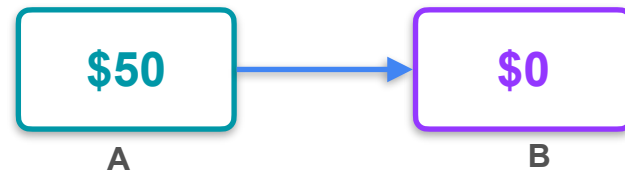
ACID Compliance



You'd be hoping



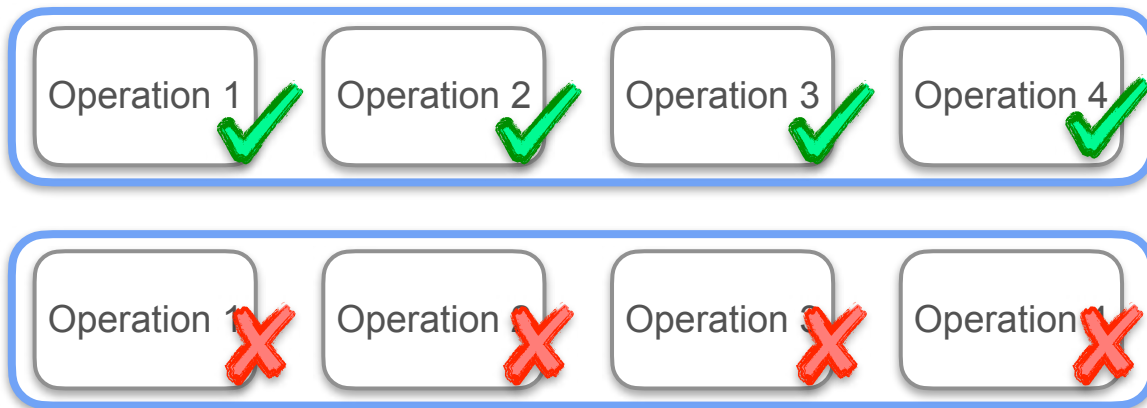
But not



Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

A transaction



Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

A transaction:
placing an order

Deducting the total cost from
the customer's account



Updating the inventory to
reflect the purchased item



Both operations must happen as a single transaction

Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

Consistency

Any changes to the data made within a transaction follow the set of rules or constraints defined by the database schema.

id	product_name	quantity
1	blender	1



Buy 2 blenders



Transaction



id	product_name	quantity
1	blender	-1



Rule: stock level ≥ 0

Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

Consistency

Any changes to the data made within a transaction follow the set of rules or constraints defined by the database schema.

ACID compliance

A

Atomicity

C

Consistency

I

Isolation

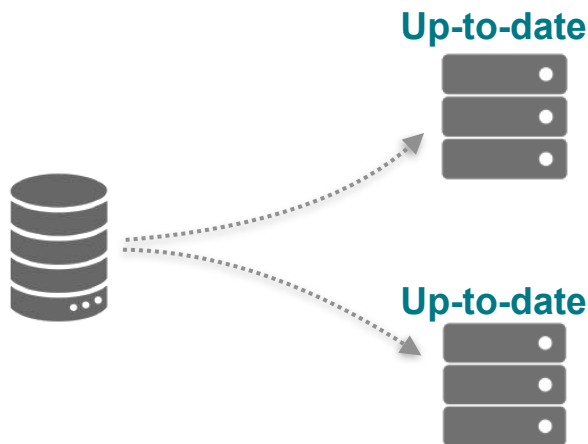
D

Durability



Strong Consistency

All nodes provide the same up-to-date



Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

Consistency

Any changes to the data made within a transaction follow the set of rules or constraints defined by the database schema.

Isolation

Each transaction is executed independently in sequential order.

id	product_name	quantity
1	blender	5

Transaction

Buy 5 blenders



Transaction

Buy 5 blenders



Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

Consistency

Any changes to the data made within a transaction follow the set of rules or constraints defined by the database schema.

Isolation

Each transaction is executed independently in sequential order.

id	product_name	quantity
1	blender	5

Transaction

Buy 5 blenders



Transaction

Buy 10
blenders



Atomicity

Atomicity ensures that transactions are **atomic**, treated as a single, indivisible unit.

Consistency

Any changes to the data made within a transaction follow the set of rules or constraints defined by the database schema.

Isolation

Each transaction is executed independently in sequential order.

Durability

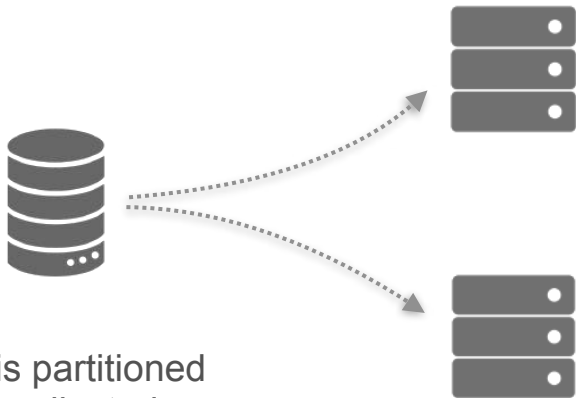
Once a transaction is completed, its effects are permanent and will survive any subsequent system failures.

*Essential for maintaining the
reliability of the database*



ACID Compliance

The ACID principles guarantee that a database will maintain a consistent picture of the world.



Data is partitioned
or replicated

Strong Consistency

- Data is consistent across the entire network
- Key feature of relational databases that ensures ACID



DeepLearning.AI

Introduction to Source Systems

**Lab Walkthrough -
Interacting with Amazon
DynamoDB NoSQL Database**

Interacting with Amazon DynamoDB



Amazon DynamoDB

Apply some **C**reate, **R**ead, **U**ppdate and **D**eleete
(**CRUD**) operations

In this video,

- Overview of DynamoDB features
- Data you will work on
- DynamoDB methods that you will use to apply **CRUD** operations



Amazon DynamoDB

Key-value Database

Key-value Items

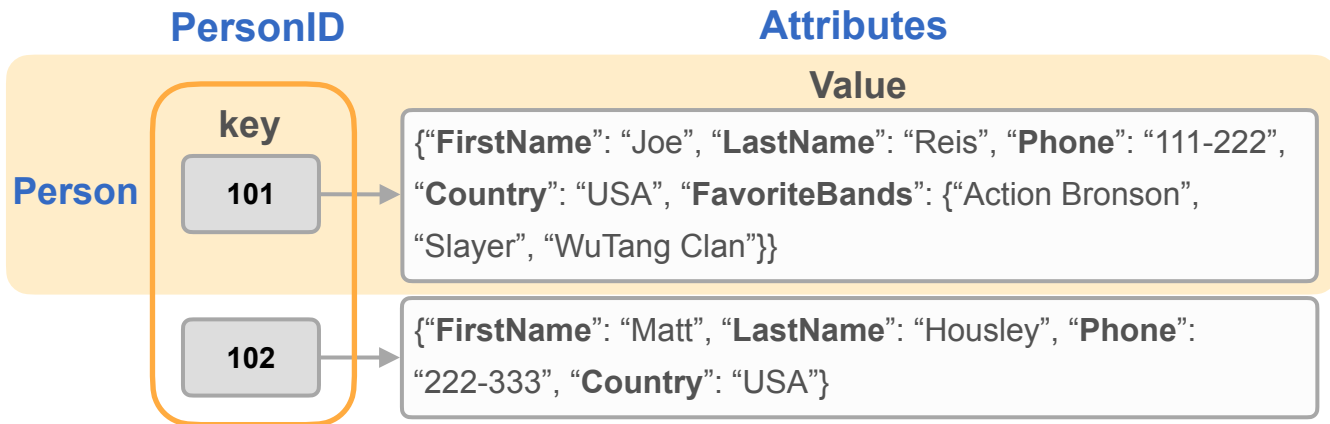


Table

- Row: attributes of one item
- Uniquely identified by the item's key.
- **Simple Primary Key:** partition key
- **Composite Primary Key:**

Simple
primary
key

Key: PersonID	Attributes				
	FirstName	LastName	Phone	Country	FavoriteBands
101	Joe	Reis	111-222	USA	{“Action Bronson”, “Slayer”, “WuTang Clan”}
102	FirstName	LastName	Phone	Country	
	Matt	Housley	222-333	USA	





Amazon DynamoDB

Key-value Database

Key-value Items



Table

- Row: attributes of one item
- Uniquely identified by the item's key.
- **Simple Primary Key:** partition key
- **Composite Primary Key:** partition key & sort key

Composite Primary Key		Attributes				
Partition Key	Sort Key					
OrderID	ItemNum	Price	Quantity	ProductType	ISBN	Title
1234	Item1	10	1	Book	45679	Data
1234	Item2	50	1	Bike	AZY	Black
1235	Item1	23	4	ProductCode 23697		
1235	Item2	1200	2	ProductType	Brand	
				Laptop	XYZ	

Schema-less: Each item can have its own distinct attributes.

Simple
primary
key

Key: PersonID	Attributes				
101	FirstName	LastName	Phone	Country	FavoriteBands
	Joe	Reis	111-222	USA	{“Action Bronson”, “Slayer”, “WuTang Clan”}
102	FirstName	LastName	Phone	Country	
	Matt	Housley	222-333	USA	

Interacting with Amazon DynamoDB



Table

Interact with the tables using Python

Boto3

AWS Software Development Kit (SDK) for Python
Allows you to create and configure AWS services using Python



Table



Table



Table



Boto3 1.34.144 documentation

🔍 Search

Feedback

Do you have a suggestion to improve this website or boto3?
[Give us feedback.](#)

[Quickstart](#)

[A Sample Tutorial](#)

[Code Examples](#)



[Developer Guide](#)



[Security](#)

[Available Services](#)



[Core References](#)



Boto3 documentation



You use the AWS SDK for Python (Boto3) to create, configure, and manage AWS services, such as Amazon Elastic Compute Cloud (Amazon EC2) and Amazon Simple Storage Service (Amazon S3). The SDK provides an object-oriented API as well as low-level access to AWS services.



Note

Documentation and developers tend to refer to the AWS SDK for Python as "Boto3," and this documentation often does so as well.

Quickstart

- [Quickstart](#)
 - [Installation](#)
 - [Configuration](#)
 - [Using Boto3](#)
- [A Sample Tutorial](#)
 - [SQS](#)
 - [Creating a queue](#)
 - [Using an existing queue](#)
 - [Sending messages](#)
 - [Processing messages](#)
- [Code Examples](#)
 - [Amazon CloudWatch examples](#)
 - [Amazon DynamoDB](#)

Interacting with Amazon DynamoDB



Table



Table



Table



Table

Interact with the tables using Python

Boto3

AWS Software Development Kit (SDK) for Python

Allows you to create and configure AWS services using Python

```
import boto3  
  
client = boto3.client('dynamodb')
```

Create

create_table

Read

scan
get_item
query

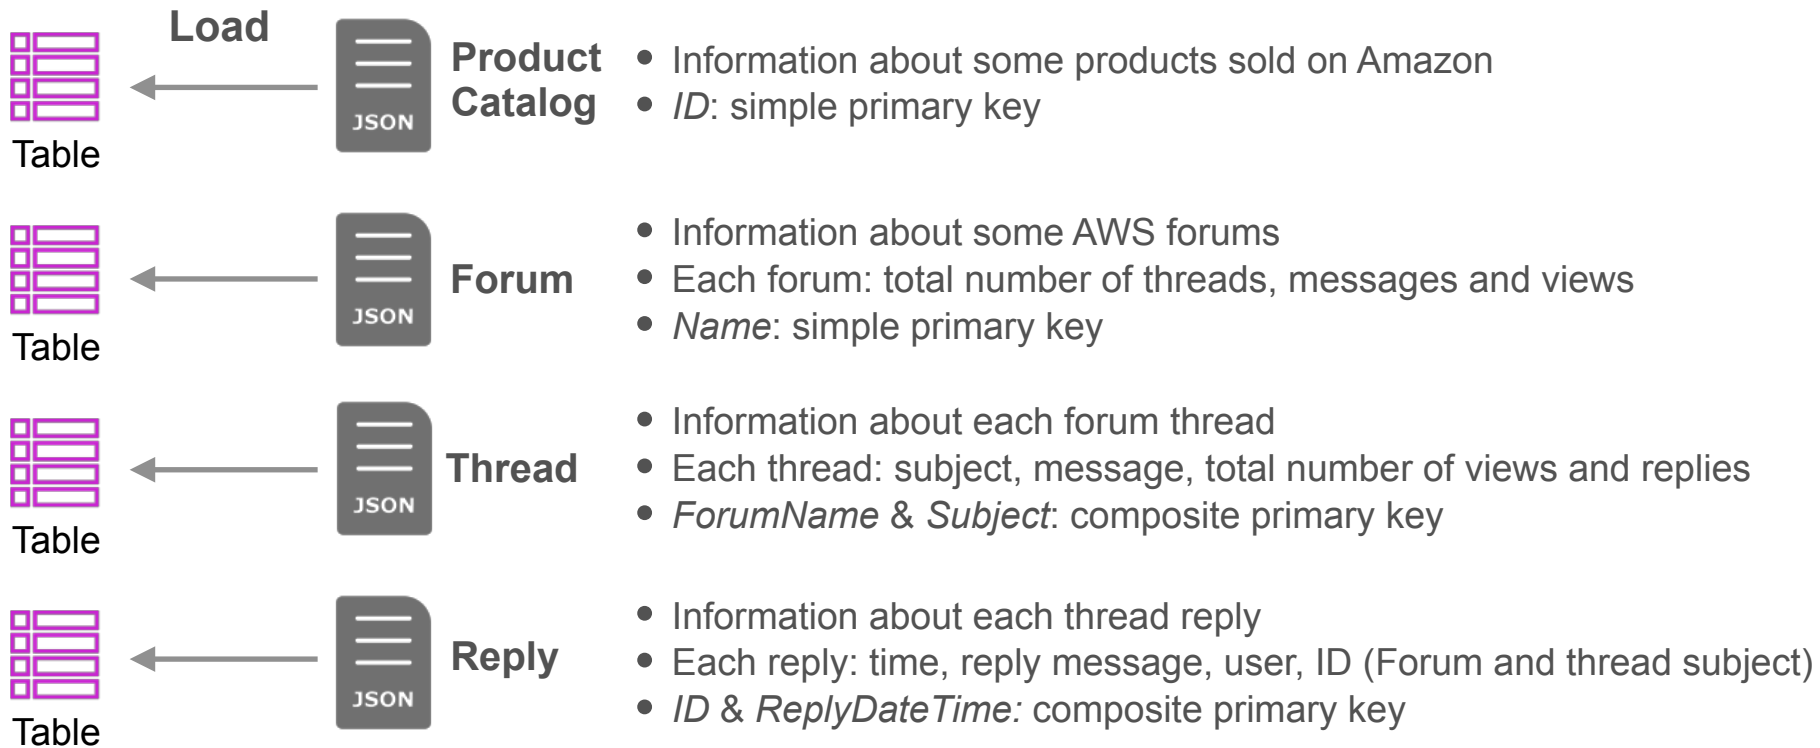
Update

put_item
write_batch_items
update_item

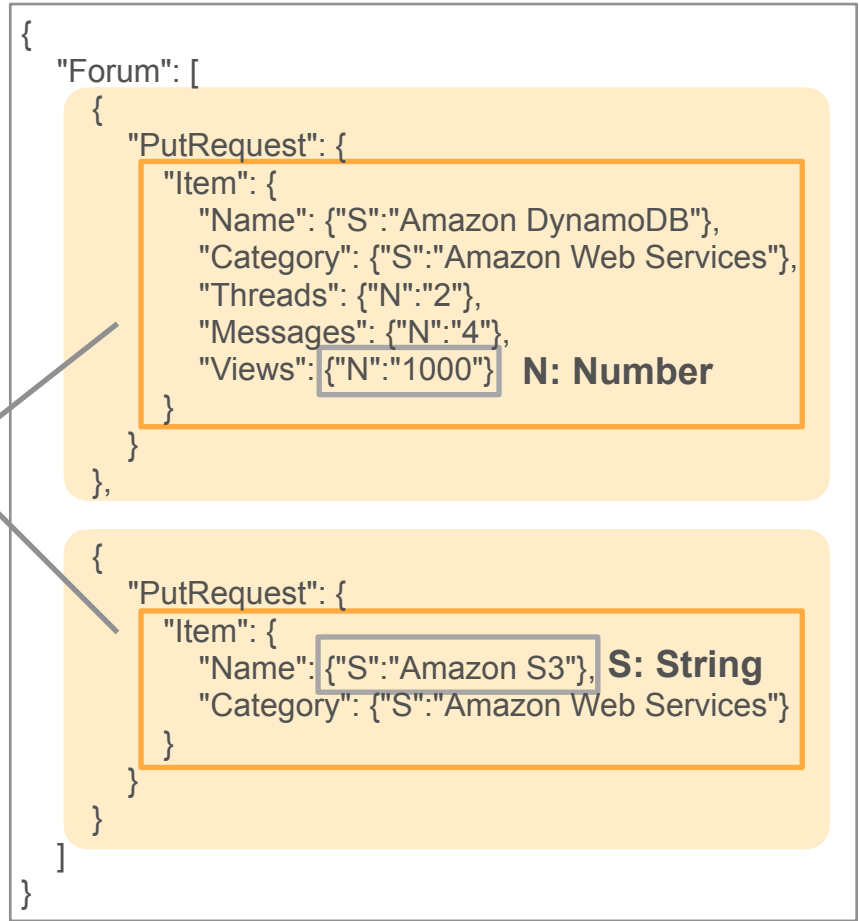
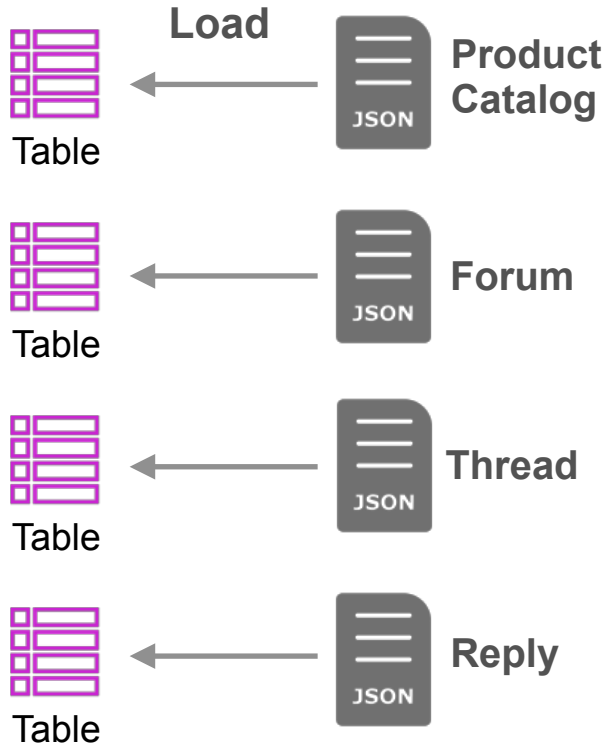
Delete

delete_item

Data



Data





DeepLearning.AI

Introduction to Source Systems

Object Storage

Object Storage



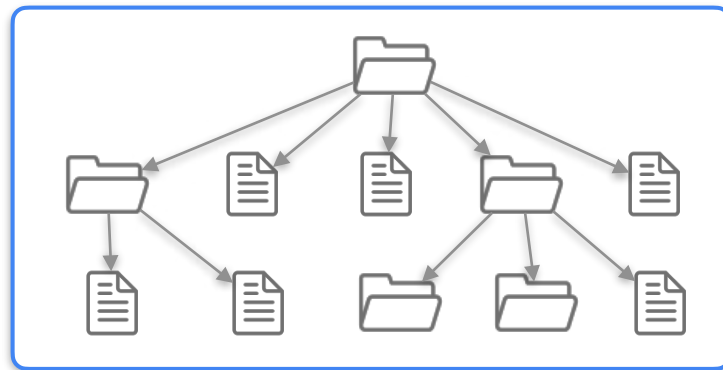
files

Object Storage



No hierarchy!

Traditional File System Hierarchy



Object Storage



Amazon S3

Amazon S3 > Buckets > mybucket1275

mybucket1275 [Info](#)

[Objects](#) | [Properties](#) | [Permissions](#) | [Metrics](#) | [Management](#) | [Access Points](#)

Objects (3) [Info](#)

[Refresh](#) [Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions](#) **Create folder**

[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	output-2024-03-01/	Folder	-	-	-
<input type="checkbox"/>	output-2024-03-02/	Folder	-	-	-
<input type="checkbox"/>	output-2024-03-03/	Folder	-	-	-

Object Storage



files

Object Storage



No hierarchy!

Object Storage



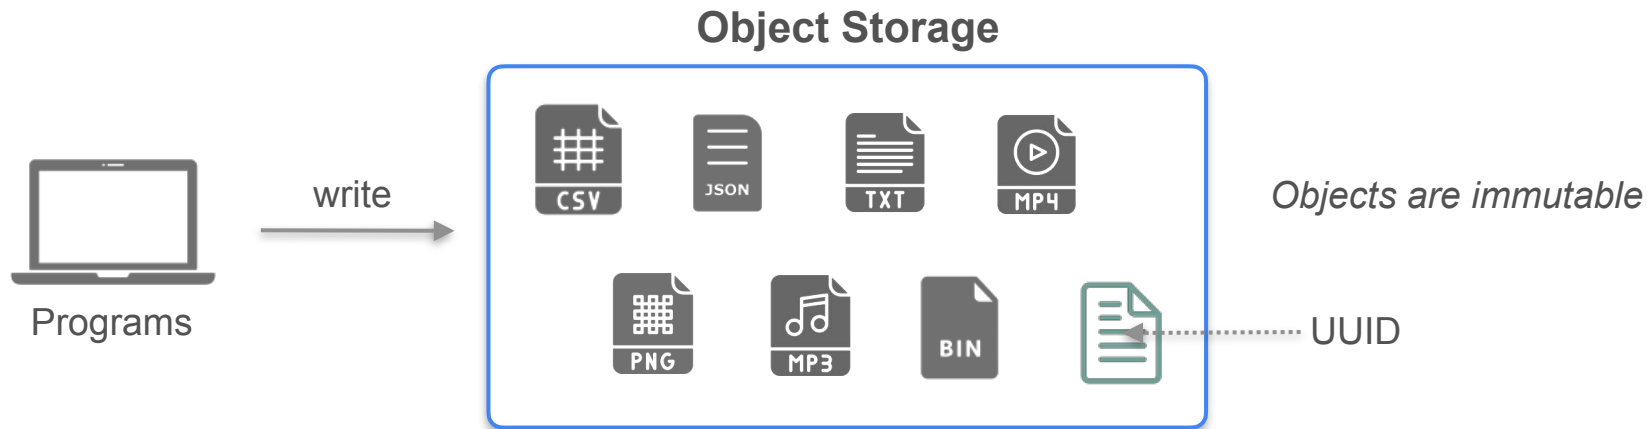
files

Object Storage



- Storing semi-structured and unstructured data
- Serving data for training machine learning models

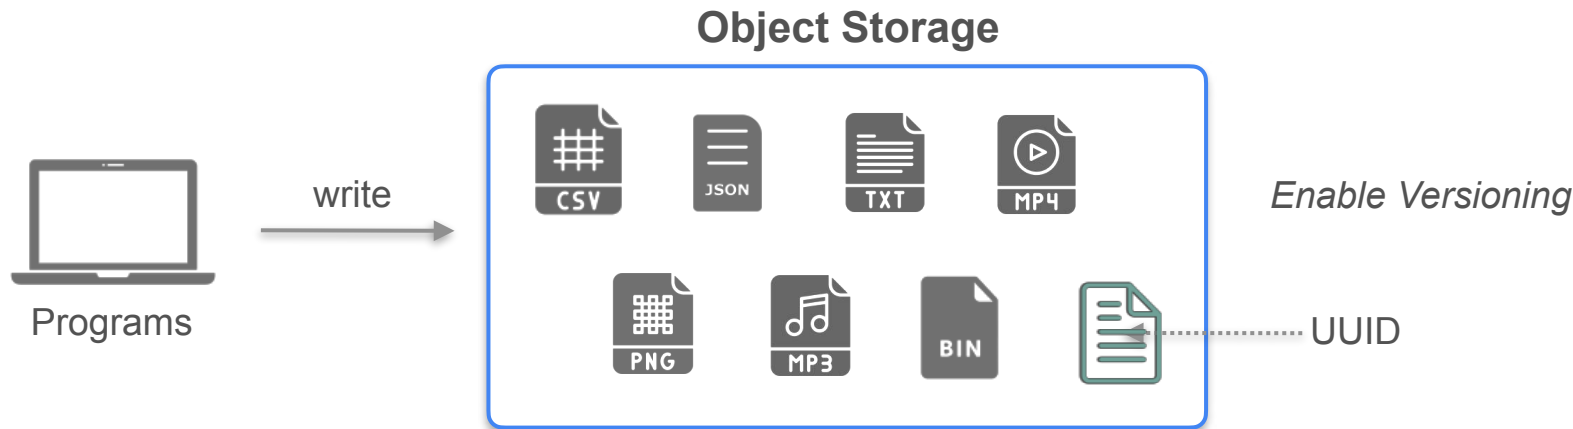
Object Storage



For each object,

- Universal Unique Identifier or UUID (key)
- Metadata: creation date, file type, owner

Object Storage

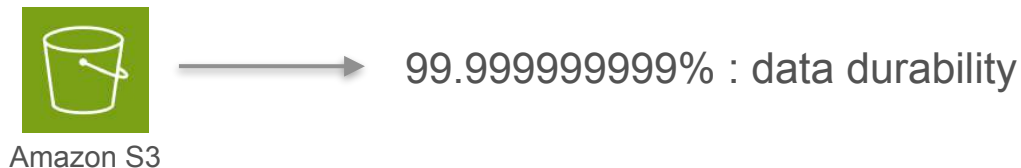


For each object,

- Universal Unique Identifier or UUID (key)
- Metadata: creation date, file type, owner, version

Why Use Object Storage?

- Store files of various data formats without a specific file system structure
- Easily scale out to provide virtually limitless storage space
- Replicate data across several availability zones



- Cheaper than other storage options

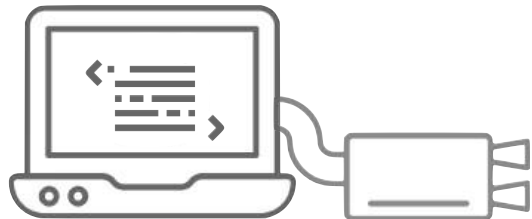


DeepLearning.AI

Introduction to Source Systems

Logs

Logs



Software
Application

Logs

01-01-2025:10.30	67945	success	user added a product x to their cart
01-01-2025:10.32	38910	fail	invalid values typed for product quantity
01-01-2025:10.38	17462	fail	customer table corrupted

Exhaust / Byproduct

Monitoring or Debugging a system

- User activity:
 - Signing in
 - navigating to a particular page
- An update to a database
- An error from a procedure

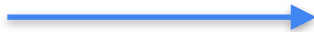
Log

An append-only sequence of records ordered by time, capturing information about events that occur in systems.

Rich data source

Web Server Logs

Detailed user activity data



Analysis of user behavior patterns



Database System Logs



Track changes in source database

Security System Logs

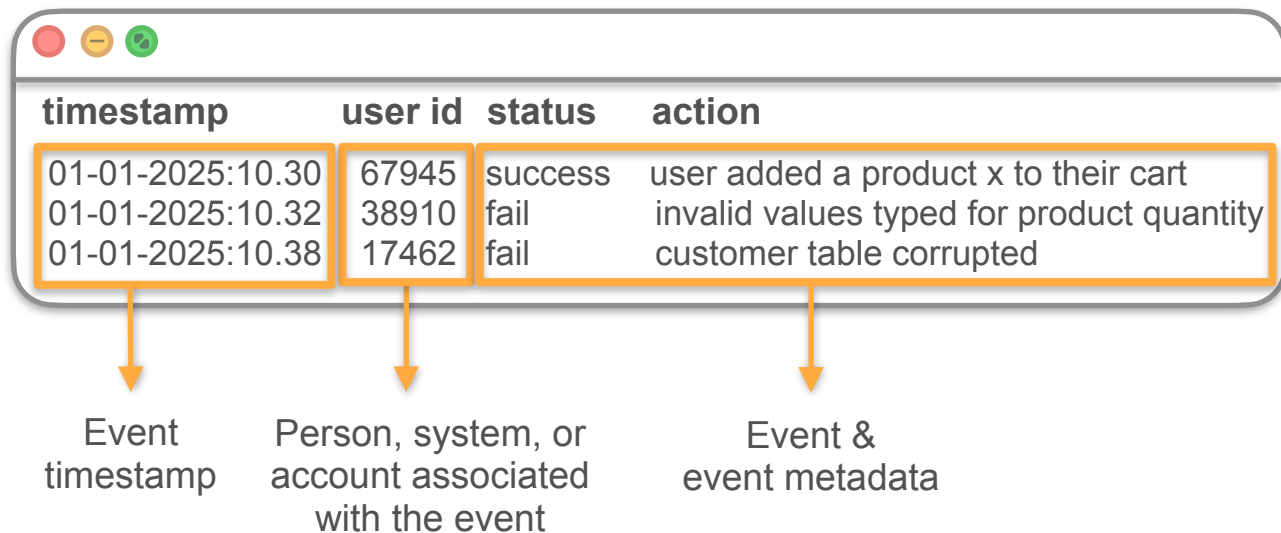


Machine Learning *anomaly detection*

Downstream use cases

Log

An append-only sequence of records ordered by time, capturing information about events that occur in systems.



The diagram shows a log window with a table of events. The table has four columns: timestamp, user id, status, and action. Three rows of data are shown, each with a timestamp, a user id, a status, and an action. Arrows point from the first three columns to labels below the table: 'Event timestamp' from the timestamp column, 'Person, system, or account associated with the event' from the user id column, and 'Event & event metadata' from the status and action columns.

timestamp	user id	status	action
01-01-2025:10.30	67945	success	user added a product x to their cart
01-01-2025:10.32	38910	fail	invalid values typed for product quantity
01-01-2025:10.38	17462	fail	customer table corrupted

Event timestamp

Person, system, or account associated with the event

Event & event metadata

Log

An append-only sequence of records ordered by time, capturing information about events that occur in systems.



timestamp	user id	status	action
01-01-2025:10.30	67945	success	user added a product x to their cart
01-01-2025:10.32	38910	fail	invalid values typed for product quantity
01-01-2025:10.38	17462	fail	customer table corrupted

```
{  
  "user id": 67945,  
  "action": "user added a product x to their cart",  
  "status": "success",  
  "time-stamp": 01-01-2025:10.30  
}
```

[00101011 11000101 11001001 11000101 110001001]

user id	action	status	timestamp
67945	user added a product x to their cart	success	01-01-2025:10.30
38910	invalid values typed for product quantity	fail	01-01-2025:10.32
17462	customer table corrupted	fail	01-01-2025:10.38

Log Levels

A tag to categorize the event (log level)

- “debug”
- “info”
- “warn”
- “error”
- “fatal”

user id	action	status	timestamp	level
67945	user added a product x to their cart	success	01-01-2025:10.30	Info
38910	invalid values typed for product quantity	fail	01-01-2025:10.32	error
17462	customer table corrupted	fail	01-01-2025:10.38	fatal



DeepLearning.AI

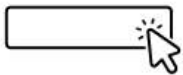
Introduction to Source Systems

Streaming Systems

Terminology

Event

Something that happened in the world or a change to the state of a system.



User clicking
on a link

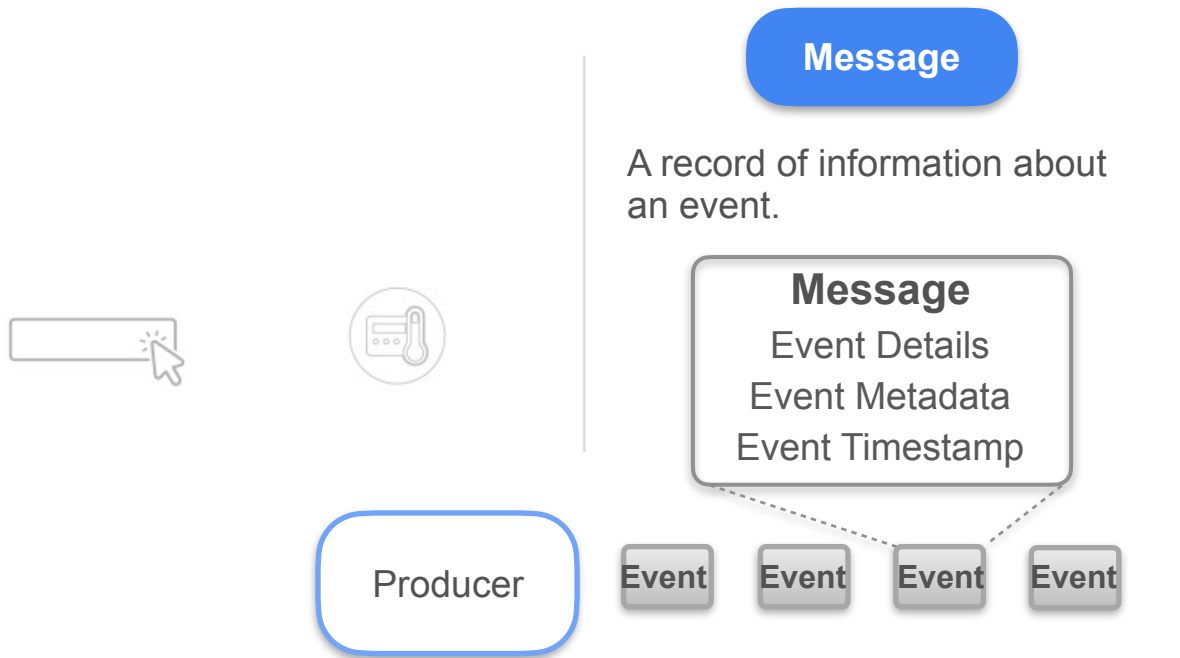


Sensor measuring a
temperature change

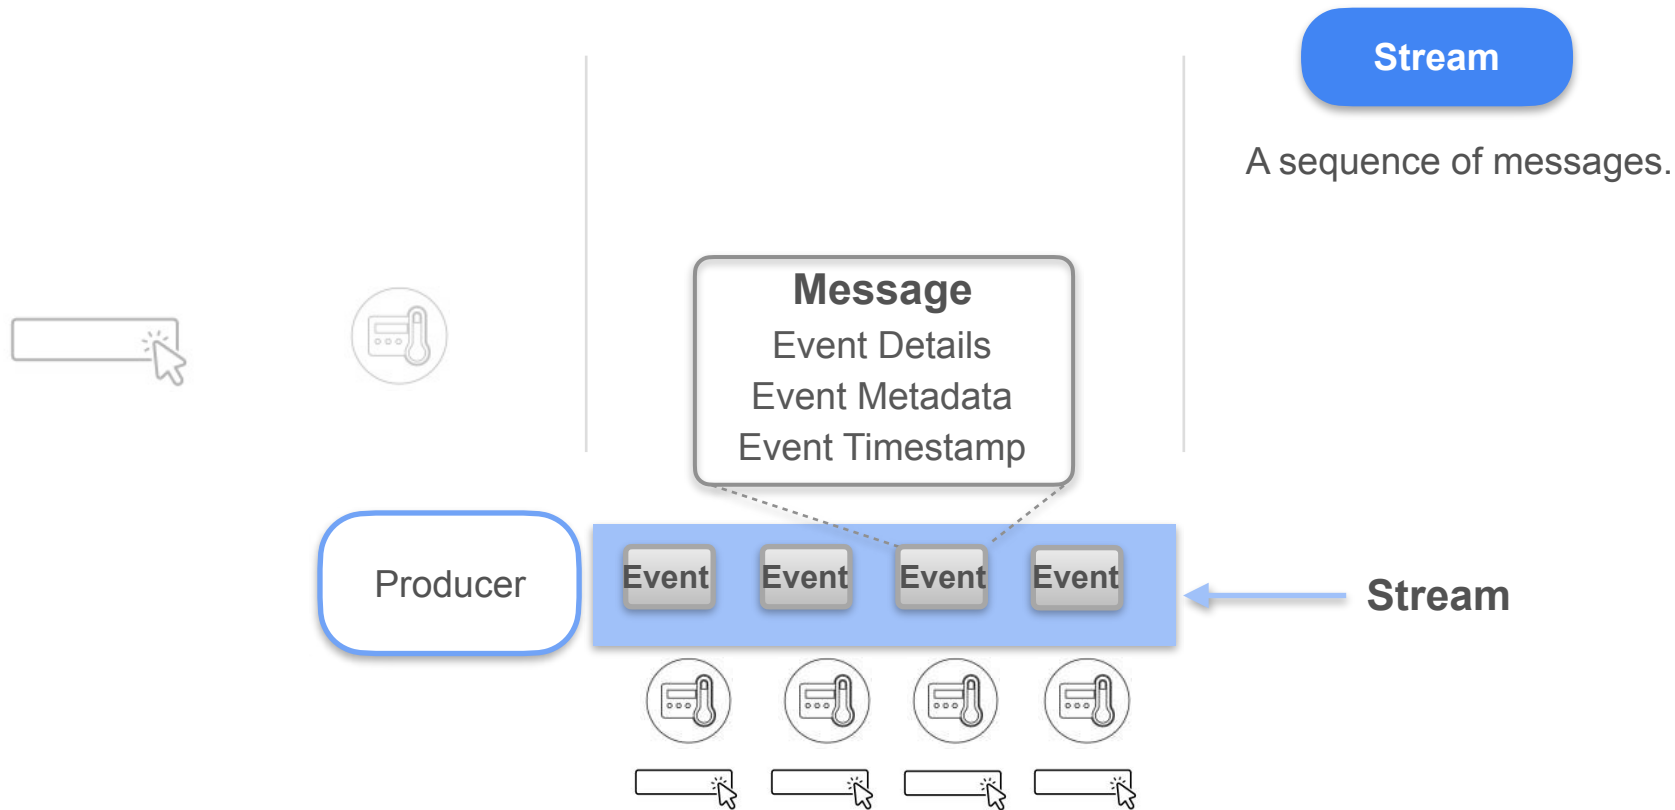
PEvent;er

Data: record of events

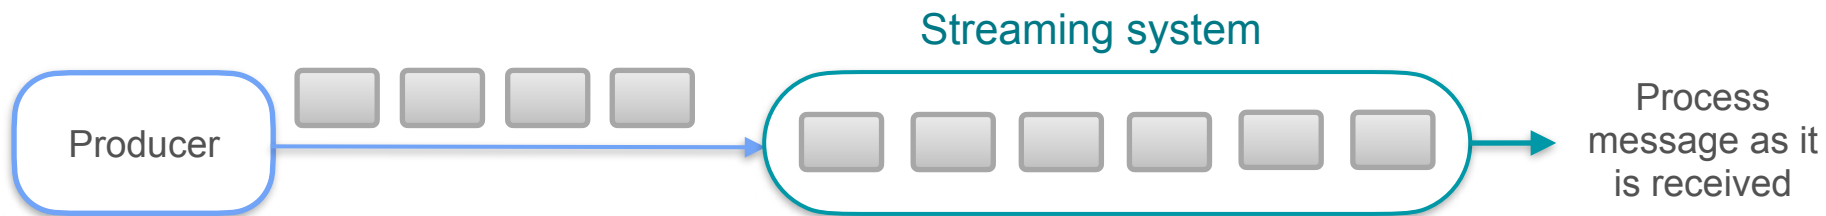
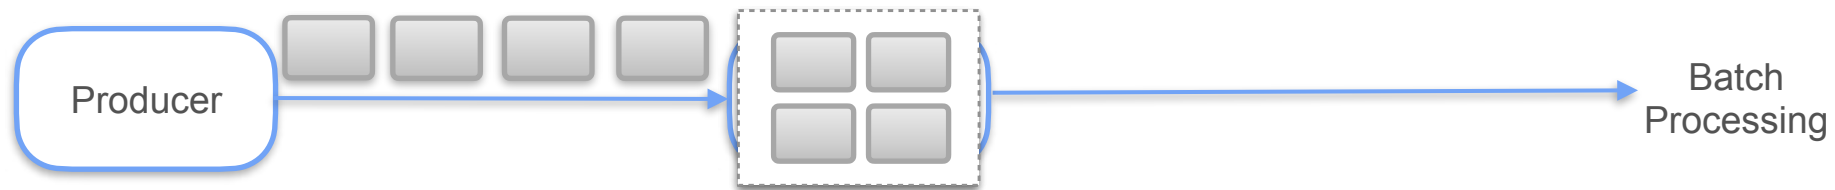
Terminology



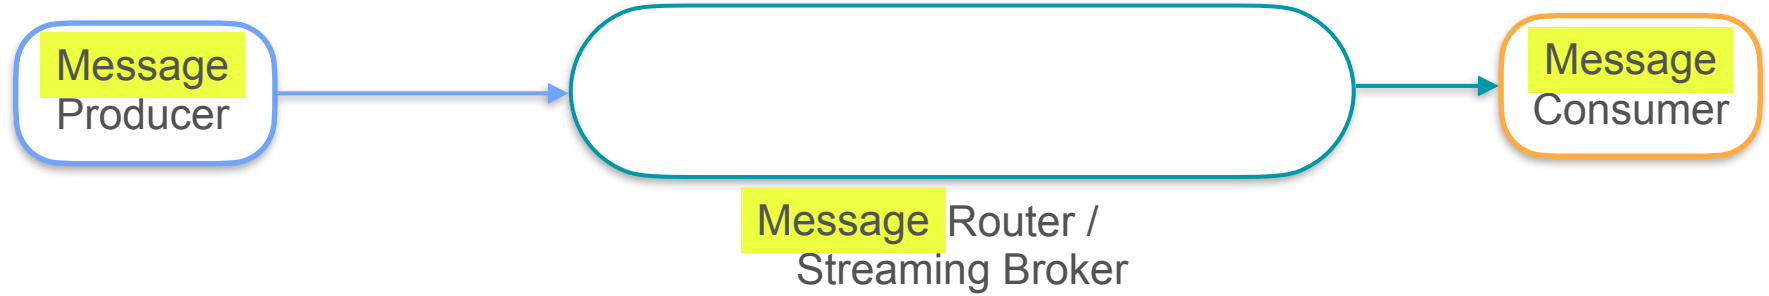
Terminology



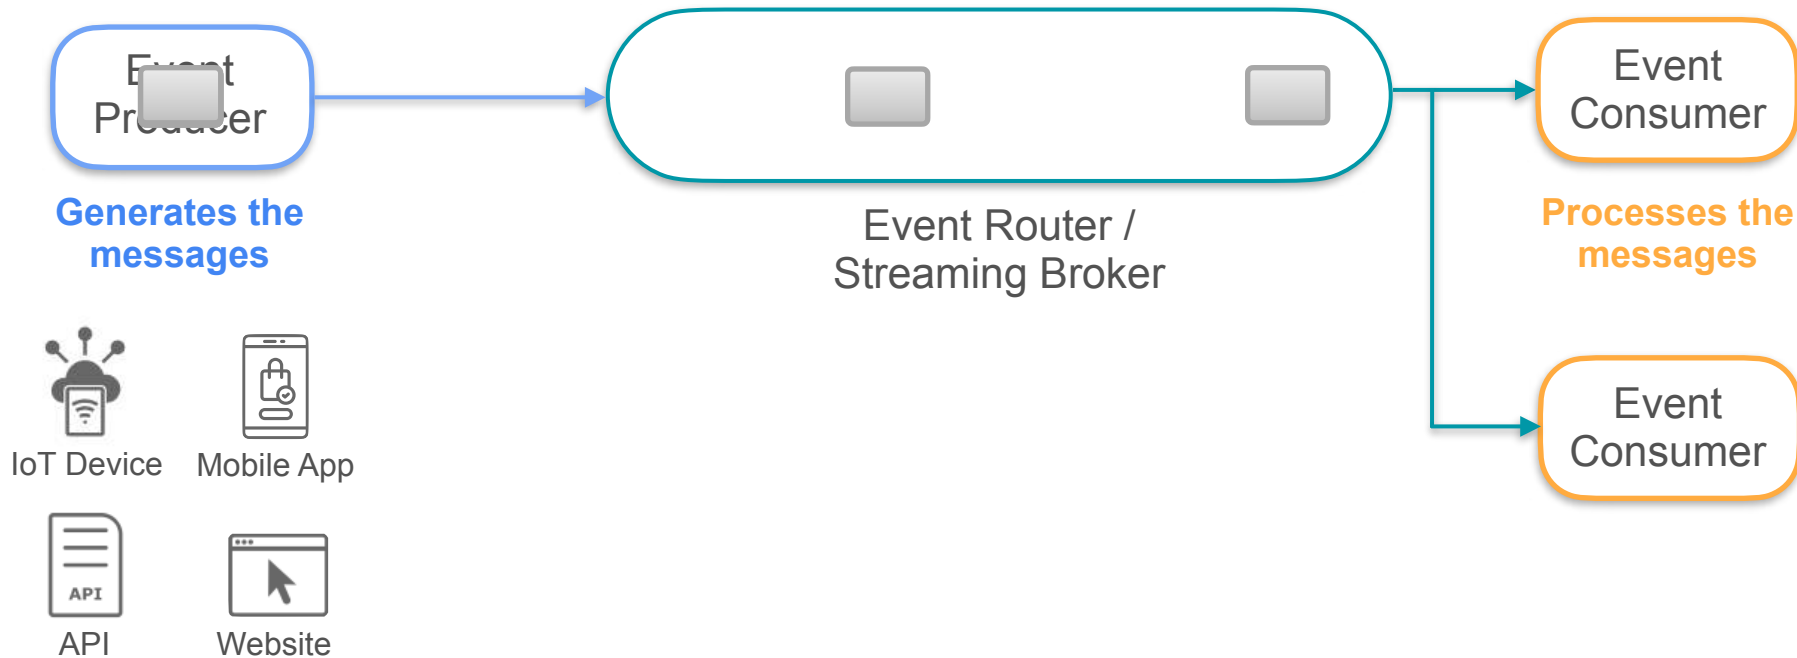
Stream Processing



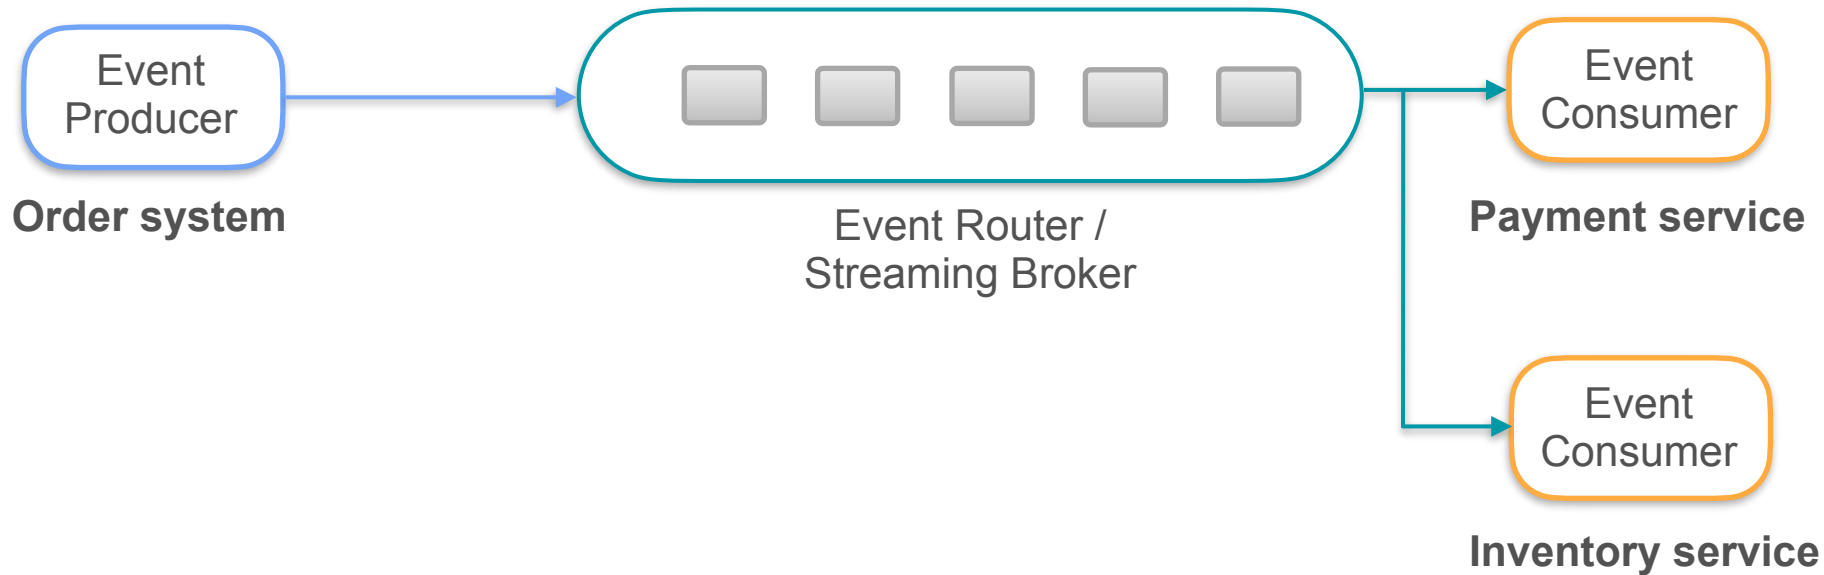
Streaming System



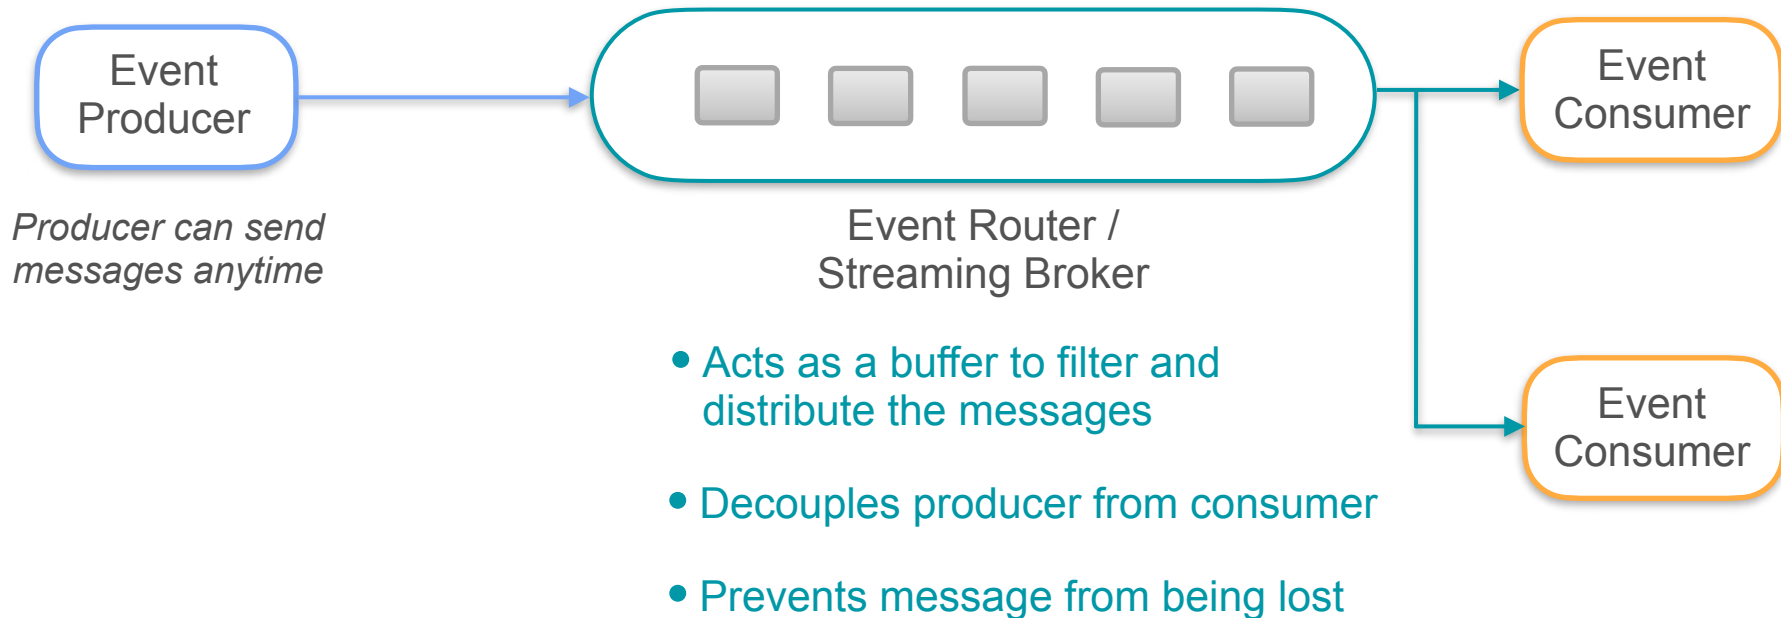
Streaming System



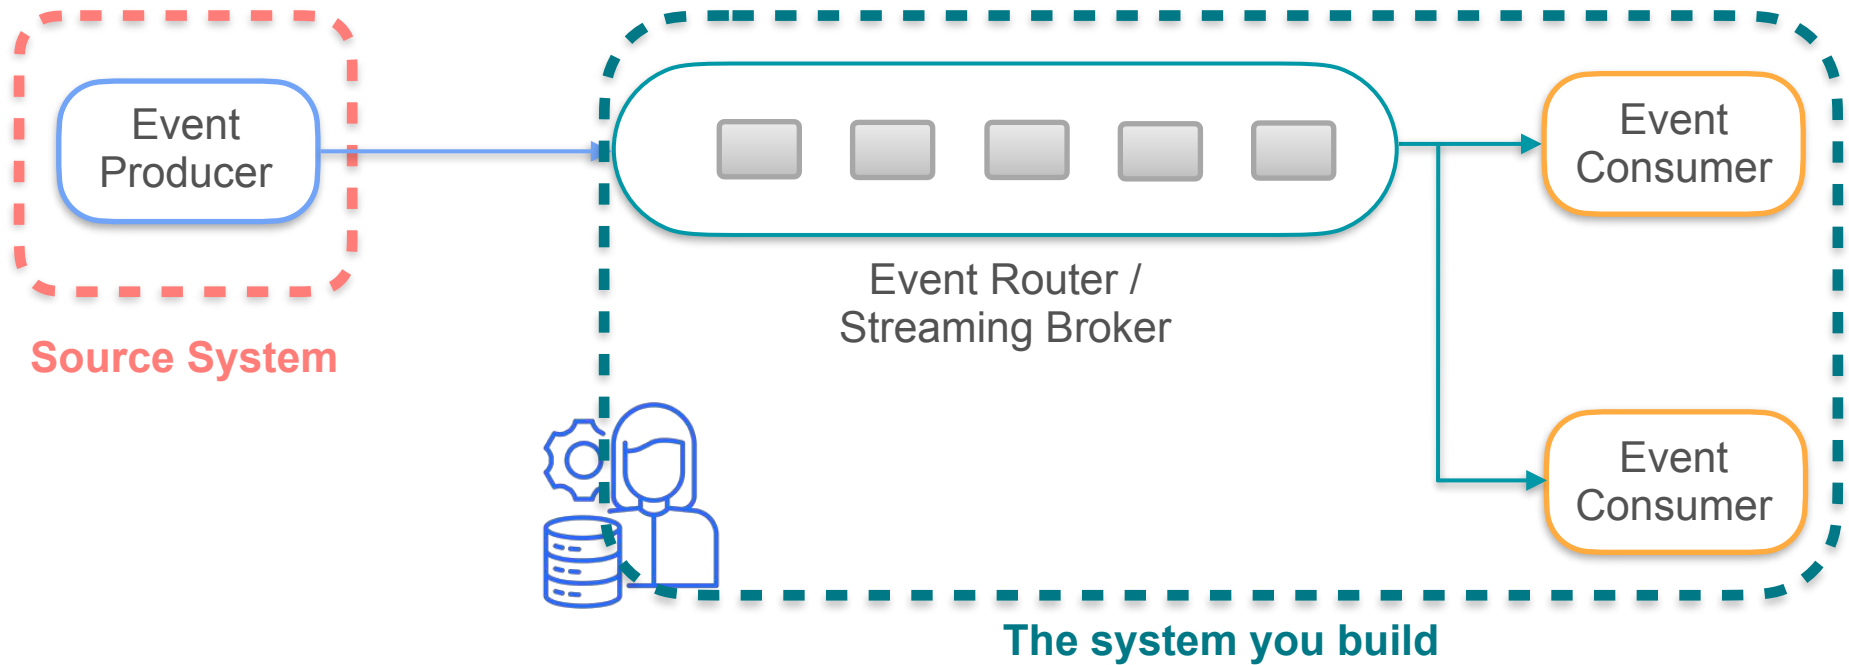
Streaming System



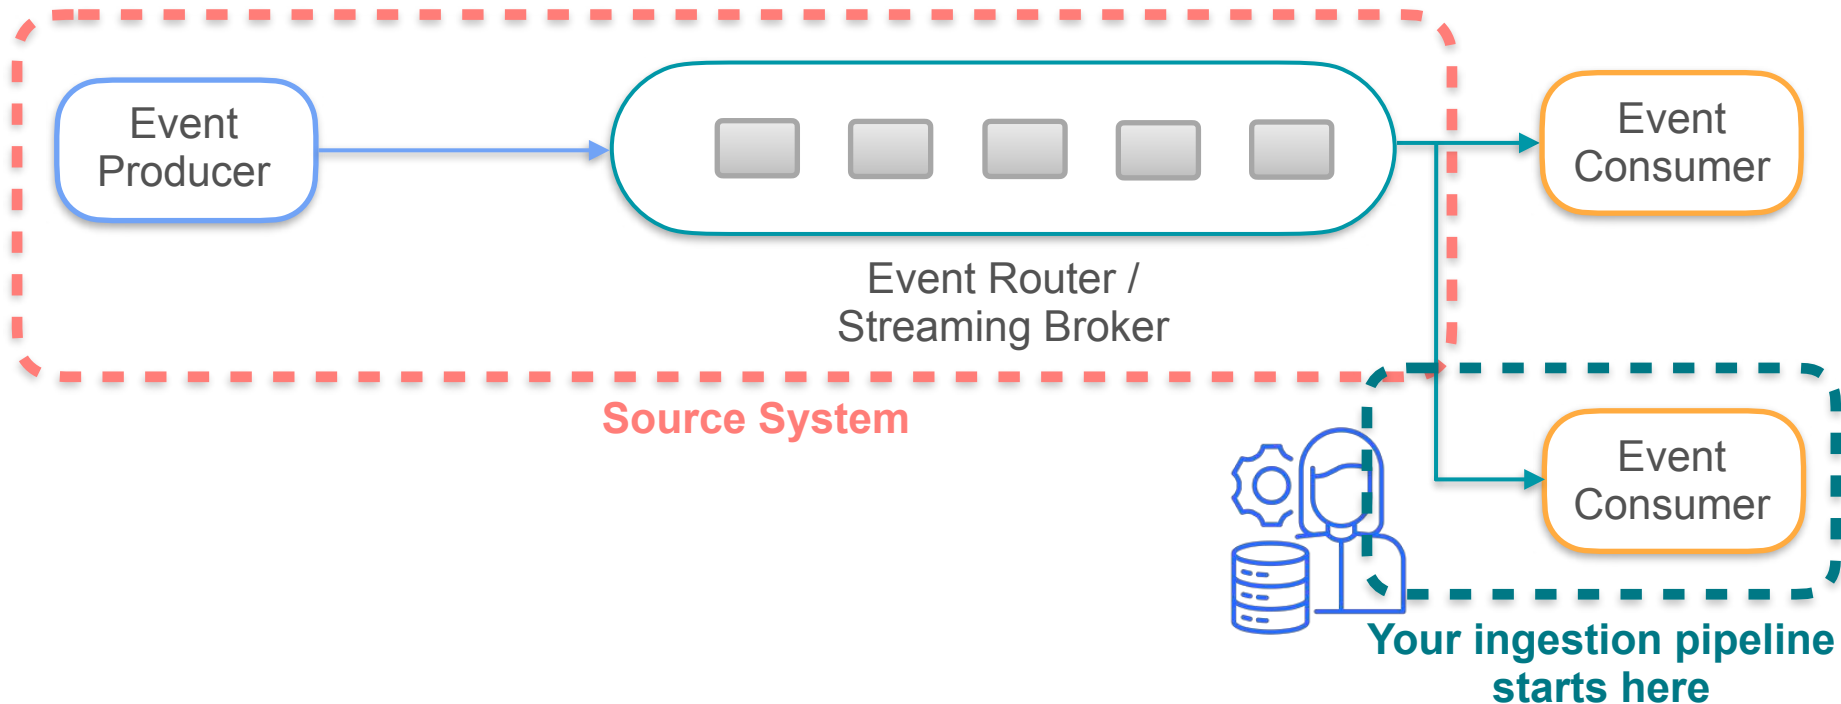
Streaming System



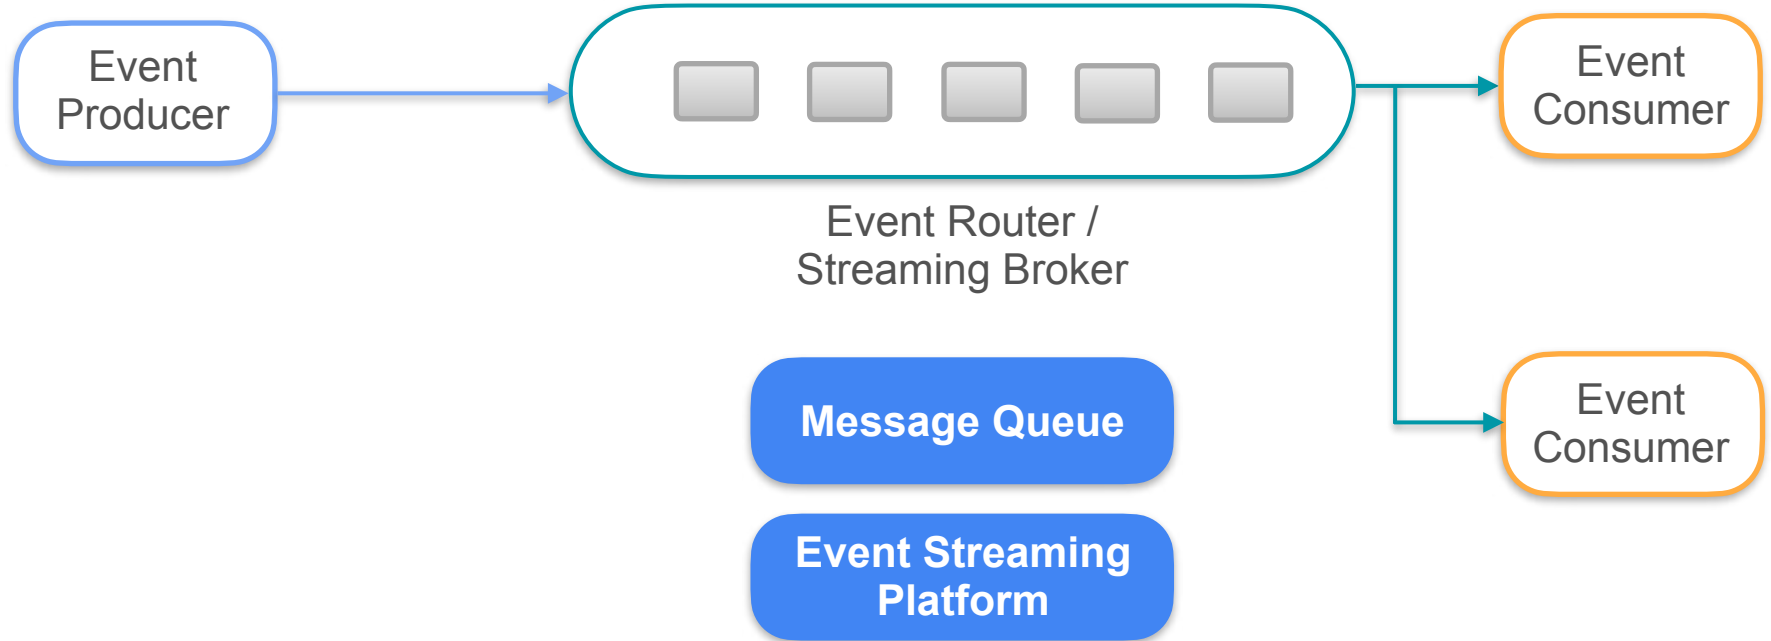
Your Data System



Your Data System

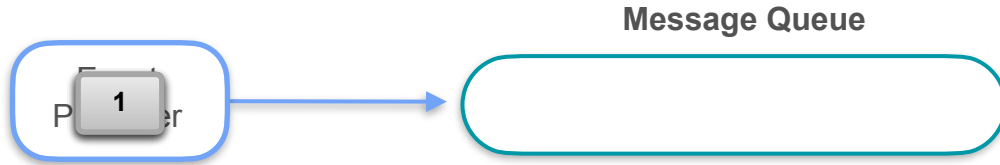


Streaming System



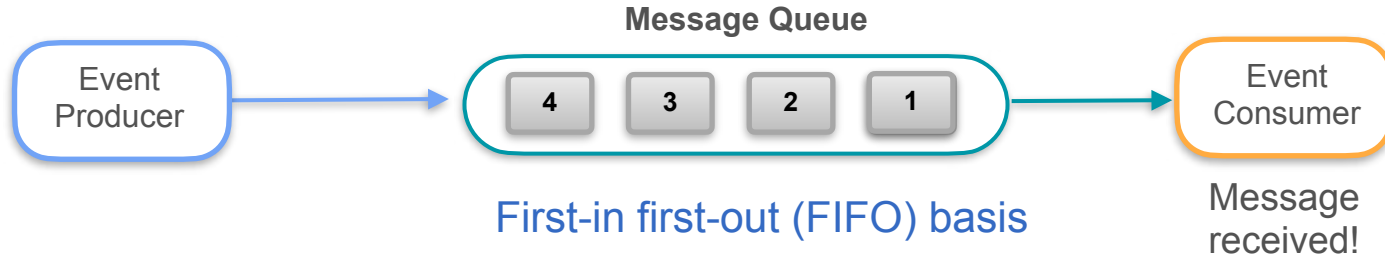
Message Queue

A queue/buffer that accumulates messages



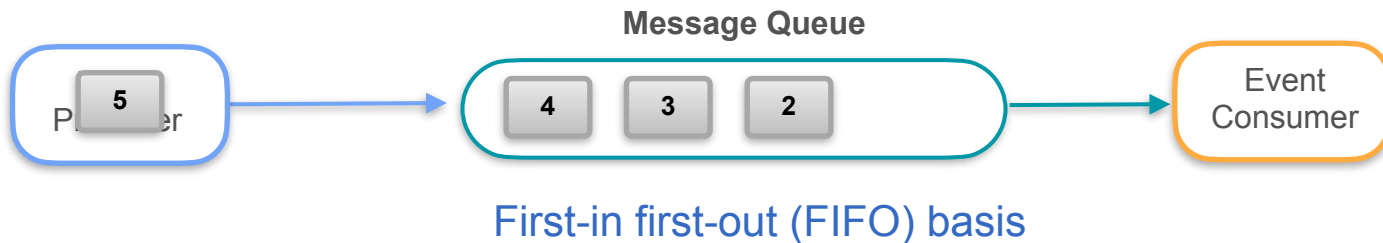
Message Queue

A queue/buffer that accumulates messages and delivers those messages to consumers asynchronously.



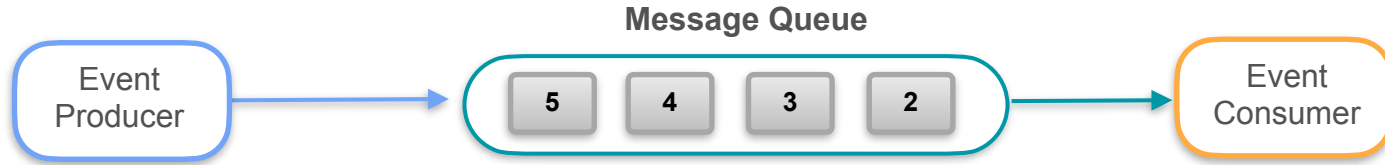
Message Queue

A queue/buffer that accumulates messages and delivers those messages to consumers asynchronously.



Message Queue

A queue/buffer that accumulates messages and delivers those messages to consumers asynchronously.



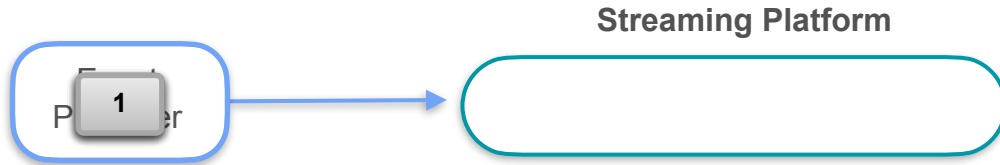
First-in first-out (FIFO) basis
Temporary storage



Amazon Simple Queue
Service (Amazon SQS)

Event Streaming Platform

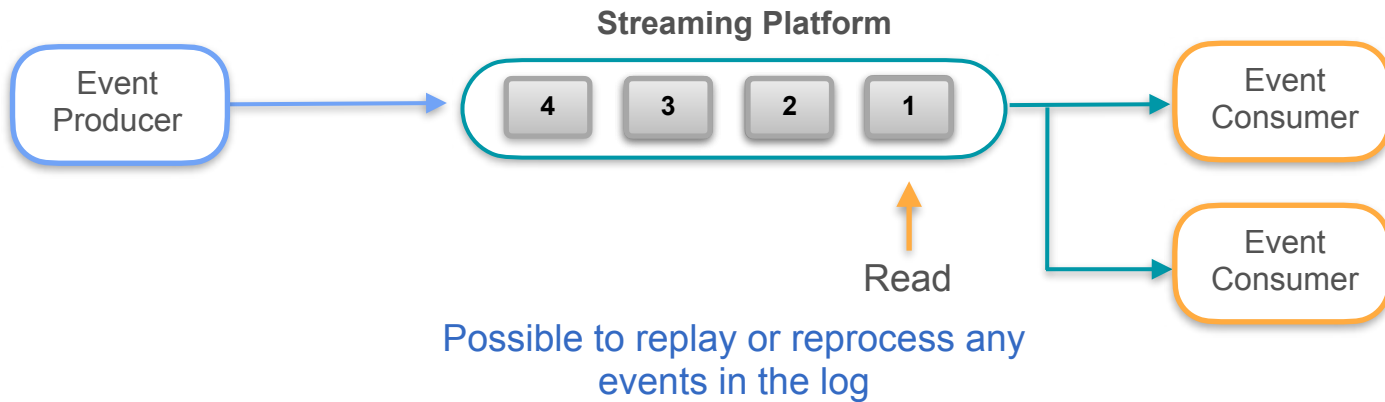
Log: Append-only record of events



Amazon Kinesis Data Streams

Event Streaming Platform

Log: Append-only record of events



Amazon Kinesis Data Streams

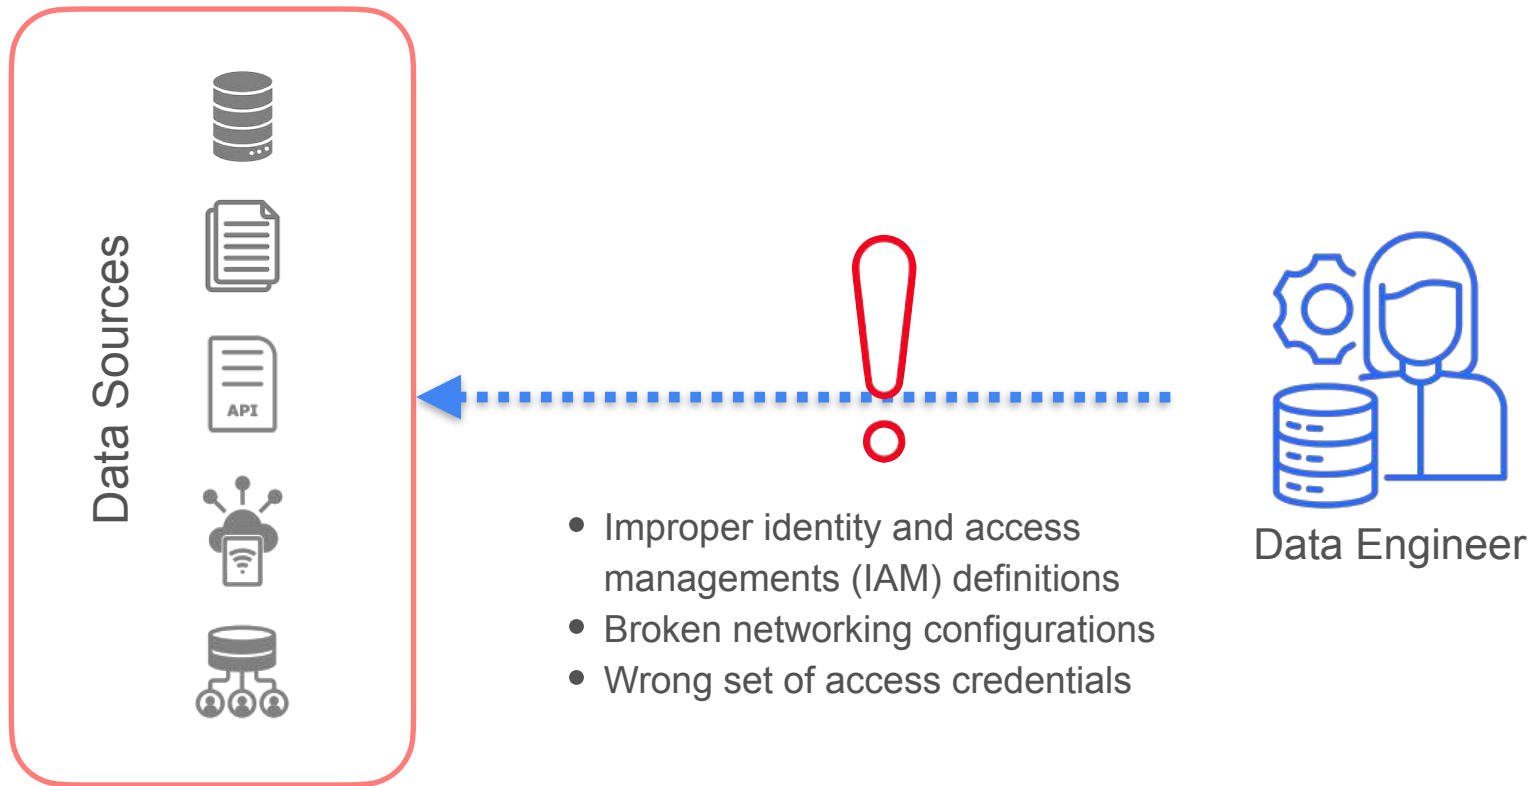


DeepLearning.AI

Interacting with Source Systems

Lesson Overview

Connecting to Source Systems



Lesson's Plan

1

Ways of connecting to source systems

2

IAM roles and permissions

*Key to controlling and managing access to
cloud-based data sources*



Role



Permissions

Lesson's Plan

1

Ways of connecting to source systems

2

IAM roles and permissions

*Key to controlling and managing access to
cloud-based data sources*

3

Basics of networking

VPCs and Subnets, Gateways, Routing, Security groups



Role



Permissions

Lesson's Plan

1

Ways of connecting to source systems

2

IAM roles and permissions

Key to controlling and managing access to cloud-based data sources

3

Basics of networking

VPCs and Subnets, Gateways, Routing, Security groups

4

Lab exercise: put your skills to the test

Your job: troubleshoot and figure out the cause of the problem



Role



Permissions

Real world scenario





DeepLearning.AI

Interacting with Source Systems

Connecting to Source Systems

Connecting to Source Systems

```
def create_client():  
    dynamodb_client = boto3.client("dynamodb")  
  
    return dynamodb_client
```

boto3: AWS Software Development Kit (SDK) for Python

Connecting to Source Systems

Running this command in Cloud9 IDE

```
mysql --host=<MySQLEndpoint> --user=<DatabaseUserName> --password=<Password> --port=3306
```

Programmatic Way

SDK (boto3)



IDE (Cloud 9)

```
connect_to_source.py x (+)
1 import pymysql
2 import boto3
3
4 ENDPOINT="....."
5 PORT="3306"
6 USER="jane_doe"
7 REGION="us-east-1"
8 DBNAME="mydb"
9
10 #gets the credentials from .aws/credentials
11 session = boto3.Session(profile_name='default')
12 client = session.client('rds')
```



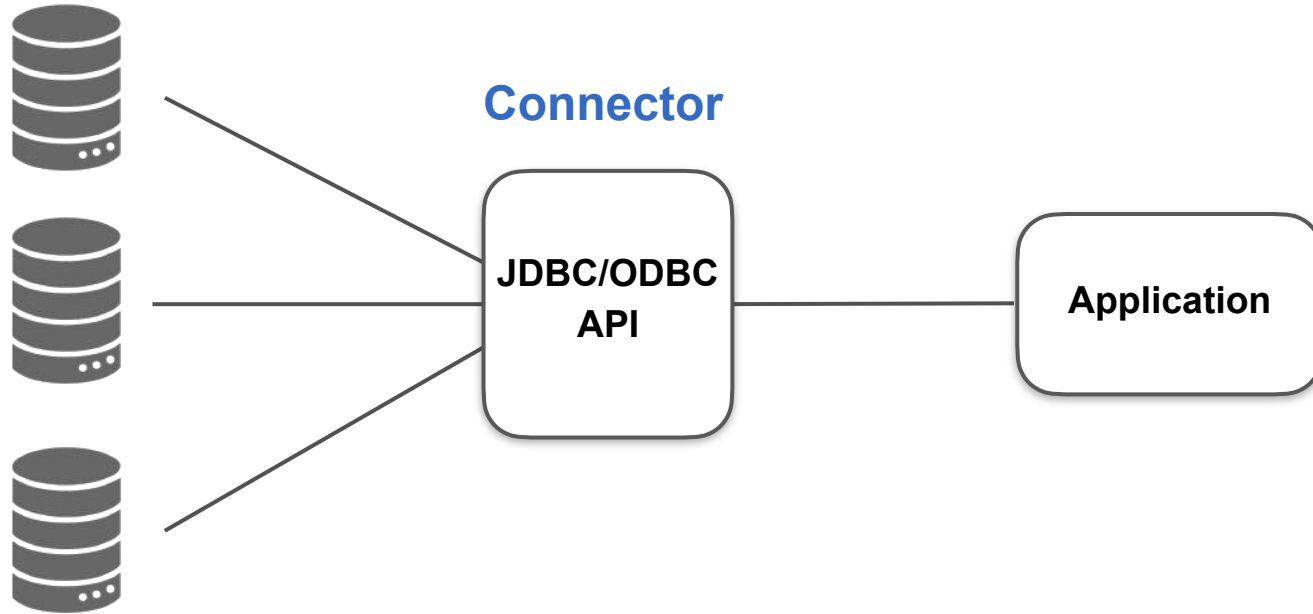
Jupyter Notebook

```
import pymysql
import boto3

ENDPOINT="....."
PORT="3306"
USER="jane_doe"
REGION="us-east-1"
DBNAME="mydb"

#gets the credentials from .aws/credentials
session = boto3.Session(profile_name='default')
client = session.client('rds')
```

API Connector





DeepLearning.AI

Interacting with Source Systems

Basics of IAM and Permissions

Security on the Cloud



Encryption Methods

Identity and Access
Management (IAM)

Networking Protocols



We're only human:

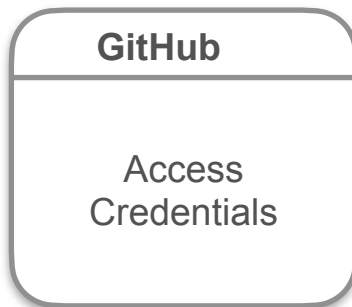
The #1 root
cause of cloud
data breaches
is human error

- Insecure storage of passwords
- IAM misconfigurations

Mistakes



Public S3 Bucket



Admin access

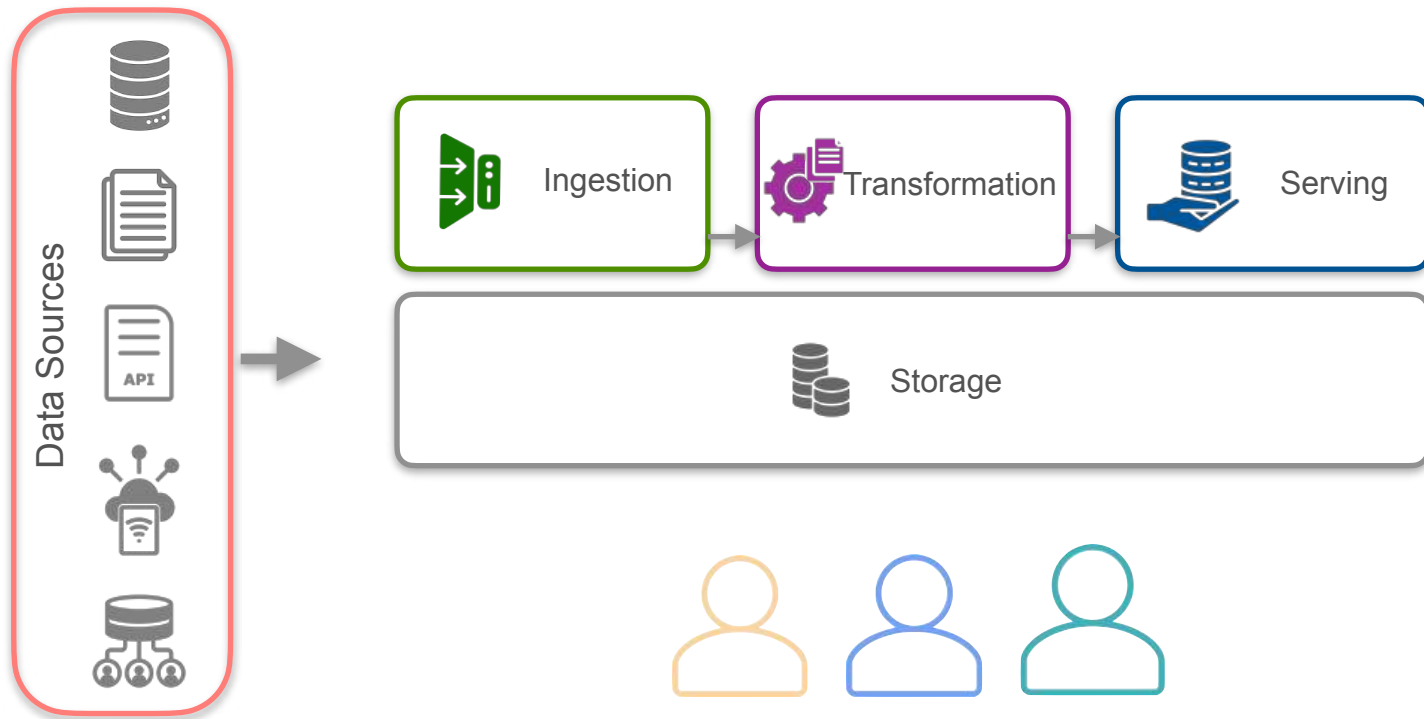
IAM

IAM is a framework for managing permissions.

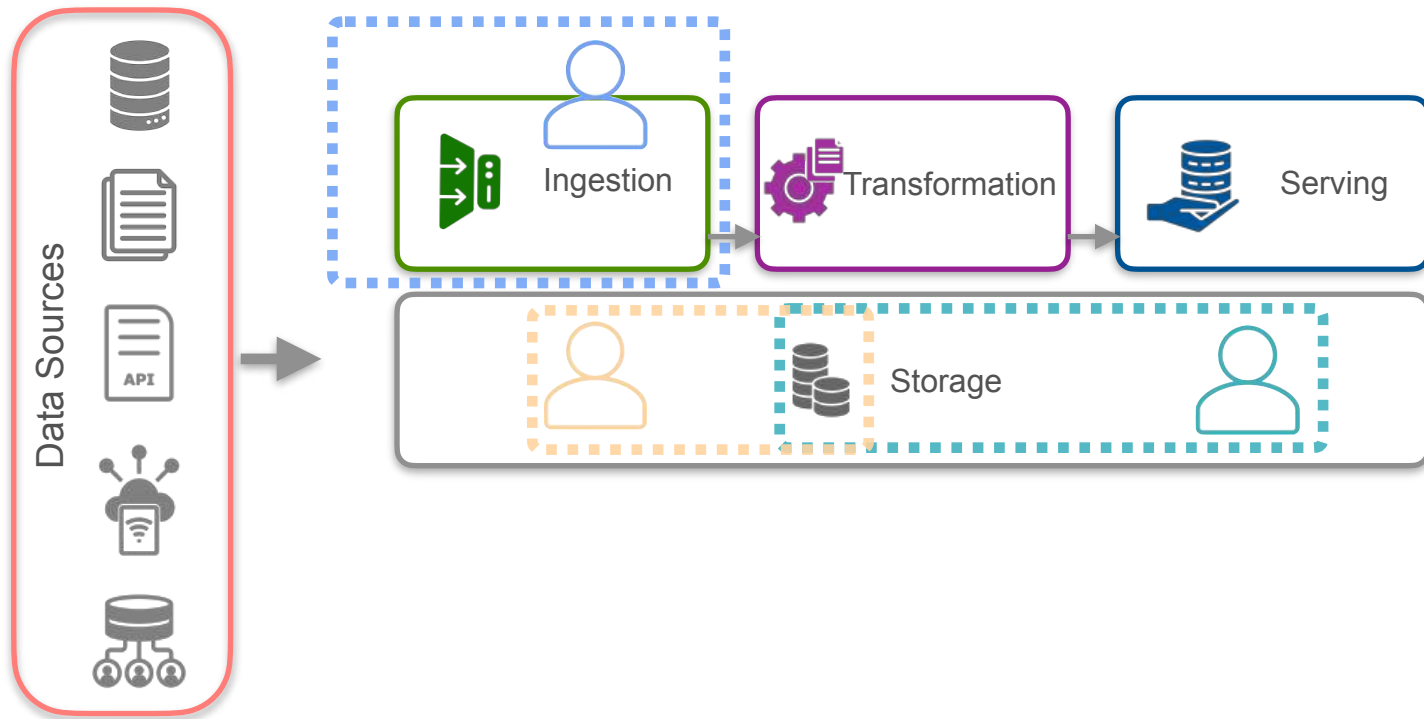
Permissions define which actions an identity
can perform on a specific set of resources



Principle of Least Privilege



Principle of Least Privilege



Principle of Least Privilege



AWS IAM



AWS Identity and Access
Management (IAM)

AWS IAM

Root User

Has unrestricted access to all resources

IAM User

Has specific permissions to certain resources

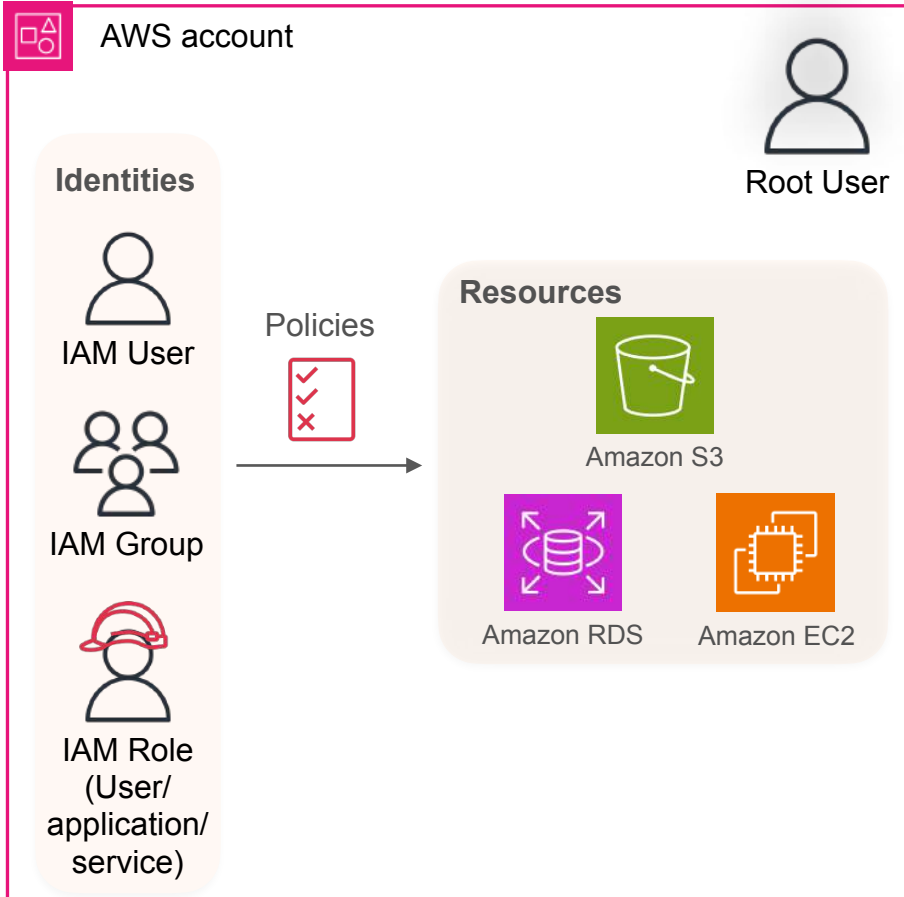
- Username & password
- Access key

IAM Group

A collection of users that inherit the same permission from the group policy

IAM Role

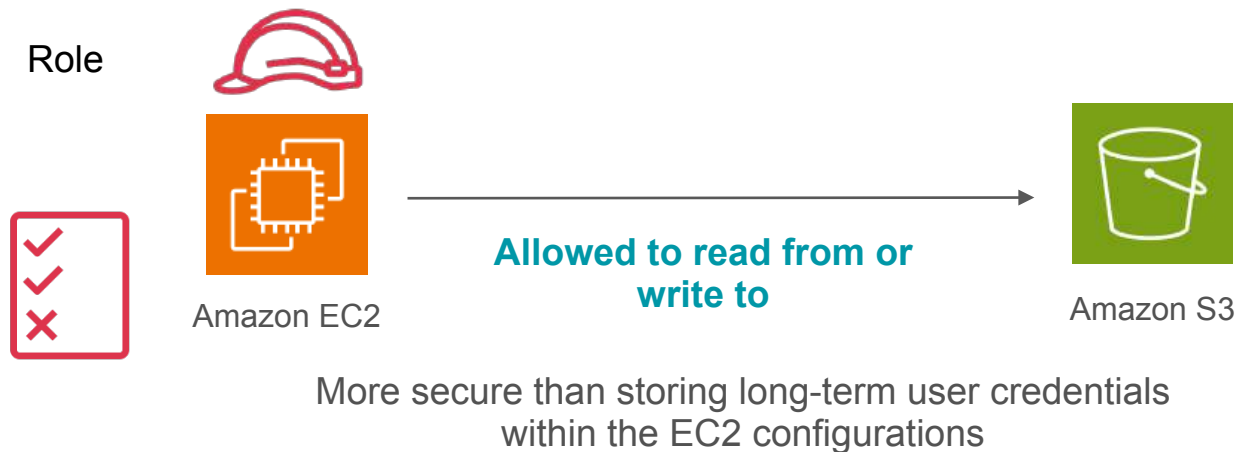
A user, application, or service that's been granted temporary permissions



AWS IAM



AWS IAM



Check if credentials have expired!

IAM Policies

permission to access the
specified S3 buckets

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "S3AccessDLAIBucket",
      "Action": [
        "s3:List*",
        "s3:Get*"
      ],
      "Effect": "Allow",
      "Resource": [
        "arn:aws:s3:::dlai-data-engineering",
        "arn:aws:s3:::dlai-data-engineering/*"
      ]
    },
    {
      "Sid": "GlueMgmt",
      "Action": [
        "glue:*"
      ],
      "Effect": "Allow",
      "Resource": [
        "arn:aws:glue:*:*:catalog",
        "arn:aws:glue:*:*:/de-c1w2*"
      ]
    }
  ]
}
```

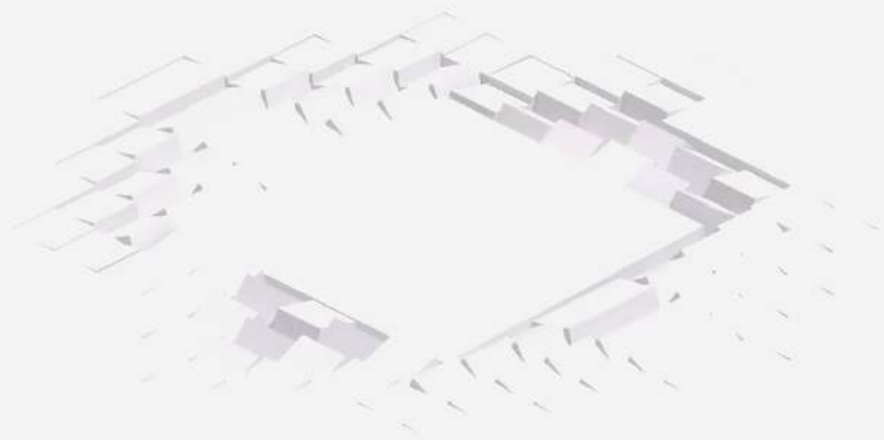
permission to access
the AWS Glue job



DeepLearning.AI

Interacting with Source Systems

Basics of Networking



AWS Cloud

What does cloud in “cloud computing” mean?

The “cloud” is made up of very real physical data centers that are spread out around the world.

Each dot
represents
a region



Screenshot from [AWS Global Infrastructure](#) (2023)

AWS Cloud



Resources are replicated across availability zones to ensure that your systems keep working even if a data center goes down.

AWS Cloud

Region considerations:

- Legal compliance
- Latency: *the closer your end users are to the region, the lower the latency*
- Availability: *the more availability zones, the better you will be able to recover from a disaster*
- Cost

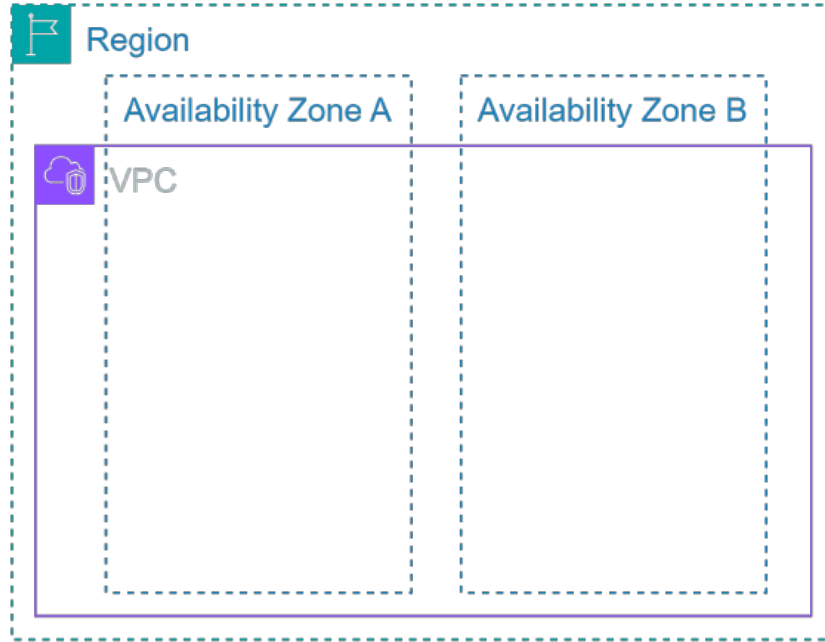
AWS Cloud

Region considerations:

- Legal compliance
- Latency
- Availability
- Cost



Virtual Private Cloud



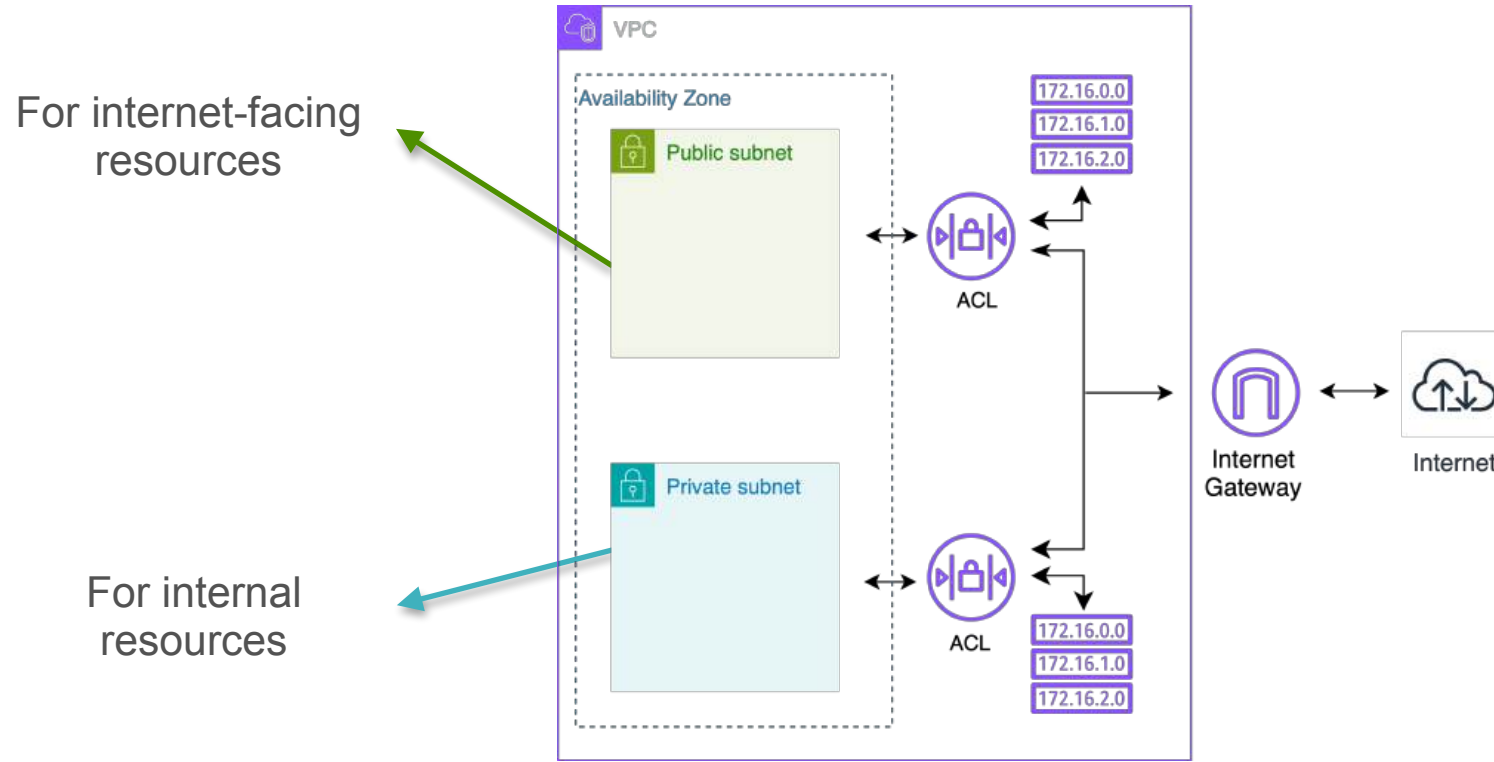
Virtual Private Cloud (VPC)

Smaller networks that span multiple availability zones

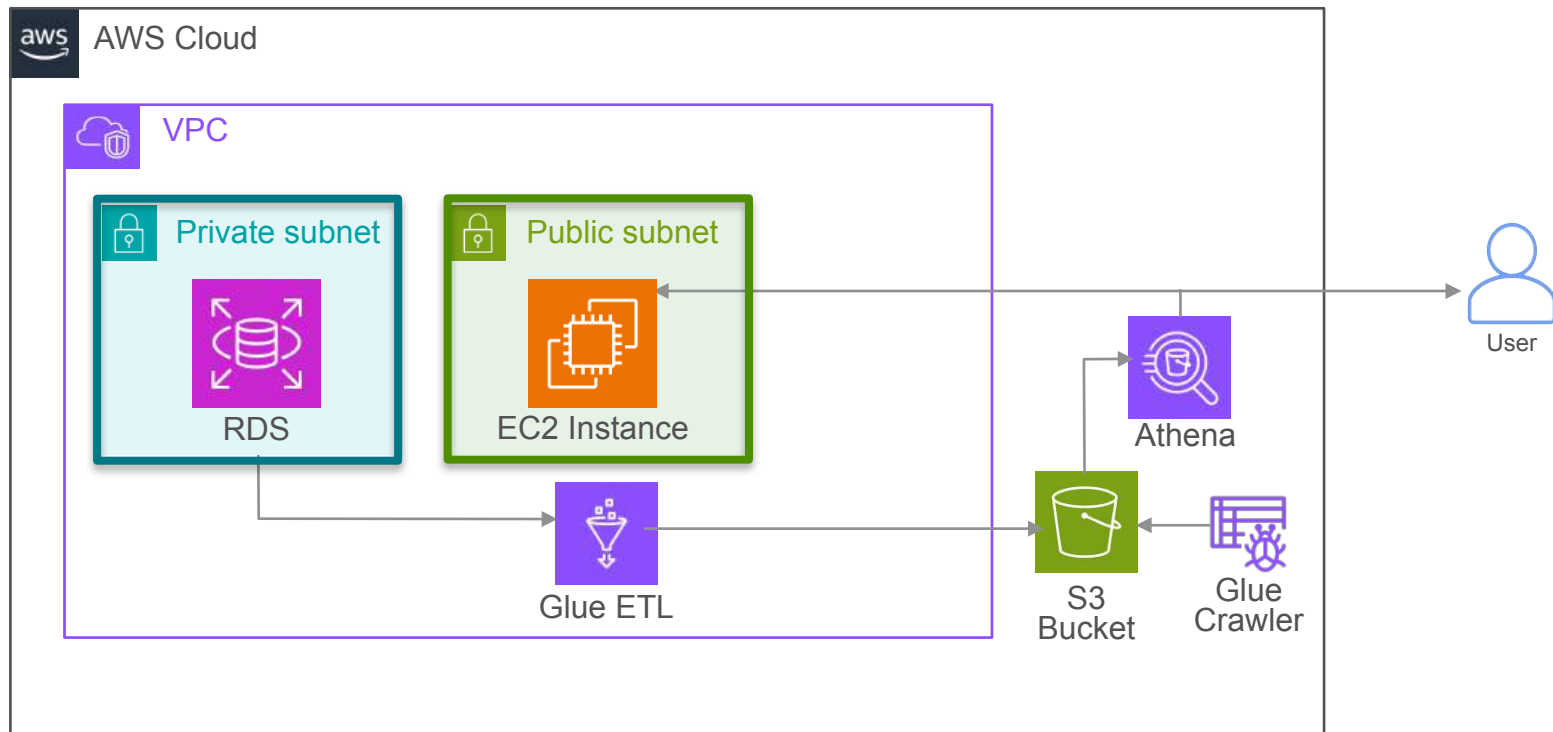
Virtual Private Cloud



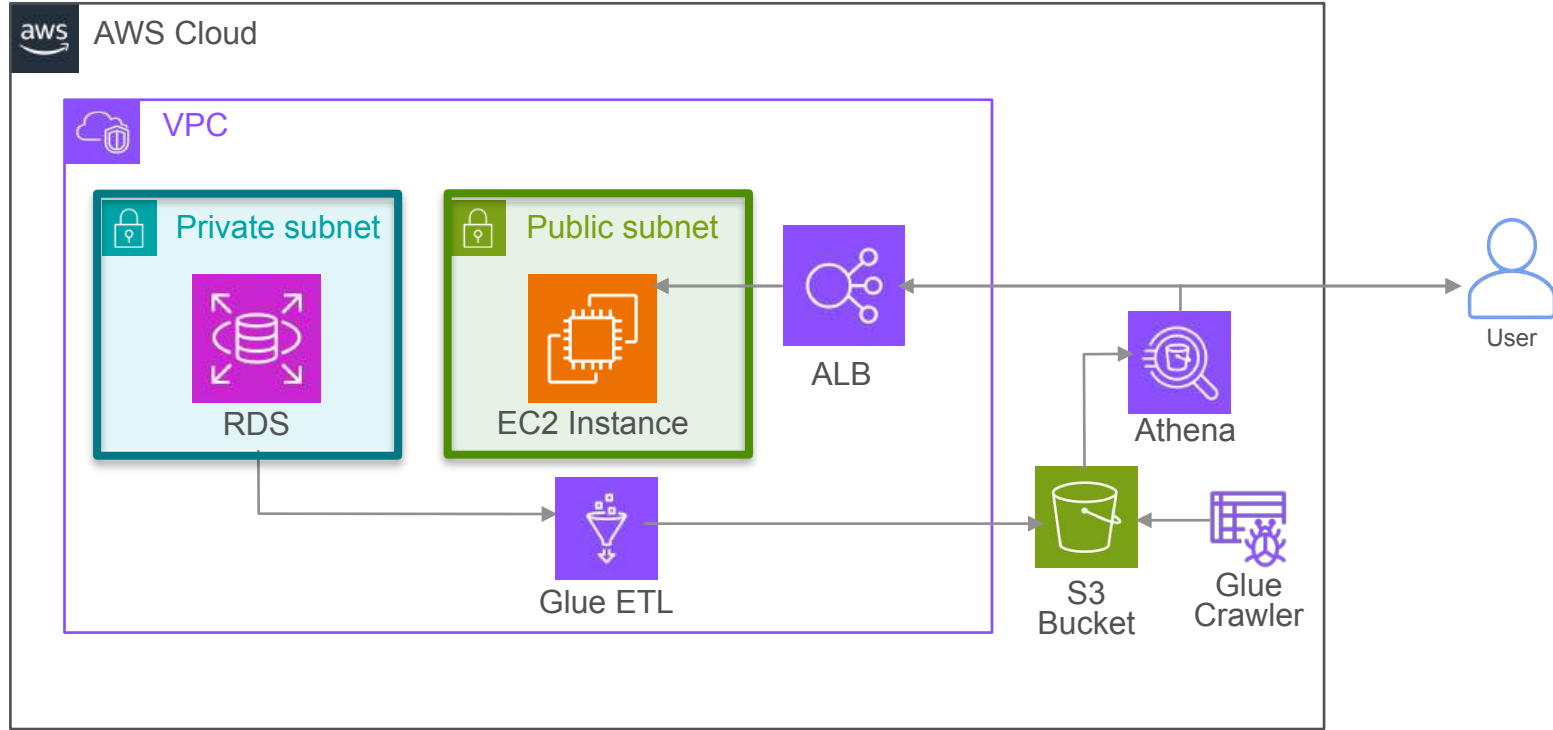
Virtual Private Cloud



Virtual Private Cloud



Virtual Private Cloud



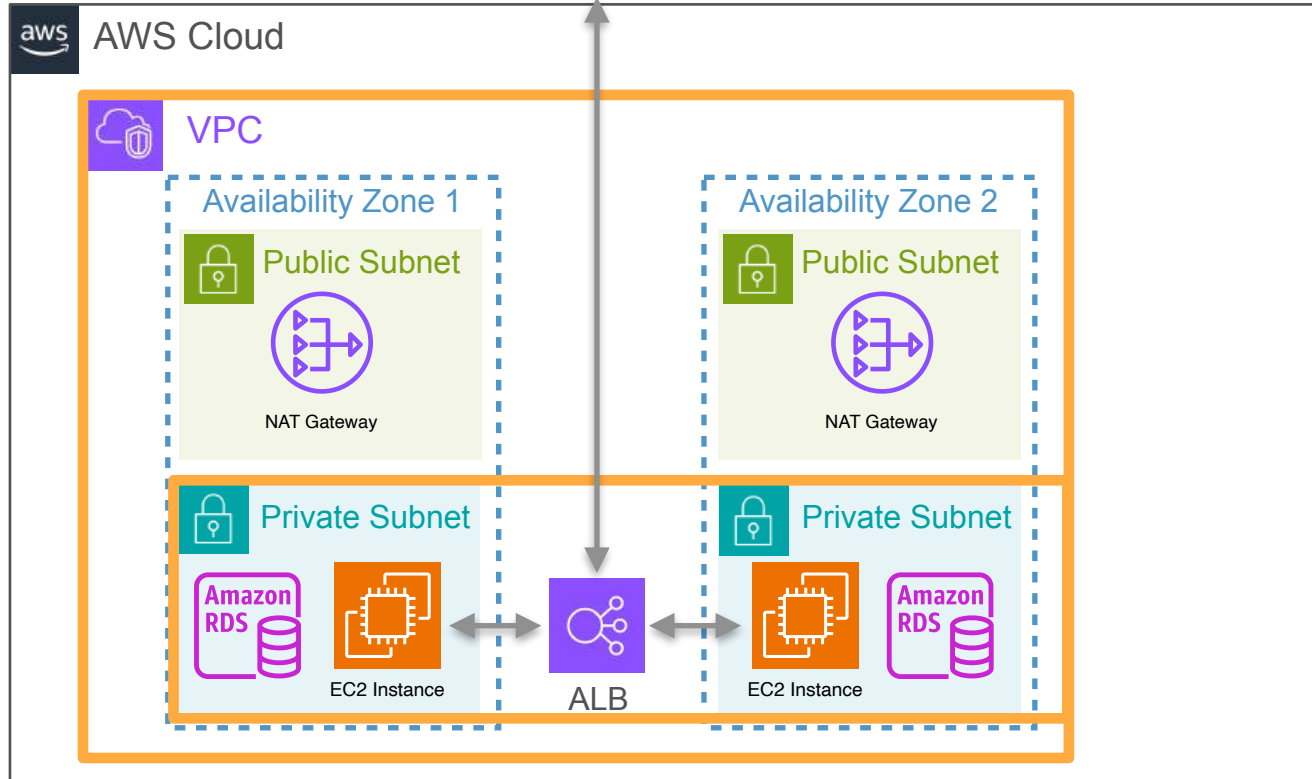


DeepLearning.AI

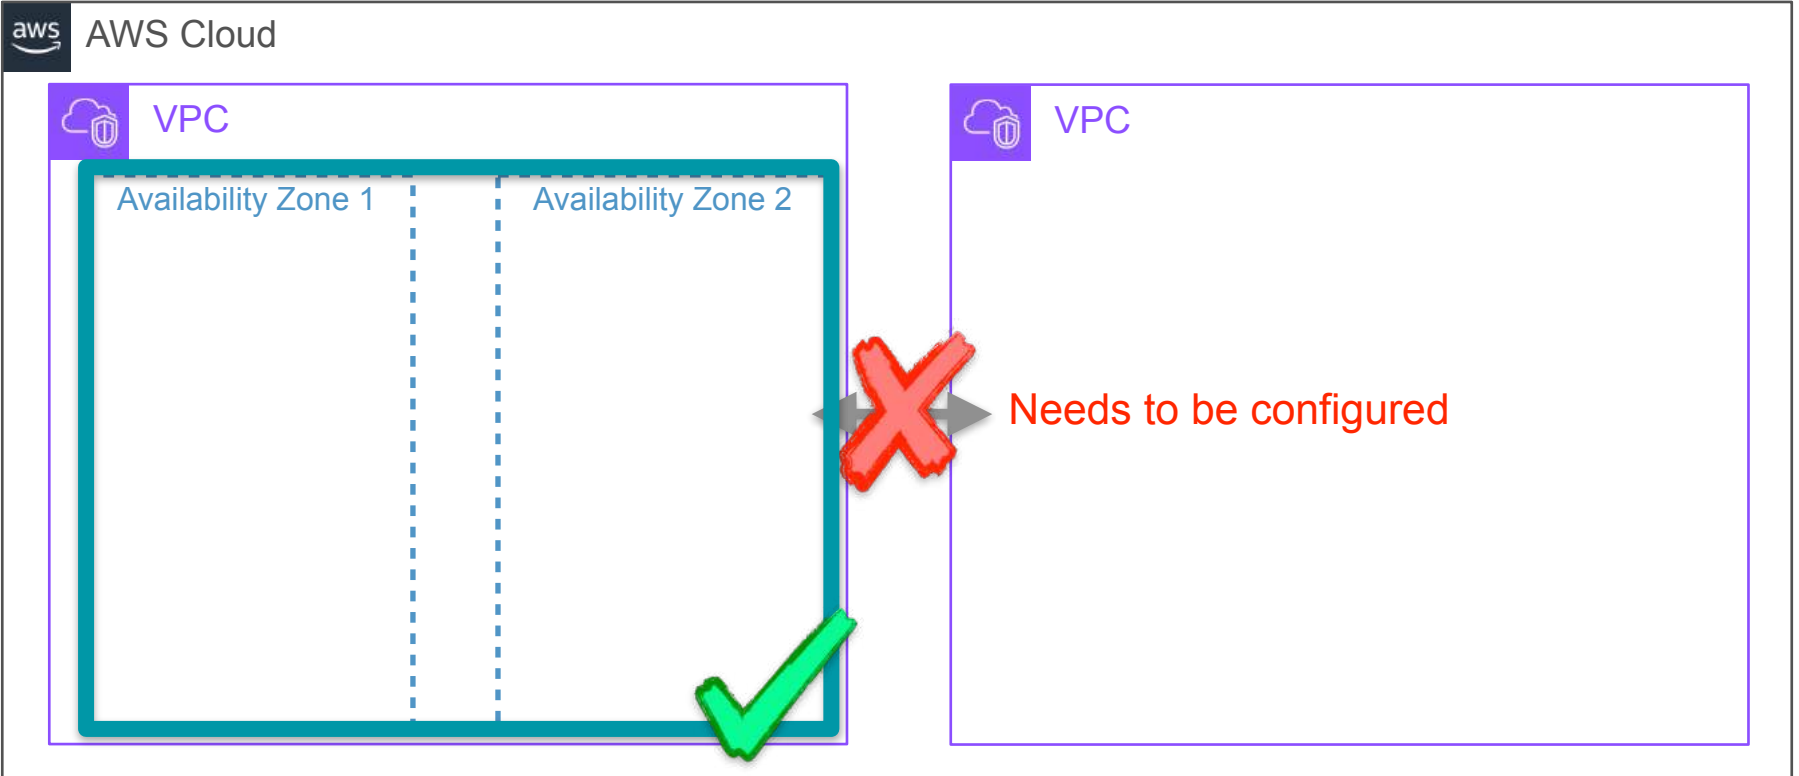
Interacting with Source Systems

AWS Networking - VPCs & Subnets

Example Scenario



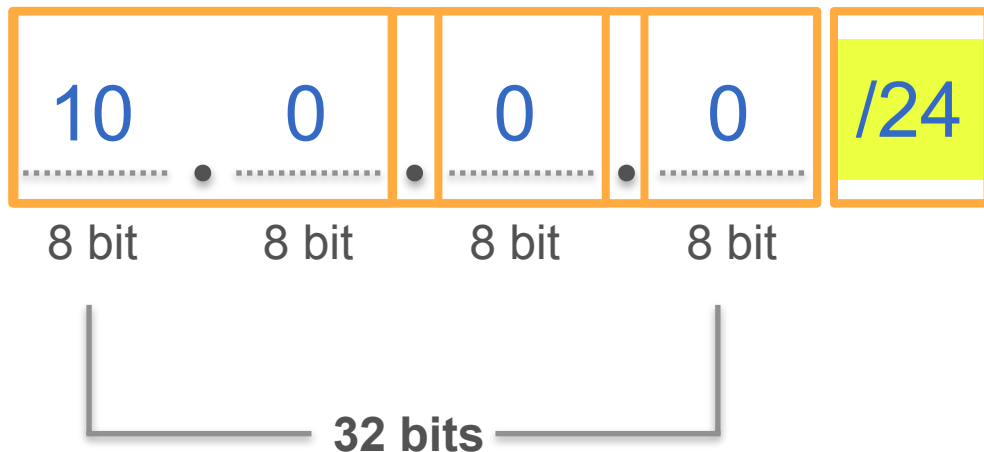
AWS Networking - VPCs & Subnets



AWS Networking - VPCs & Subnets

Define the network

Host addresses



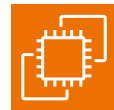
16 bits

prefix length

How many bits used for the network part of the address



10.0.__.

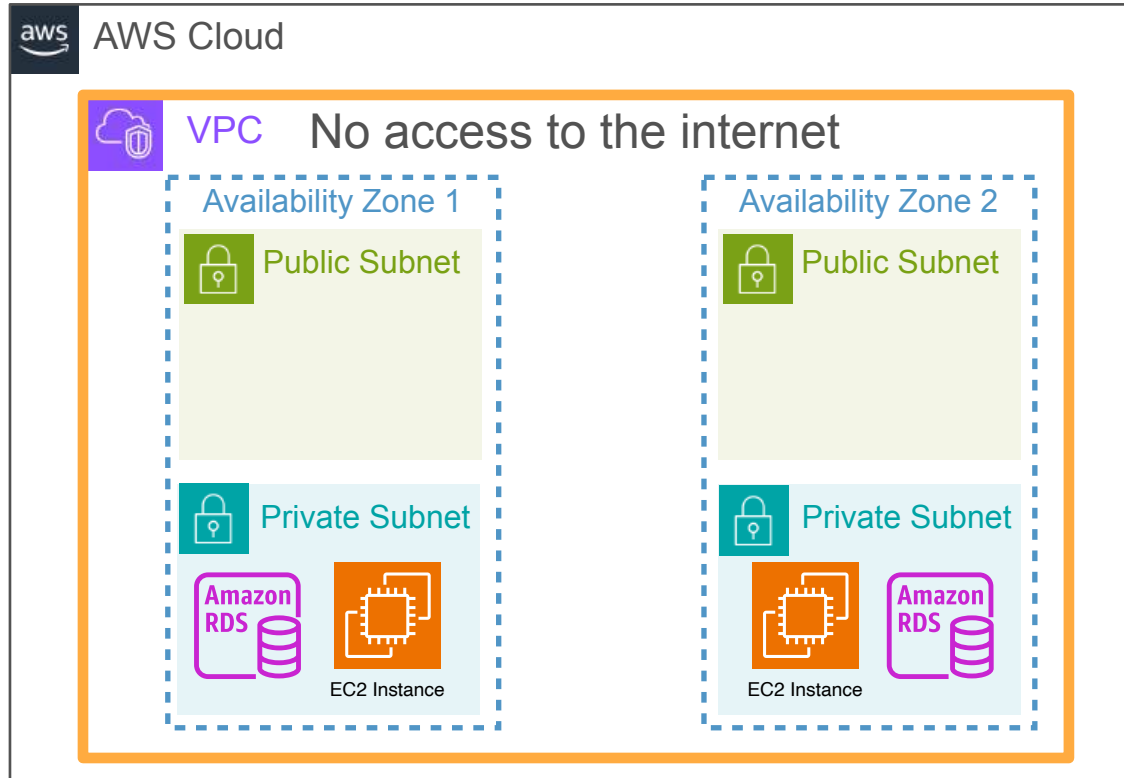


EC2 Instance

10.0.__.

0 to 255

AWS Networking - VPCs & Subnets



Closed network



DeepLearning.AI

Interacting with Source Systems

AWS Networking - Internet Gateways & NAT Gateways



AWS Cloud



VPC

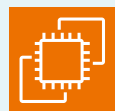
Availability Zone 1



Public Subnet



Private Subnet



EC2 Instance

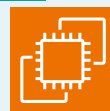
Availability Zone 2



Public Subnet

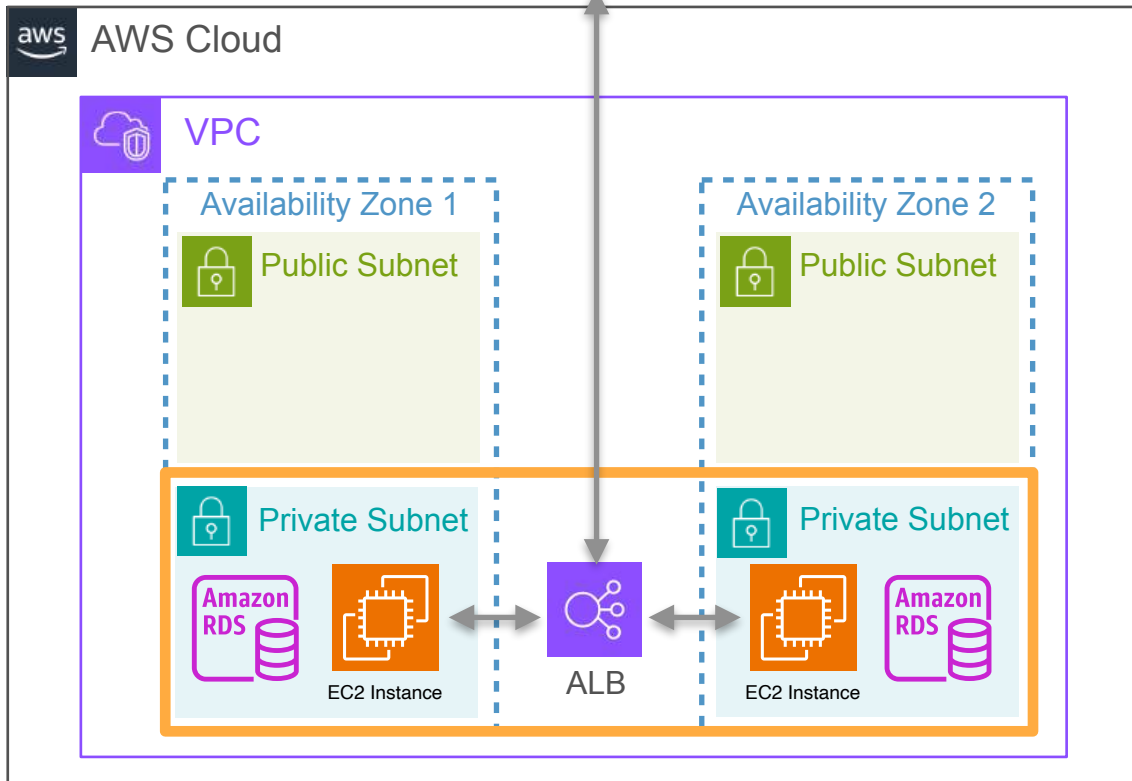


Private Subnet



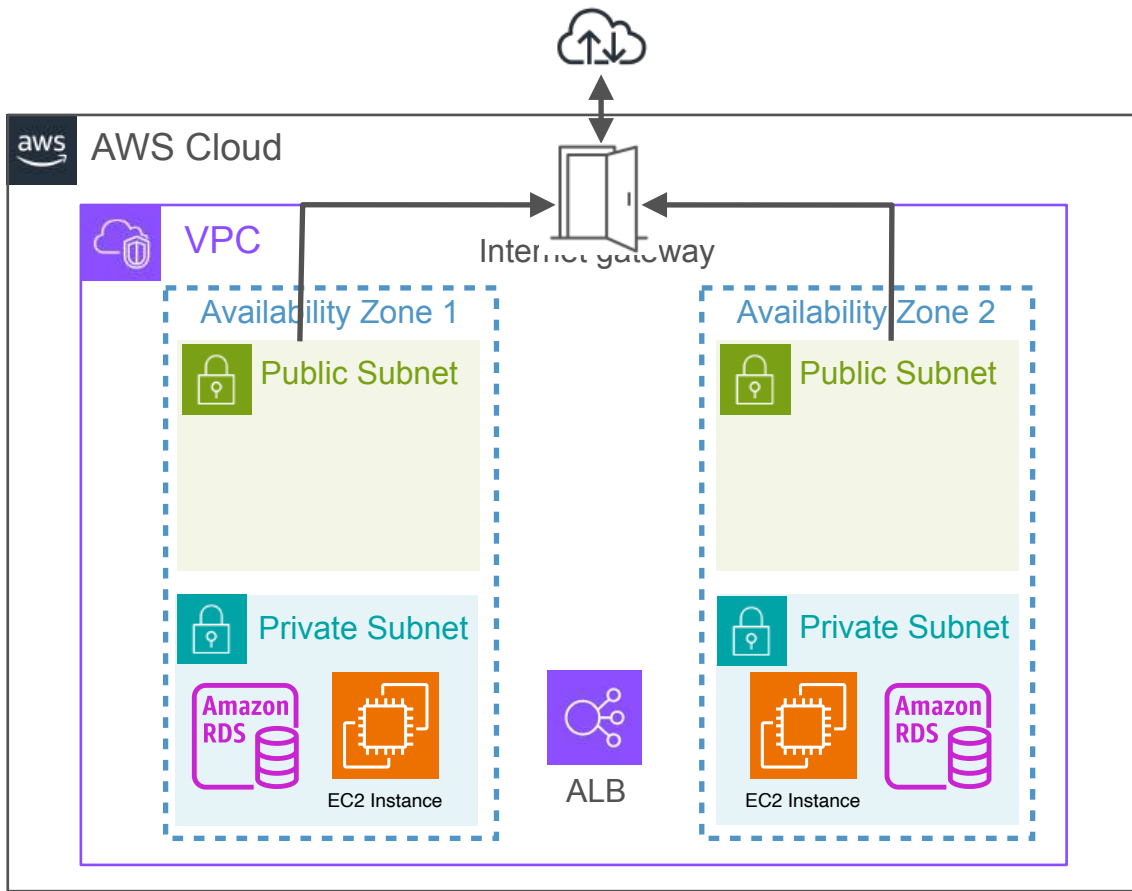
EC2 Instance

Example Scenario



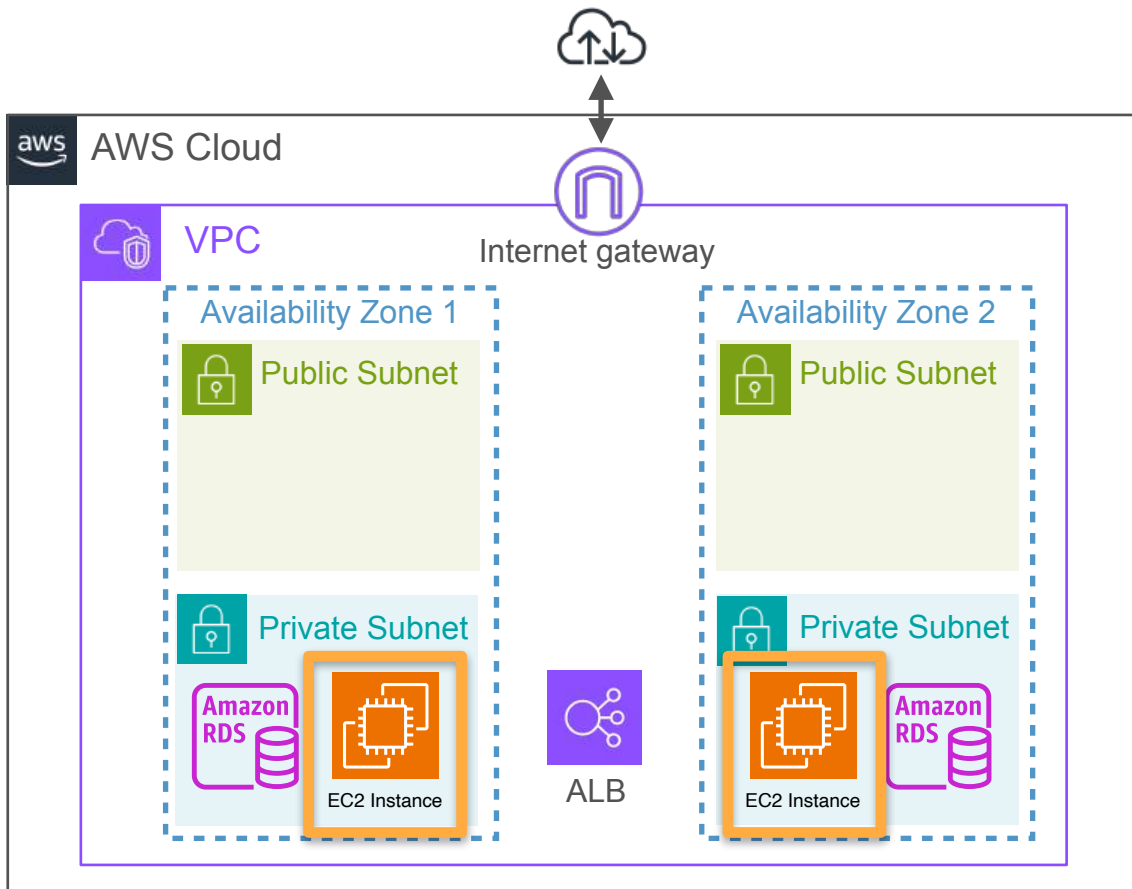
Considerations

1. Applications running on EC2 need to occasionally download updates from resources on the internet
 - Upgrades, patching
2. Need a way to submit requests to the application running on the EC2 instance



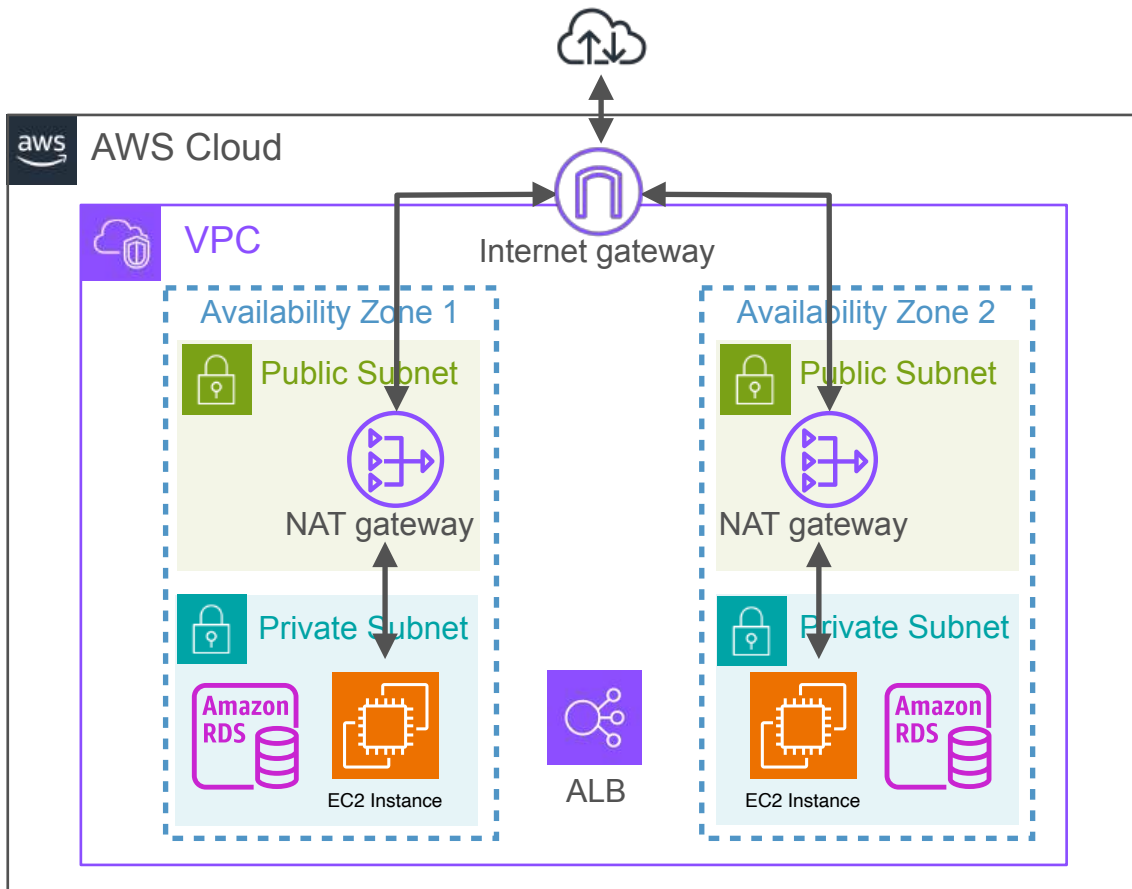
Considerations

1. Applications running on EC2 need to occasionally download updates from resources on the internet
 - Upgrades, patching
2. Need a way to submit requests to the application running on the EC2 instance



Considerations

1. Applications running on EC2 need to occasionally download updates from resources on the internet
 - Upgrades, patching
2. Need a way to submit requests to the application running on the EC2 instance



Considerations

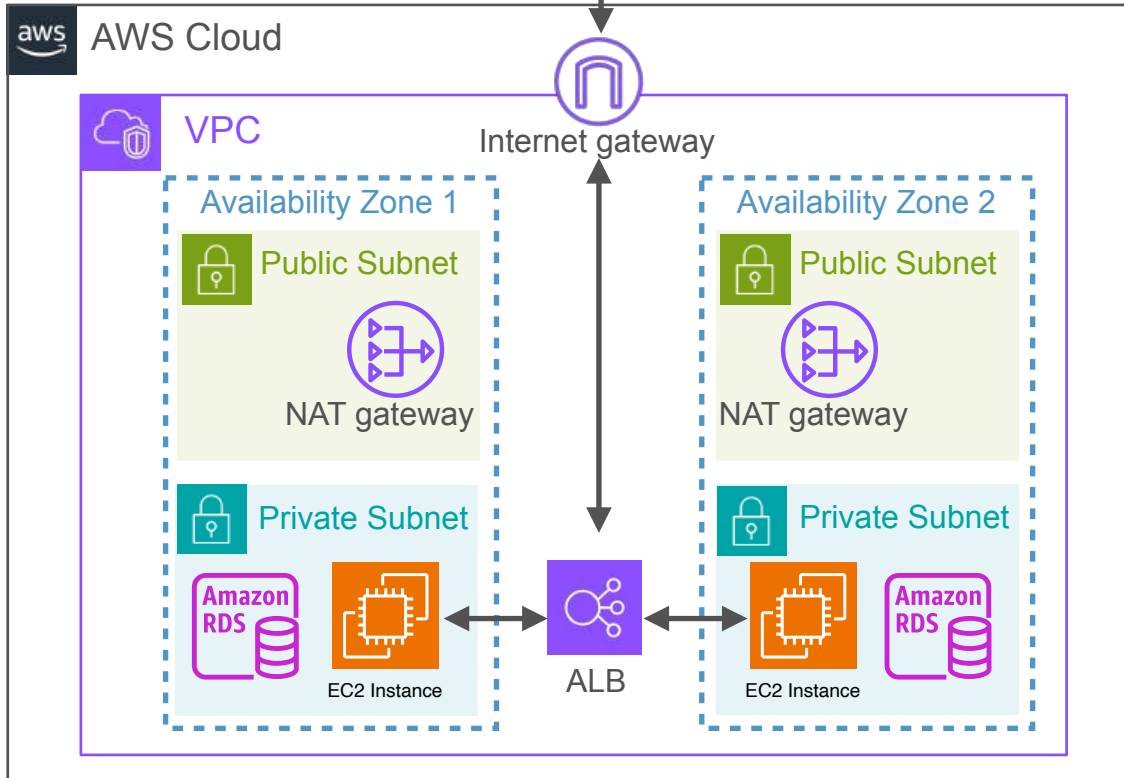
1. Applications running on EC2 need to occasionally download updates from resources on the internet
 - Upgrades, patching
2. Need a way to submit requests to the application running on the EC2 instance

**NAT
Gateway**

**Network Address
Translation Gateway**



- Allows resources in a private subnet to connect to the internet or other AWS services
- Prevents the internet from initiating connections with those resources

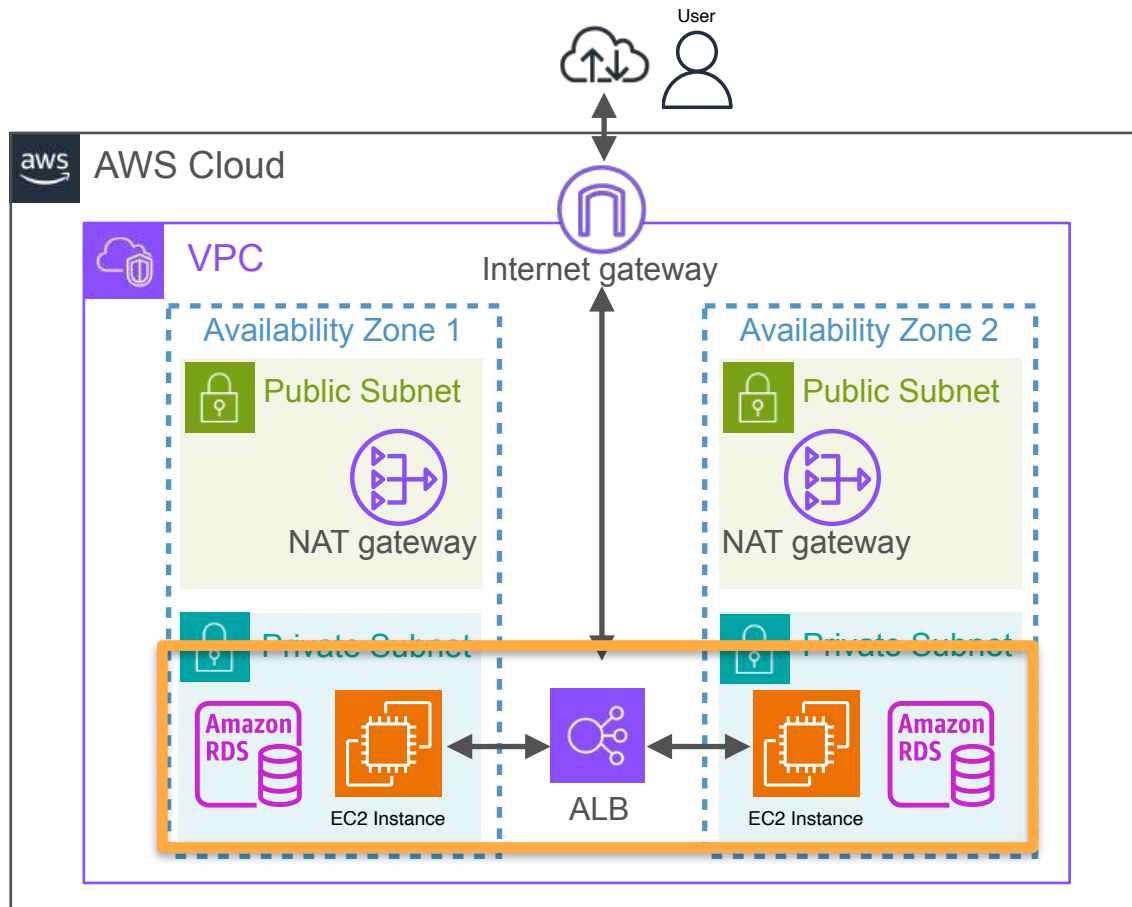


Considerations

1. Applications running on EC2 need to occasionally download updates from resources on the internet
 - Upgrades, patching
2. Need a way to submit requests to the application running on the EC2 instance

ALB:

- Distributes incoming application traffic across multiple backend targets
- Handles the load and ensures the application remains responsive and available
- Keeps those EC2 instances private

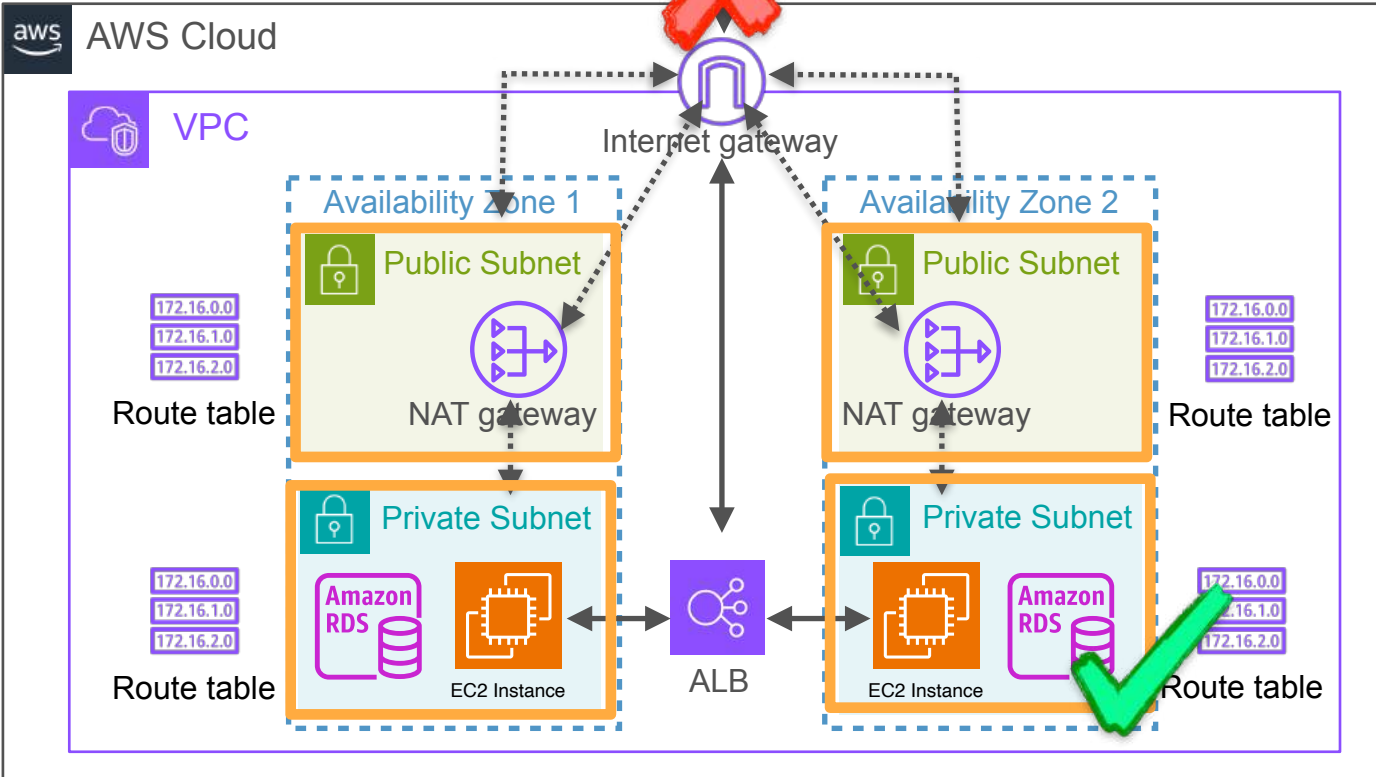




DeepLearning.AI

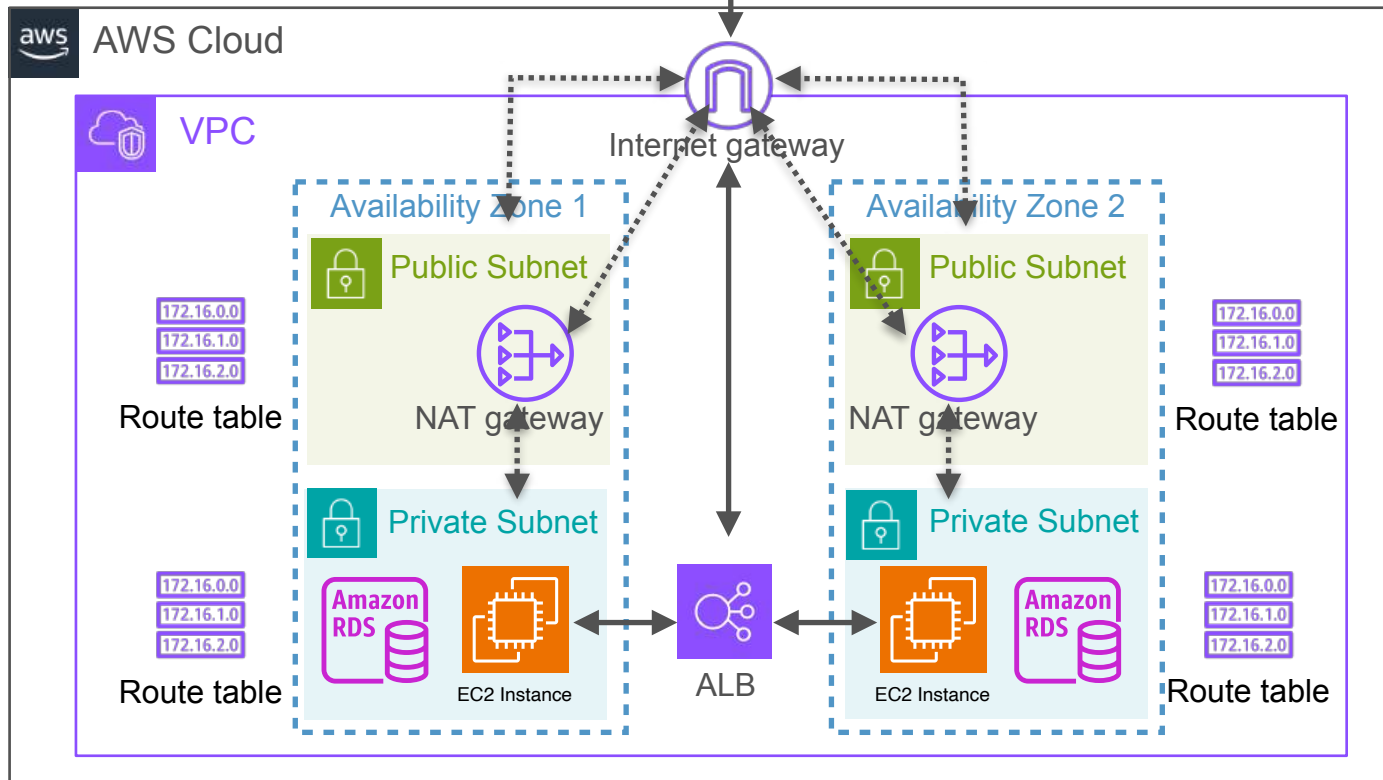
Interacting with Source Systems

AWS Networking - Route Tables



Route Tables

- Essential for directing network traffic within your VPC
- Default route table allows internal communication within the VPC

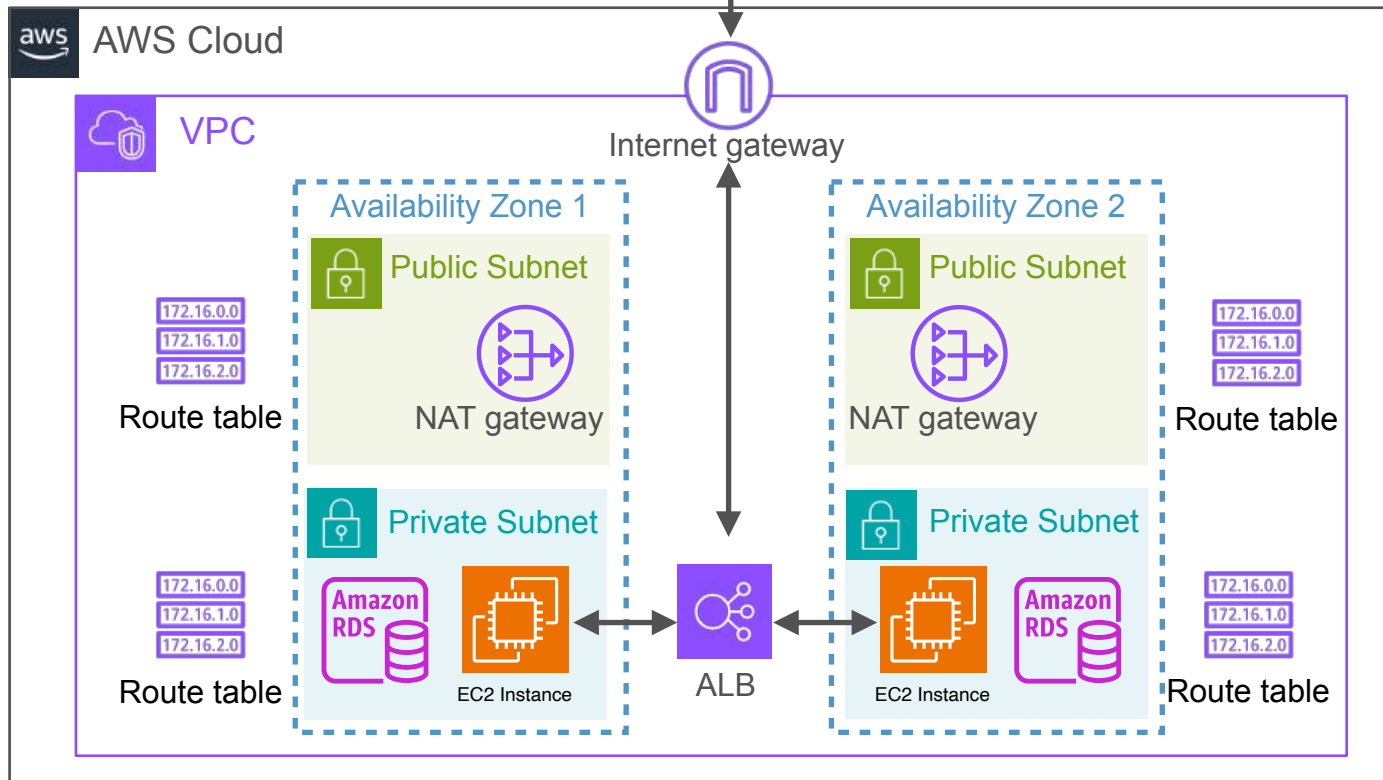




DeepLearning.AI

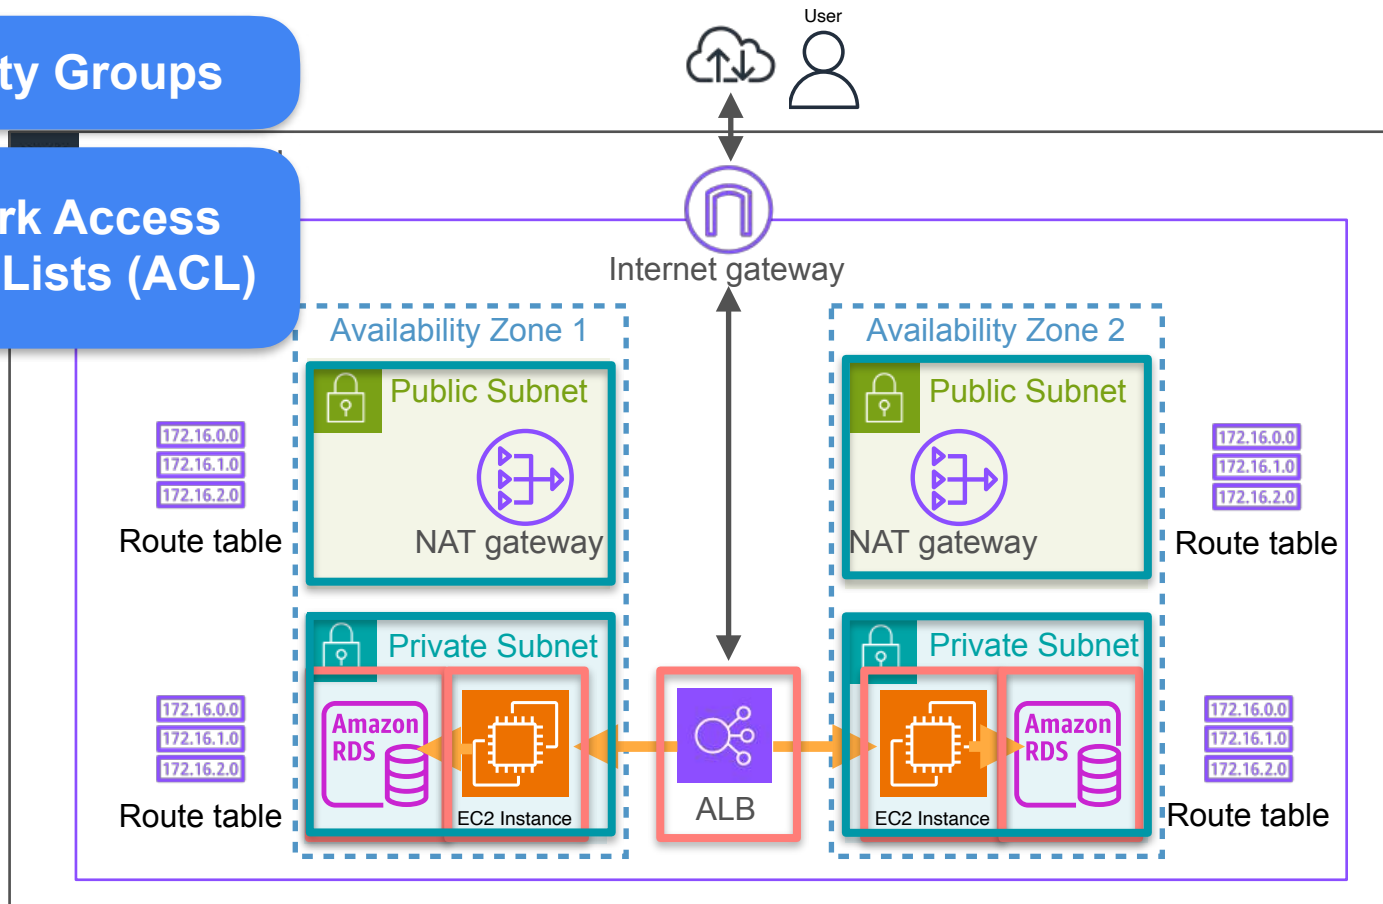
Interacting with Source Systems

AWS Networking - Network ACLs & Security Groups



Security Groups

Network Access Control Lists (ACL)



Security Groups

Instance level virtual firewalls, controlling both inbound and outbound traffic



deny



inbound traffic



outbound traffic

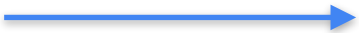


allow

inbound HTTP
requests



outbound HTTP
responses



Port 80

Inbound Rules

- Determine what types of traffic you want to allow
- Where you want to allow that traffic to come from

Security Groups are stateful

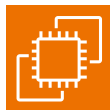
- Allow inbound traffic to an instance automatically allows the return traffic

Security Groups

Instance level virtual firewalls, controlling both inbound and outbound traffic



ALB



EC2 Instance



Security group ID: sg-123

Source	Protocol	Port
0.0.0.0/0 (internet)	HTTP	80
0.0.0.0/0 (internet)	HTTPS	443

Security group ID: sg-456

Source	Protocol	Port
sg-123 (ALB)	HTTP	80
sg-123 (ALB)	HTTPS	443

Security group ID: sg-789

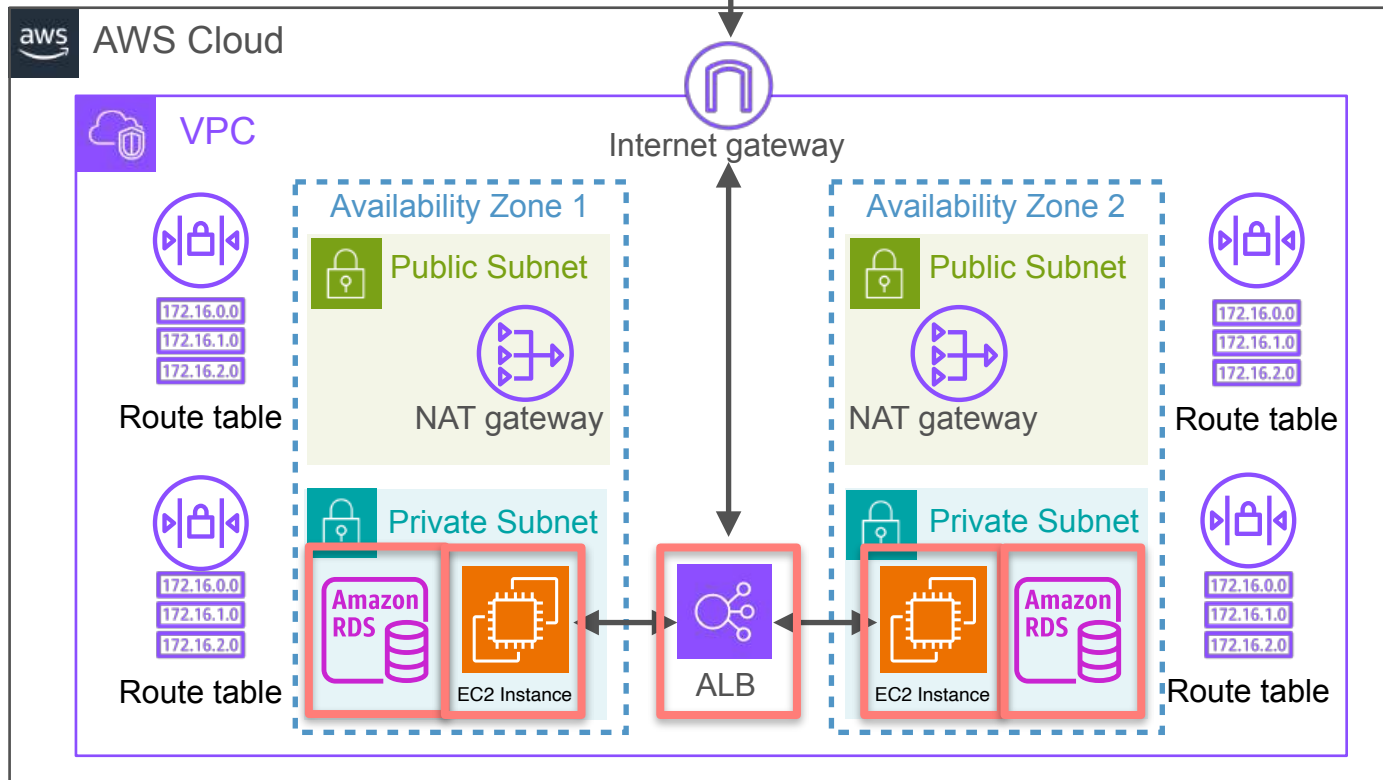
Source	Protocol	Port
sg-456 (EC2)	TCP	3306

Security Group Chaining

Network Access Control Lists (ACL)



- They provide an additional layer of security at the subnet level
- Network ACLs are stateless
- You need to define both inbound and outbound rules explicitly
- Useful for implementing security policies at the subnet level



VPCs and subnets



- Give you a way to define a private network on AWS.

Route Tables



- Direct traffic within the VPC and to the internet.

Public Subnets



Public subnet

Private Subnets



Private subnet

Internet Gateway



- Allow resources within public subnets to access the internet.

NAT Gateway



- Enable instances to initiate outbound connections securely.

Security Groups



- They act as virtual firewalls at the instance level
- They control both inbound and outbound traffic
- They are stateful

Network ACLs



- They provide an additional layer of security at the subnet level
- They are stateless, ie. require explicit rules for both inbound and outbound traffic

If you encounter connectivity issues:

1. Verify that your VPC has an internet gateway properly attached
2. Verify that the route tables have appropriate rules to direct traffic correctly
3. Verify that the route table associations with the subnets are configured correctly
4. Check security groups to make sure they have the needed rules in place
5. Review network ACLs to confirm they allow the necessary traffic
6. Double-check instance configurations to ensure they are associated with the correct security groups and subnets

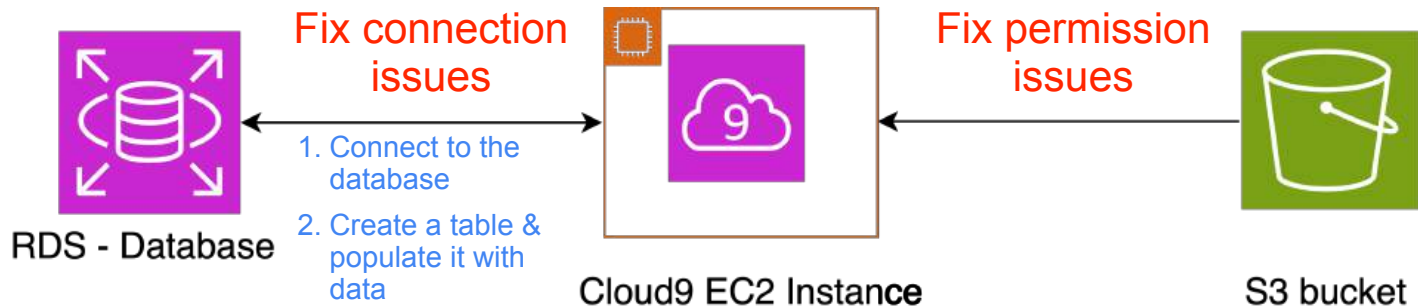


DeepLearning.AI

Connecting to Source Systems

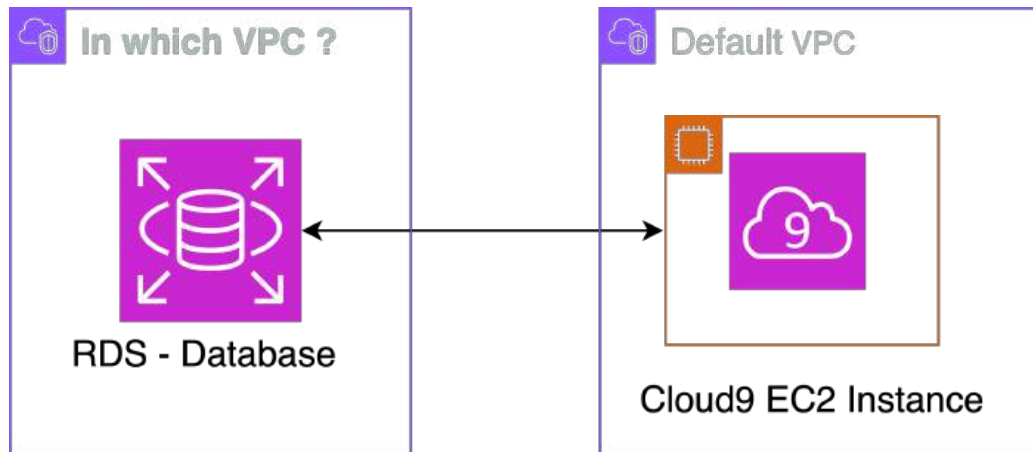
Lab Walkthrough - Database Connectivity and Troubleshooting on AWS

Database Connectivity and Troubleshooting on AWS



- Skip this video, jump straight into the lab and go for it
 - The lab instructions contain hints
- Or, start the lab and follow along with me as you go through this video.
 - When an issue occurs, I'll be inviting you to pause the video
 - After that, I'll show you how to fix it.

Database Connectivity and Troubleshooting on AWS





DeepLearning.AI

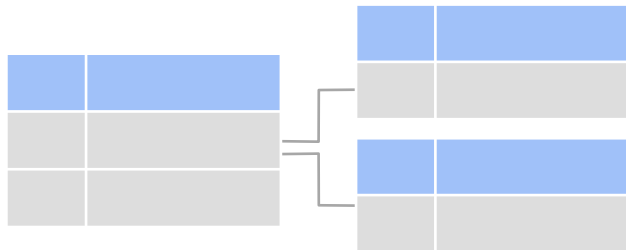
Working with Source Systems

Week 1 Summary

Week 1 Summary

Understand how source systems work

Relational databases



NoSQL databases



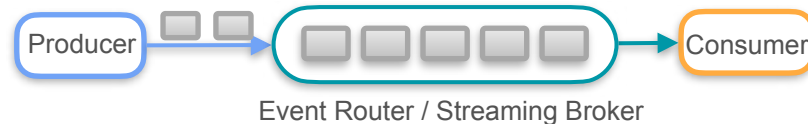
Object Storage



Logs

user id	action	status	timestamp
67945	user added a product x to their cart	success	01-01-2025:10.30
38910	invalid values typed for product quantity	fail	01-01-2025:10.32

Streaming Systems



Week 1 Summary

- How to connect to data sources
- Basics of networking
- Importance of IAM in ensuring security in source systems

Week 1 Summary

