

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI

# Data Storage and Queries

---

## **Storage Systems**



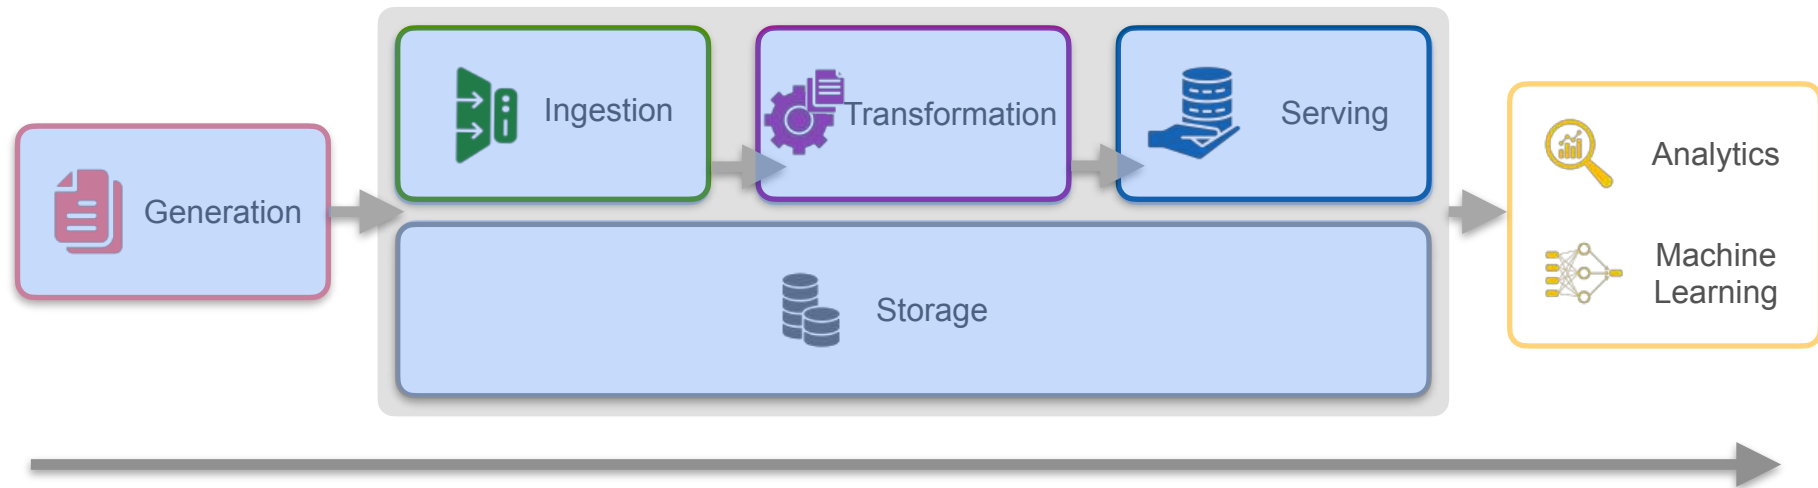
DeepLearning.AI

# Storage Systems

---

## Course 3 Overview

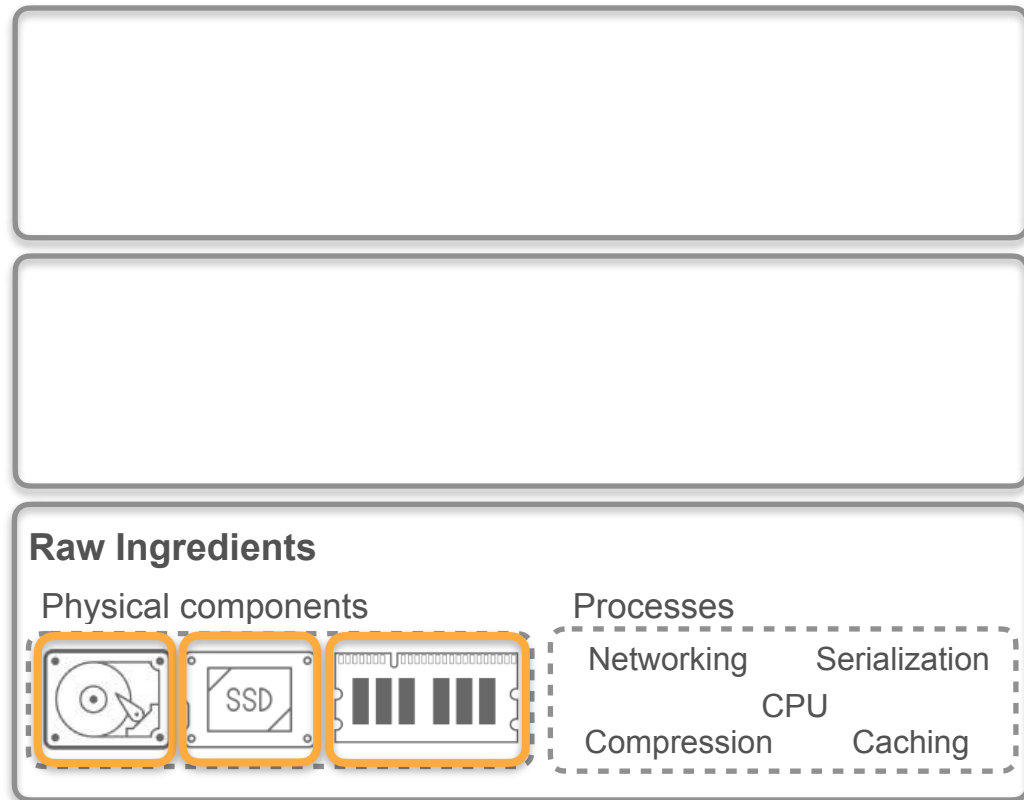
# Storage



Storage solution considerations:

- Data type
- Data size
- Data format
- Access and update pattern

# Storage Hierarchy



# Storage Hierarchy

## Management system:

Organizes data in the raw components and allows you to interact with the stored data



## OLTP Systems

Online Transactional Processing Systems  
Focus on performing read and write queries with low latency

## OLAP Systems

Online Analytical Processing Systems  
Focus on applying analytical activities on data (e.g. aggregation, summarization)

### Storage Systems



# Storage Hierarchy

## Management system:

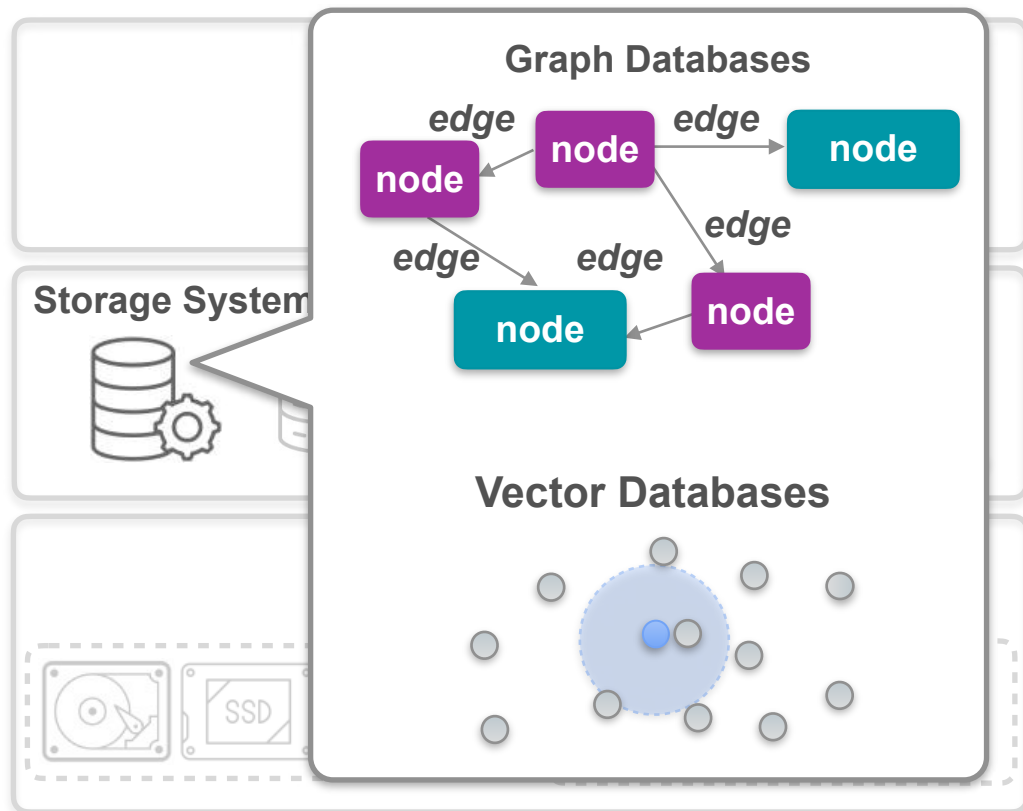
Organizes data in the raw components and allows you to interact with the stored data

### OLTP Systems

Online Transactional Processing Systems  
Focus on performing read and write queries with low latency


### OLAP Systems

Online Analytical Processing Systems  
Focus on applying analytical activities on data (e.g. aggregation, summarization)



# Storage Hierarchy

## Storage Abstractions





# Course 3 Week 1



## Trade-offs between storage cost and performance

- Cloud storage paradigms (block, object and file storage)
- Data storage in databases
  - Row vs column-oriented databases
  - Graph and vector databases
- Characteristics of physical components
- Serialization and compression

### Storage Abstractions




### Storage Systems



### Raw Ingredients

#### Physical components



#### Processes

Networking

Serialization

CPU

Compression

Caching



# Course 3 Week 2

How to choose the appropriate  
abstractions for storing your data

## Storage Abstractions




## Storage Systems



## Raw Ingredients

### Physical components



### Processes

Networking

Serialization

CPU

Compression

Caching

# Course 3 Week 3



## Queries

- How queries work
- How different storage solutions affect query performance
- Techniques for improving query performance

### Storage Abstractions




### Storage Systems



### Raw Ingredients

#### Physical components



#### Processes

Networking

Serialization

CPU

Compression

Caching



DeepLearning.AI

# Storage Systems

---

**Raw Ingredients:  
Physical Components of Data Storage**

# Raw Storage Ingredients

## Persistent Storage Medium

Magnetic disk



Solid-state storage



## Volatile Memory


RAM



CPU cache




# Magnetic Disks



Hard Disk Drives (HDDs)

# Magnetic Disks



**Track + Sector  
= Address**

## **Write:**


Encode binary data by changing the magnetic field

## **Read:**

Converts magnetic field into binary data

Hard Disk Drives (HDDs)

# Solid State Drives



Solid-State Drives (SSDs)



SSDs read and write data  
much faster



# Performance Comparison

	Magnetic Disk	SSD
Latency	4 milliseconds	
IOPS (Input/output operations per second)	Hundreds	

**Commercial magnetic disk drive:**  
Rotates at 7200 revs/min



# Performance Comparison



	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands



**Solid-State Drives (SSDs)**




Electrical charges

# Performance Comparison



	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands
Data Transfer Speed (number of bytes read/ written from disk to memory in a second)	Up to 300 MB/s	4 GB/s


# Improving Performance



Data transfer speed limited by network performance


# Improving Performance

## Distributed Storage



Data transfer speed limited by network performance

## Partitioning



Slicing SSDs into partitions

# Performance Comparison




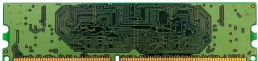



	Magnetic Disk	SSD
Latency	4 milliseconds	0.1 milliseconds
IOPS (Input/output operations per second)	Hundreds	Tens of thousands
Data Transfer Speed (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s
Cost	\$0.03–0.06/GB	\$0.08–0.10/GB

2-3 times cheaper

# Volatile Memory Ingredients

*\*\*Note: these metrics can vary*

	  <b>Magnetic Disk</b>	 <b>SSD</b>	 <b>RAM</b> (Random Access Memory)	 <b>CPU Cache</b>
<b>Latency</b>	4 milliseconds	0.1 milliseconds	0.1 microseconds	
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands	Millions	
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	
<b>Cost</b>	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	

30-50 times more expensive

# Volatile Memory Ingredients

*\*\*Note: these metrics can vary*




	<b>Magnetic Disk</b>	<b>SSD</b>	<b>RAM</b> (Random Access Memory)	<b>CPU Cache</b>
<b>Latency</b>	4 milliseconds	0.1 milliseconds	0.1 microseconds	1 nanosecond
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands	Millions	/
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	1 TB/s
<b>Cost</b>	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	/



# CPU Cache Use Cases

- CPU caching
- Store frequently and recently accessed data in a fast access layer

## Examples






DeepLearning.AI

# Storage Systems

---

**Raw Ingredients:  
Processes Required for Data Storage**

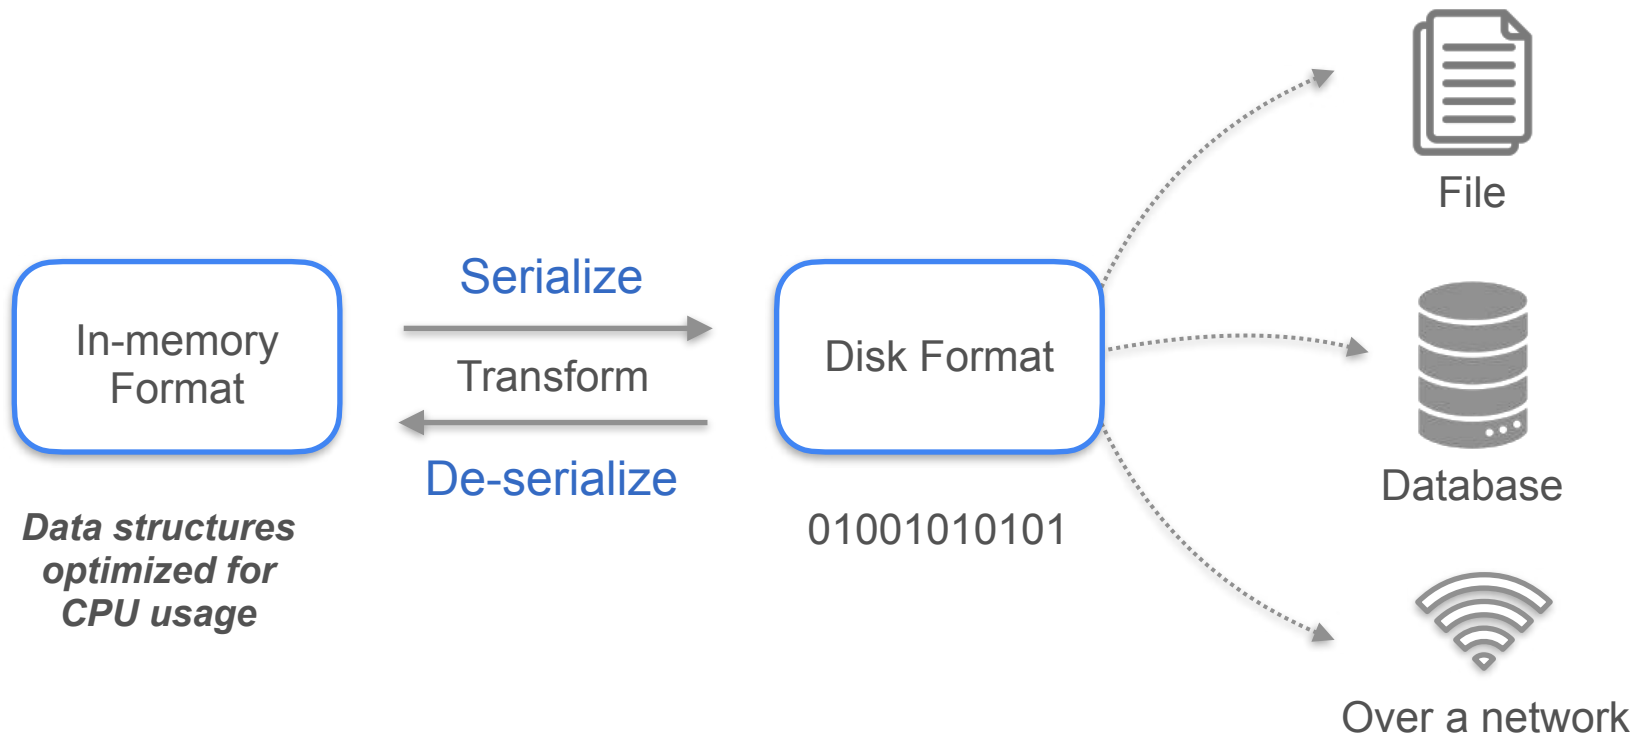
# Networking and CPU — “Raw Ingredients” of Storage Systems



## Enhance:

- read and write performance
- data durability
- data availability

# Serialization



# Serialization



Transactional operations

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q



Physical Storage



## Row-Based Serialization

bytes representing the 1st object	bytes representing the 2nd object	...	bytes representing the last object
-----------------------------------	-----------------------------------	-----	------------------------------------

## Column-Based Serialization

--	--	--	--

# Serialization

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Analytical queries

Physical Storage



## Row-Based Serialization

bytes representing the 1st row	bytes representing the 2nd row	...	bytes representing the last row
--------------------------------	--------------------------------	-----	---------------------------------

## Column-Based Serialization

bytes representing the 1st	key	bytes representing the 2nd	key	...	bytes representing the last	key
----------------------------	-----	----------------------------	-----	-----	-----------------------------	-----

# Serialization Formats

## Human-Readable Textual Formats



## Binary Formats



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling





# Serialization Formats

## Human-Readable Textual Formats



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



# Serialization Formats

## Human-Readable Textual Formats



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs



# Serialization Formats



## Binary Formats



- Column-based format
- For efficient storage and big data processing



# Serialization Formats



## Binary Formats



- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

## Binary Formats

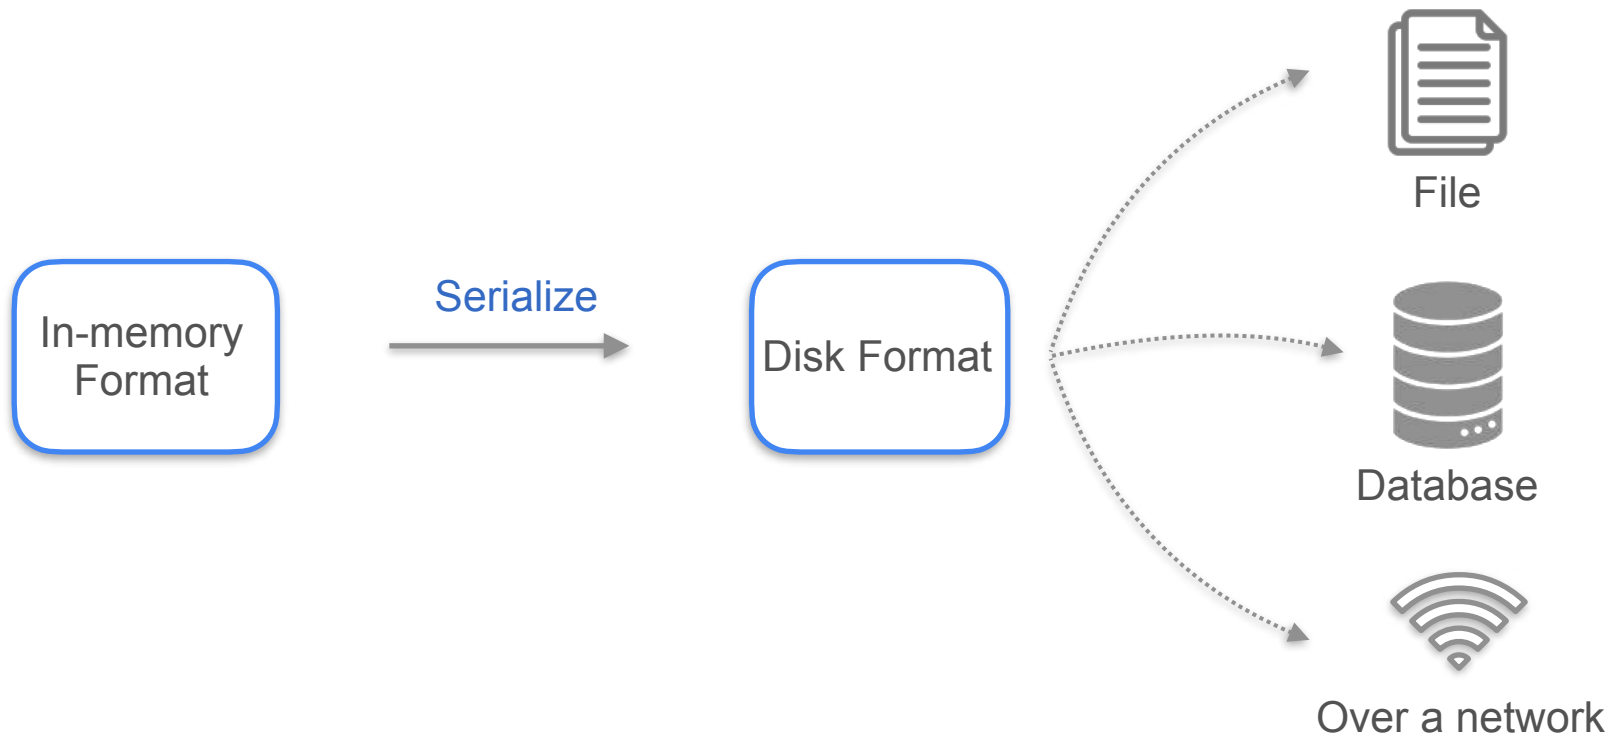


- Column-based format
- For efficient storage and big data processing

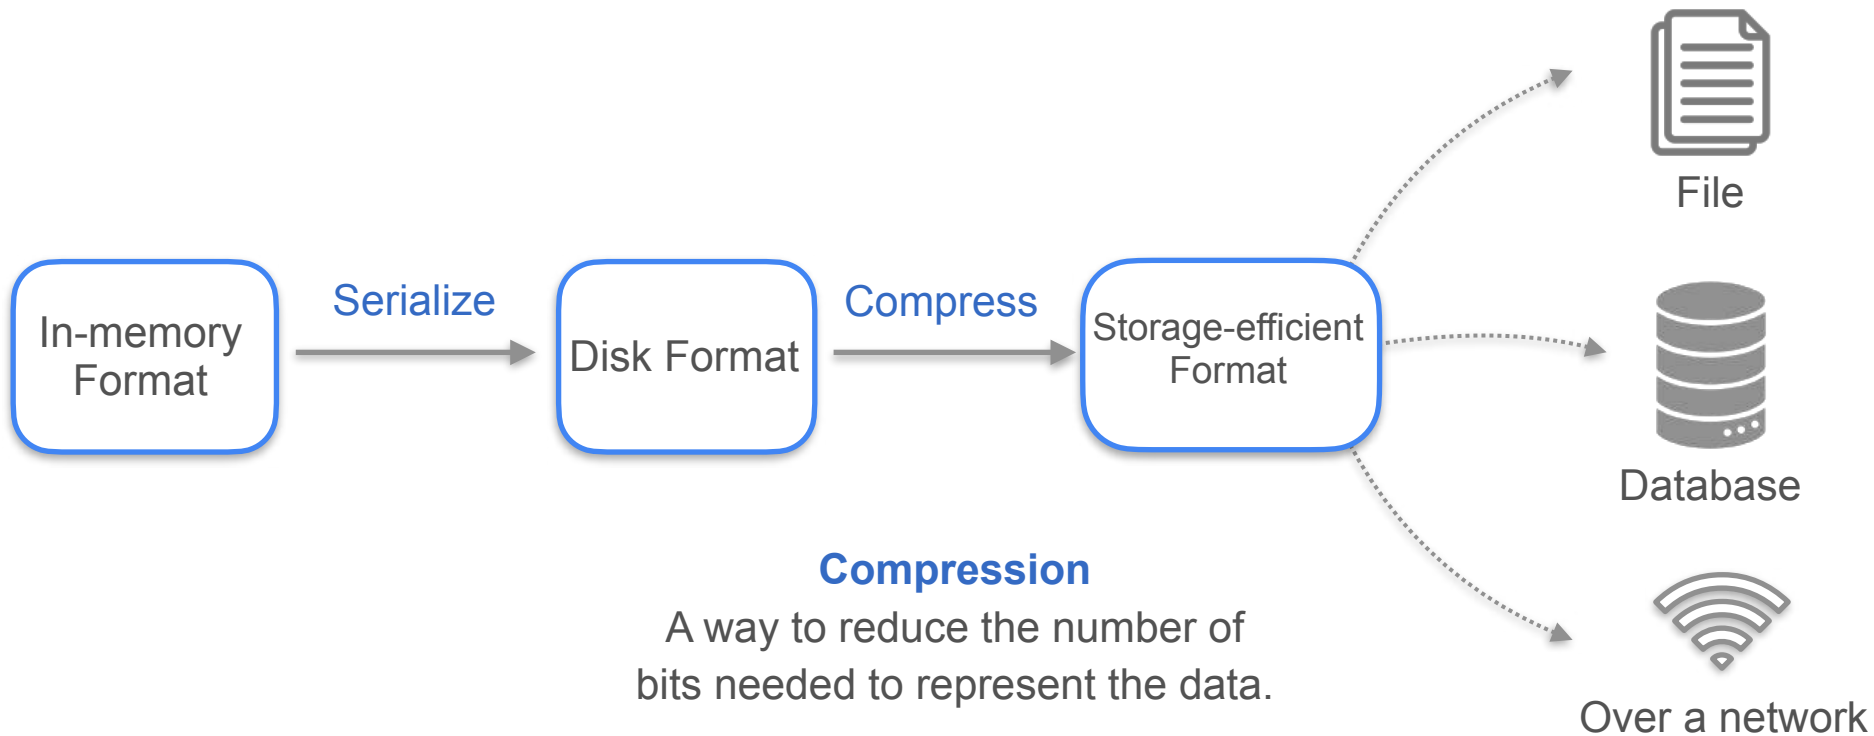


- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

# Serialization

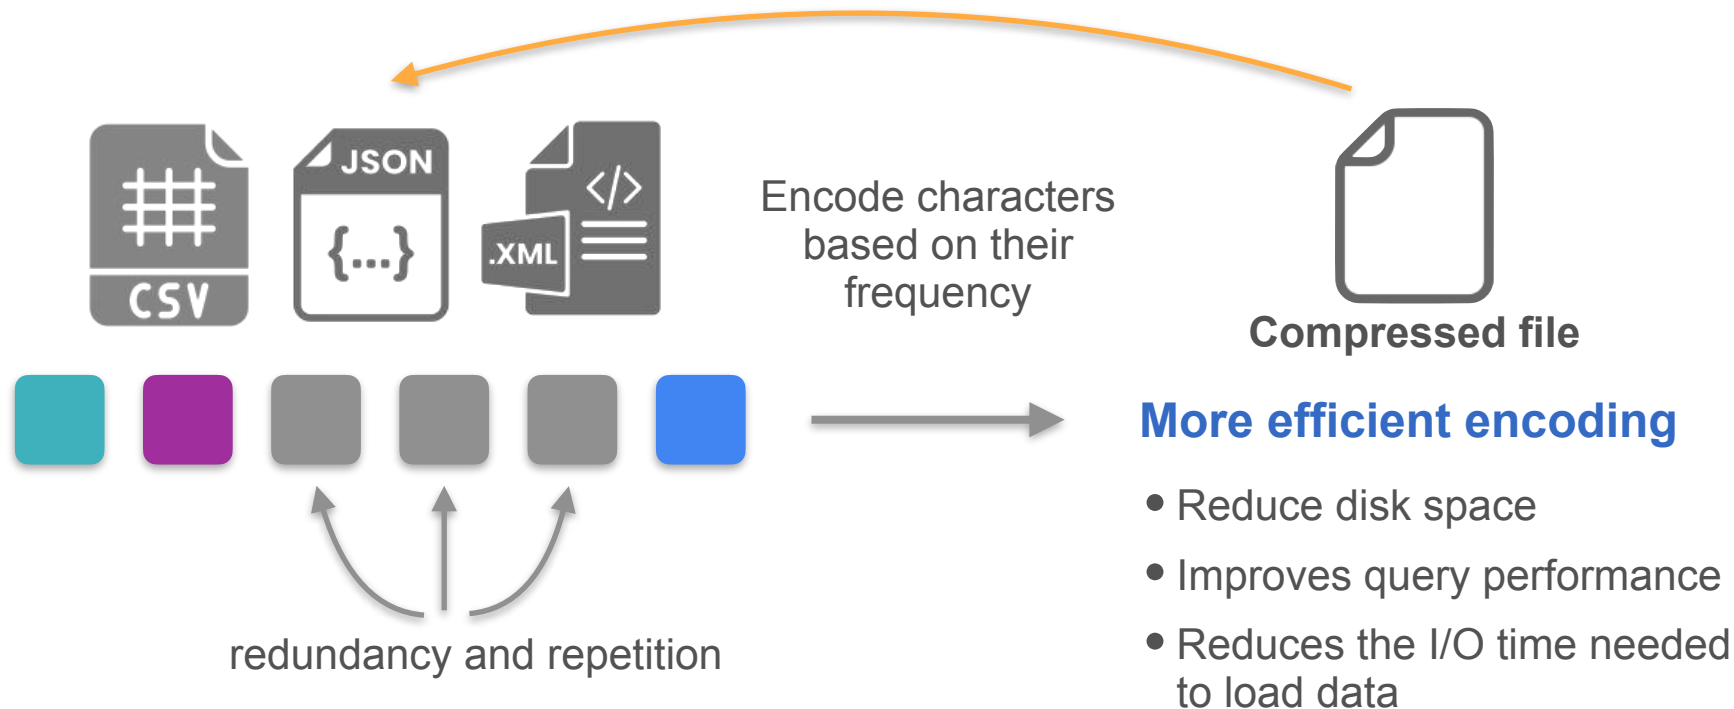


# Serialization



# Compression

Compression ratio







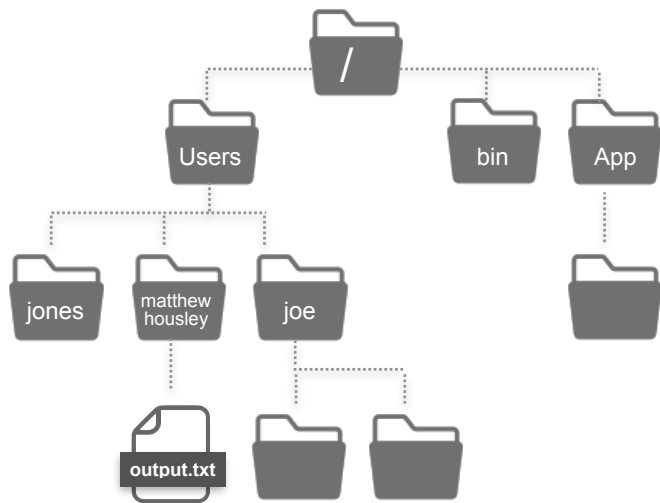
DeepLearning.AI

# Storage Systems

---

**Cloud Storage Options:  
Block, Object and File storage**

# File Storage



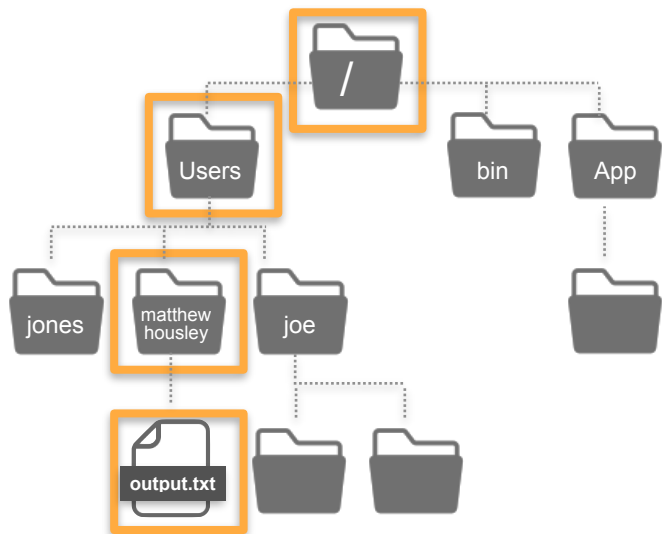
## File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

# File Storage



/Users/matthewhousley/output.txt


## File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

# File Storage Use Cases




## Cloud File Storage Service




Amazon Elastic File System  
(EFS)

- Provides you access to shared files over a network
- Networking, scaling, and configuration are handled by the cloud vendor

# Block Storage



stored on




## Block Storage

Divides files into small, fixed-size blocks of data and stores them on disk

- Each block has a unique identifier
- You can efficiently retrieve and modify data in individual blocks
- You can distribute blocks of data across multiple storage disks
  - Higher scalability
  - Stronger data durability

# Block Storage




Lookup Table

File Piece	Block Identifier
First piece	1232
Second piece	1234
Third piece	1236
Fourth piece	1238

# Block Storage Use Cases

- Ideal for frequent access and modification
- Enables OLTP systems to perform small and frequent read and write operations with low latency
- Provides persistent storage for virtual machines



Attach a root storage device backed by a block storage volume

## Default storage for EC2



Amazon Elastic Block Store (EBS)

1. SSD for latency-sensitive workloads
2. Magnetic disks to store infrequently-accessed data

# Object Storage

## Object Storage

Stores immutable files as data objects in a flat structure

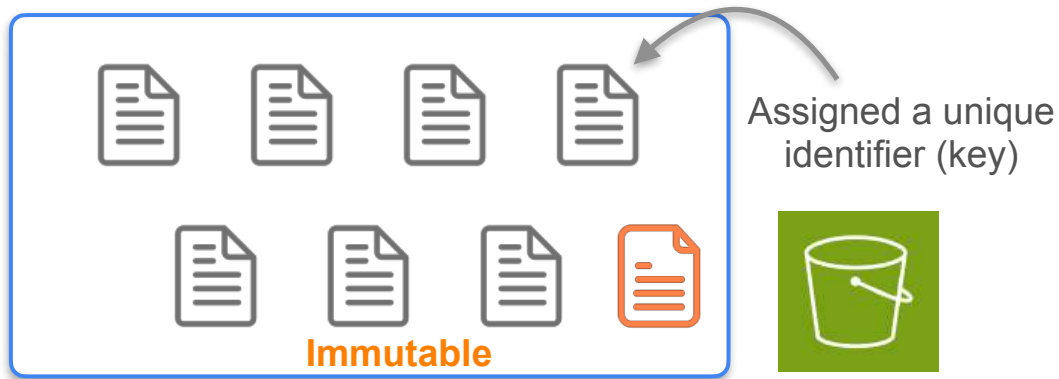




# Object Storage

## Object Storage

Stores immutable files as data objects in a flat structure




s3://o'reilly-data-engineering-book/data-example.json

**The Bucket**

**Object key**

- Each object is immutable
- To update the file you have to re-write the entire object
- Can scale horizontally and support performant parallel operations



# Object Storage Use Cases



Ideal for...	Not ideal for...
<ul style="list-style-type: none"><li>• Storage layer of cloud data warehouses or data lakes</li><li>• Storing data needed in OLAP systems</li><li>• Machine learning pipelines<ul style="list-style-type: none"><li>• Raw text</li><li>• Images</li><li>• Videos</li><li>• Audio</li></ul></li></ul>	<ul style="list-style-type: none"><li>• Not good at supporting transactional workloads</li></ul>



# Cloud Storage Options

File Storage	Block Storage	Object Storage
<ul style="list-style-type: none"><li>• Supports data sharing</li><li>• Easy to manage with low performance and scalability requirements</li></ul>	<ul style="list-style-type: none"><li>• Supports transactional workloads</li><li>• Allows frequent read and write operations with low latency</li></ul>	<ul style="list-style-type: none"><li>• Supports analytical queries on massive datasets</li><li>• Offers high scalability and parallel data processing</li></ul>



DeepLearning.AI

# Storage Systems

---

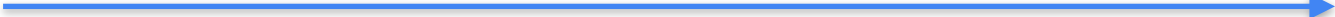
## **Storage Tiers – Hot, Warm, & Cold Data**

# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency			
Example			
Storage Medium			
Storage Cost			
Retrieval Cost			

# Storage Tiers

## Hot Storage



<b>Access Frequency</b>	Very frequent		
<b>Example</b>	Product recommendation application		
<b>Storage Medium</b>	SSD & Memory		
<b>Storage Cost</b>	High		
<b>Retrieval Cost</b>	Low		

# Storage Tiers

## Warm Storage

<b>Access Frequency</b>		Less frequent	
<b>Example</b>		Regular reports and analyses	
<b>Storage Medium</b>		Magnetic disks or hybrid storage systems	
<b>Storage Cost</b>		Medium	
<b>Retrieval Cost</b>		Medium	

# Storage Tiers

Cold Storage

<b>Access Frequency</b>			Infrequent
<b>Example</b>			Archive
<b>Storage Medium</b>			Low-cost magnetic disks
<b>Storage Cost</b>			Low
<b>Retrieval Cost</b>			High



# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
<b>Access Frequency</b>	Very frequent	Less frequent	Infrequent
<b>Example</b>	Product recommendation application	Regular reports and analyses	Archive
<b>Storage Medium</b>	SSD & Memory	Magnetic disks or hybrid storage systems	Low-cost magnetic disks
<b>Storage Cost</b>	High	Medium	Low
<b>Retrieval Cost</b>	Low	Medium	High



# AWS Storage Tiers



Amazon S3

**Access Frequency**

**Hot Storage**



S3 Express  
One Zone



S3 Standard



S3 Standard-  
IA



S3 One  
Zone-IA



S3 Glacier  
Flexible  
Retrieval



S3 Glacier  
Deep  
Archive



S3 Glacier  
Instant  
Retrieval




DeepLearning.AI

# Storage Systems


---

## **Distributed Storage Systems**

# How Distributed Storage Systems Work




# How Distributed Storage Systems Work



Magnetic disk




SSD




**Total Capacity =**      +      +


# How Distributed Storage Systems Work



## Single Machine Storage Architecture




# How Distributed Storage Systems Work




- Higher fault tolerance and data durability
- High availability

# How Distributed Storage Systems Work



- Higher fault tolerance and data durability
- High availability
- Process many read and write operations in parallel
- Fast data access





# Advantages of Distributed Storage Systems

## Distributed Storage Architecture



Object Storage



Cloud  
Data Warehouse



# Methods for Distributing Data

**Replication**

**Partitioning**

# Methods for Distributing Data



**Replication**



High availability and performance

**Partitioning**

# Methods for Distributing Data

**Replication**



High availability and performance

**Partitioning**



**Sharding**



partition or shard


# Methods for Distributing Data

**Replication**



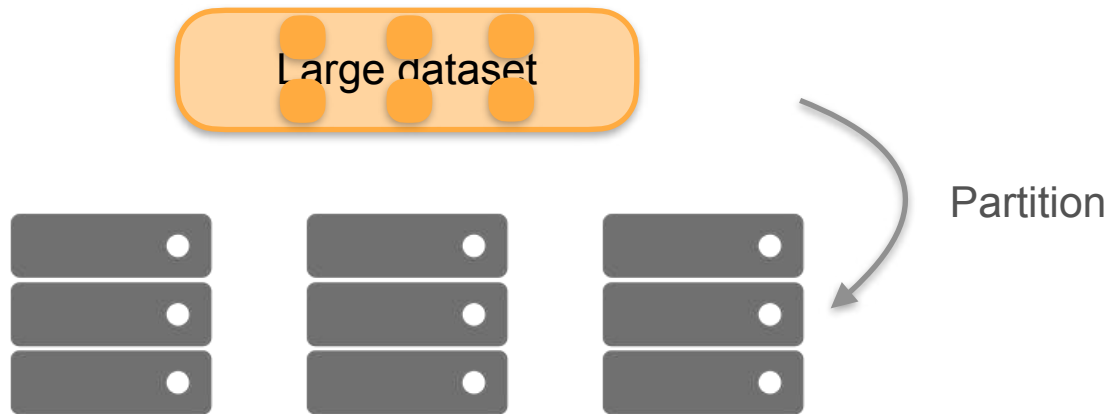
High availability and performance

**Partitioning**

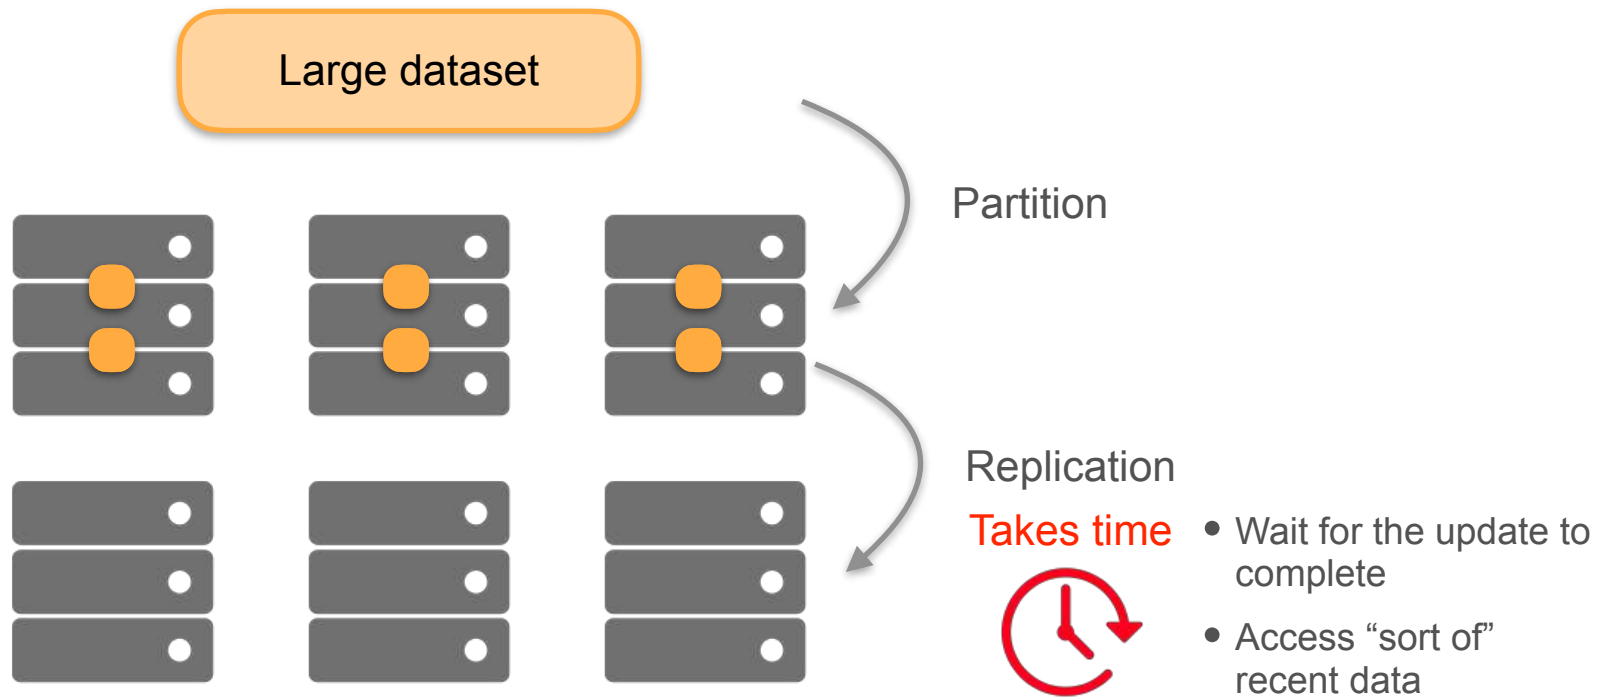


**Sharding**

# Methods for Distributing Data




# Methods for Distributing Data



# Distributed Storage Considerations – CAP Theorem

## The CAP theorem





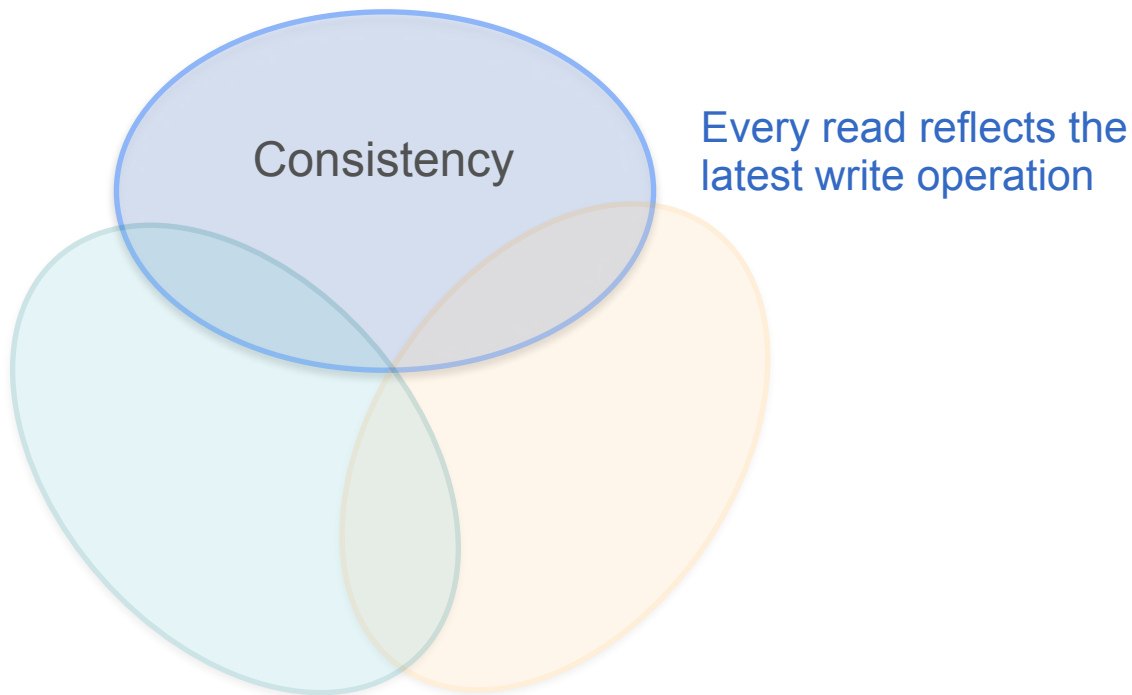
# Distributed Storage Considerations – CAP Theorem

## ACID



### Consistency

Any change to data must follow the set of rules defined by database schema

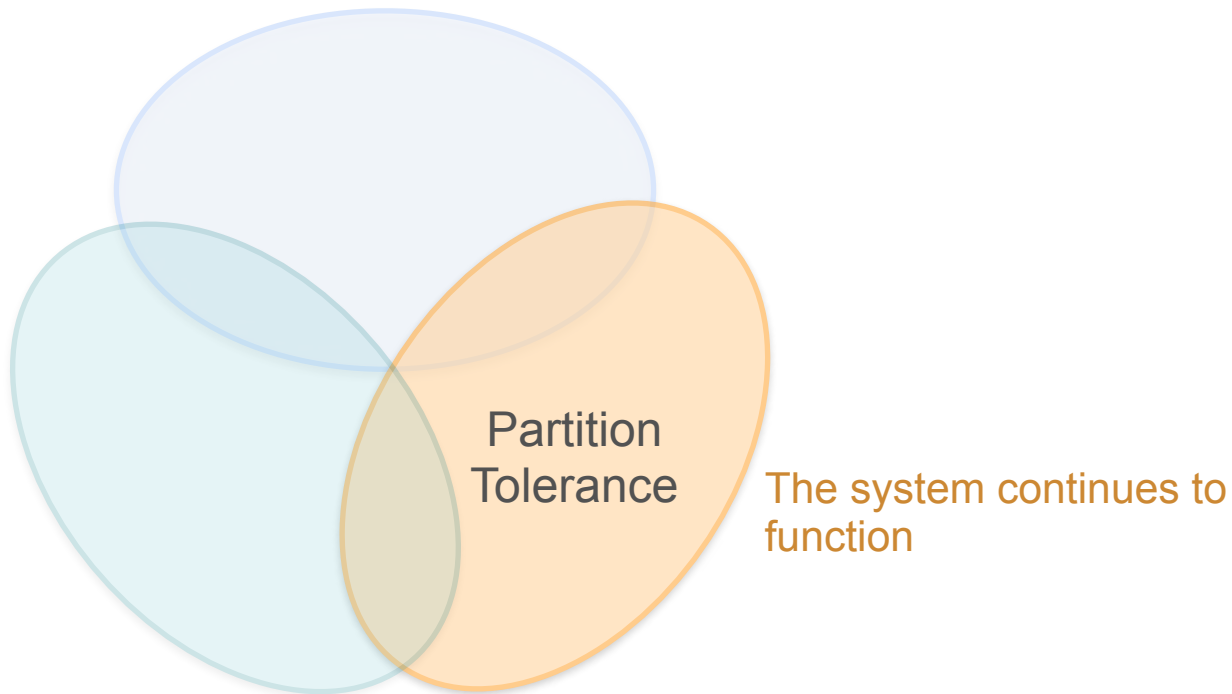


# Distributed Storage Considerations – CAP Theorem




Every request will receive a response

# Distributed Storage Considerations – CAP Theorem



# Distributed Storage Considerations – CAP Theorem

## The CAP theorem



# Distributed Storage Considerations – CAP Theorem

## Scenario:


Accessing a node that's still being updated

## Option 1:

Cancel the request


## Option 2:

Proceed with the read operation



# Distributed storage considerations – ACID vs BASE

**RDBMS**




**ACID**  
compliant

**NoSQL Databases**



**BASE**  
principles



# Distributed storage considerations – ACID vs BASE

## ACID

compliant

**A**tomicity

**C**onsistency

**I**solation

**D**urability

## BASE

principles

**B**asically **A**vailable

**S**oft-state

**E**ventual Consistency

# Scenario

## Course 1



Data Scientist



### Main database instance (Strong consistency)



#### Read-replica of the prod database

- Ingest
- Transform
- Store
- Serve

### Read Replicas in RDS (Eventual consistency)



- Track changes in main database
- Update their own data



Amazon RDS





DeepLearning.AI

## Lab Walkthrough

---

# Comparing Cloud Storage Options

**Object Storage**

**File Storage**

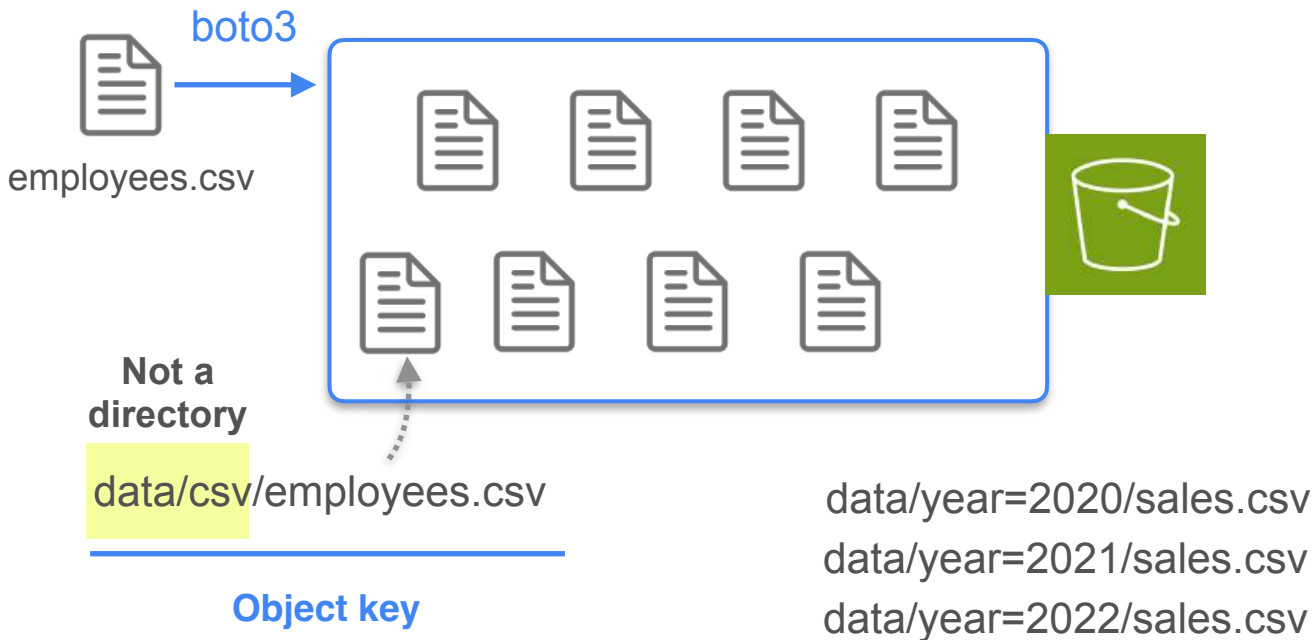
**Block Storage**

**Memory**

## Object Storage

### Flat Structure

### Immutability



# Object Storage

## Flat Structure

## Immutability

2. Modify employees' data

 boto3  
employees.csv



1. Enable versioning

Not a  
directory

data/csv/employees.csv

Object key

3. Use  
list\_object\_version

```
Object Key: data/csv/employees.csv  
Object Version Id: q4A5B9CQZ10.u7.YKA7LabC_5GcBA.uZ  
Is Latest: True  
Last Modified: 2024-08-12 19:57:00+00:00
```

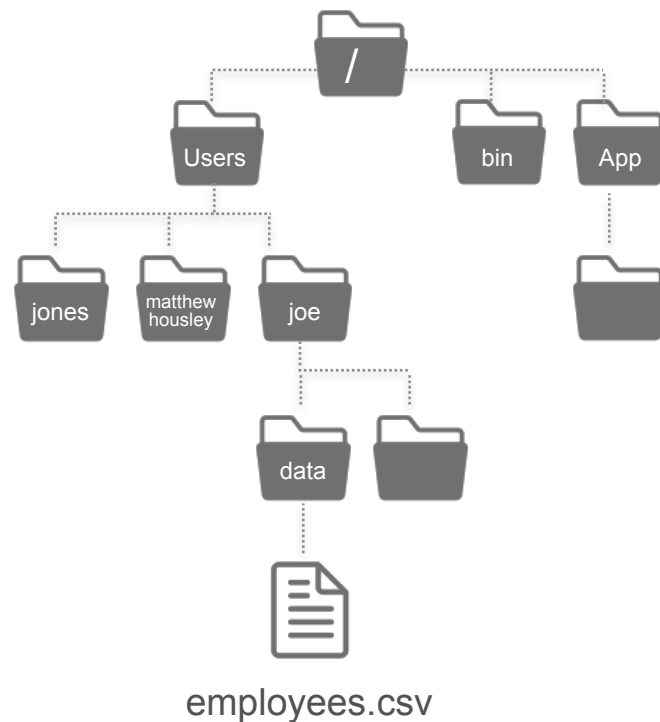
```
Object Key: data/csv/employees.csv  
Object Version Id: WOOPNaVQTFHB13CkIiRATRDCZt30Wq9k  
Is Latest: False  
Last Modified: 2024-08-12 19:39:56+00:00
```

## File Storage

1. Navigate to the “data” directory
2. Explore the directory content and metadata
3. Explore how the data is modified in place

**A directory**

data/employees.csv




## Block Storage

1. Connect to the server
2. Send a file to the server



Server that emulates the behavior of block storage



Transferring data from memory is faster than transferring data from disk.



### Use the cache-pandas package:

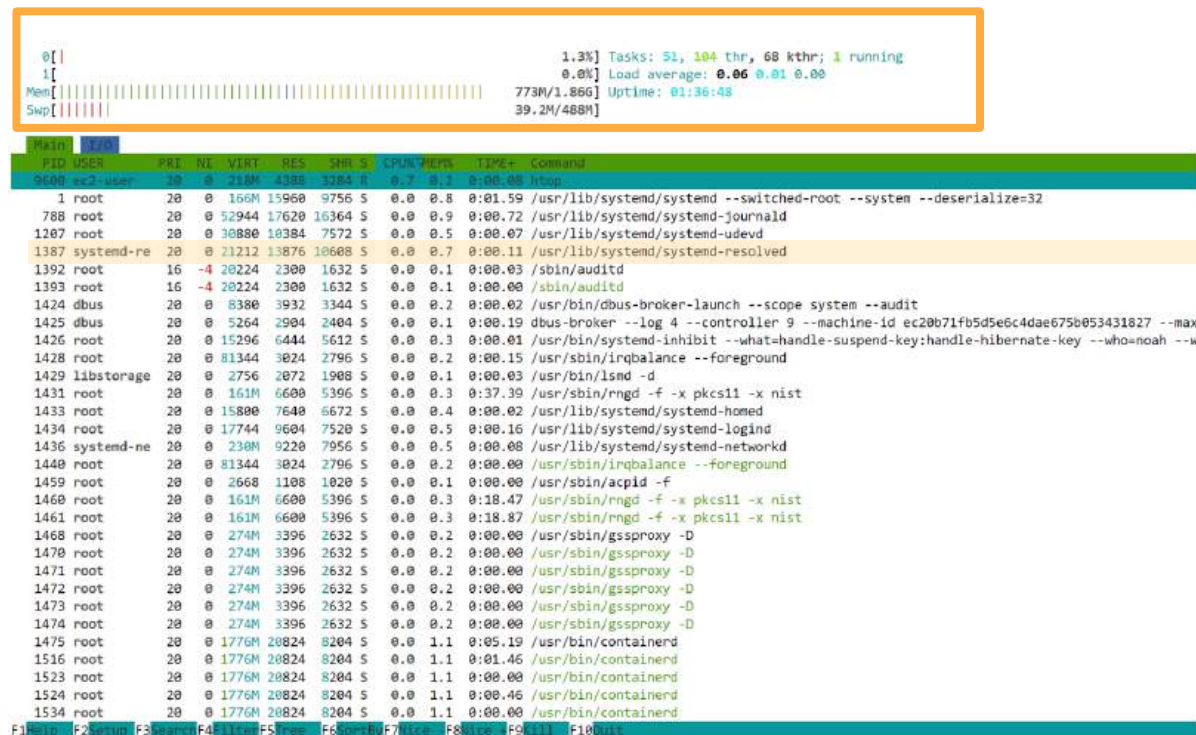

- provides the “timed\_LRU\_cache” decorator to easily cache in memory pandas DataFrames generated by functions.

```
@timed_lru_cache(seconds=100, maxsize=None)
def read_csv_to_memory(path: str) -> pd.DataFrame:
    """Read CSV function with a cache decorator."""
    return pd.read_csv(path)
```

Memory

Compare the time it takes to read the file for the first time with the time it takes to read the same data stored in memory.

# Monitor your memory storage capacity using htop command:



Memory



**Object Storage**

**File Storage**

**Block Storage**

**Memory**

Explore the features of these storage options.



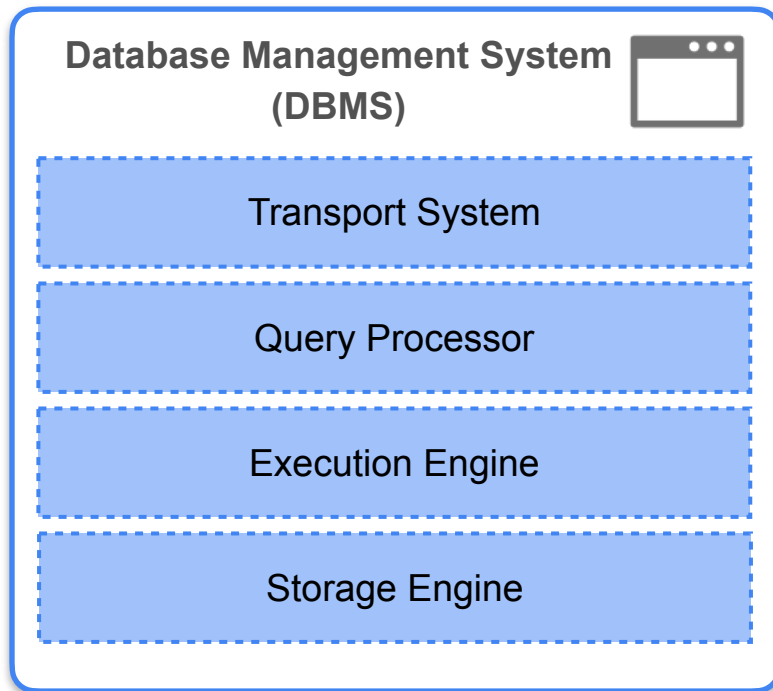
DeepLearning.AI

# Storage Systems

---

## **How Databases Store Data**

# Database Management System



# Database Management System

## Storage Engine

- Serialization
- Arrangement of data on disk
- Indexing

## Modern Storage Engines

- Support the performance characteristics of SSDs
- Handle modern data types and structures
- Offer robust columnar storage support

## Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

.....

## Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

### Index

A data structure that helps you efficiently locate data

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

.....

scan all rows

## Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

\*\*\*Scanning all rows:  $O(n)$

Binary search on rows:  $O(\log n)$

## Index


A data structure that helps you efficiently locate data

### Index table

Country	Row Address
Brazil	###
Canada	###
Canada	###
Canada	###
Canada	###
Canada	###
Mexico	###
USA	###
USA	###
USA	###

Use binary search to locate the USA rows

# In-Memory Storage Systems



- Excellent transfer speed and low latency
- Volatile
- Used to present data for ultra-fast retrieval:
  - Caching applications
  - Real-time bidding
  - Gaming leaderboards

## 1. Memcached

- Key-value store to cache database query results or API calls
- Used when it's acceptable for data to be lost

## 2. Redis

- Key-value store that supports more complex data types
- Supports high-performance applications that can tolerate minor data loss





DeepLearning.AI

# Storage Systems

---

## Row vs Column Storage

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

↓ Stores data  
row by row

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
bytes representing the 1st row					bytes representing the 2nd row					bytes representing the last row					

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

**Row Storage is perfect for OLTP**

*Perform read and write operations with low latency*

← Locate this order

↓ Stores data  
row by row

**Physical Storage**

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q



Stores data  
row by row

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----

**Analytical queries** focus on summarizing or aggregating columns

- Total revenue?
- Most popular product?
- Average quantity?

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

```
SELECT SUM(price)
FROM my_table
```

1 million rows

30 columns

100 bytes per entry

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q	...
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----	-----

# Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 million rows

30 columns

100 bytes per entry

```
SELECT SUM(price)
FROM my_table
```



CPU

## Physical Storage



# Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 million rows  $\times$  30 columns  $\times$  100 bytes per entry = 3 GB

Data transfer speed: 200 MB/s

Total transfer time?  $\frac{3\text{GB or } 3000\text{ MB}}{200\text{ MB/s}} = 15\text{ s}$

```
SELECT SUM(price)
FROM my_table
```



# Row-Oriented Storage


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 billion rows X 30 columns X 100 bytes per entry = 3000 GB

Data transfer speed: 200 MB/s

Total transfer time?  $\frac{3000 \text{ GB}}{200 \text{ MB/s}} = 4 \text{ hours !}$

```
SELECT SUM(price)
FROM my_table
```





# Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

↓ Stores data  
Column by column

## Physical Storage

1	2	3	4	40	23	45	50	45865	90234	12558	45682	...
---	---	---	---	----	----	----	----	-------	-------	-------	-------	-----

bytes representing 1st column   bytes representing 2nd column   bytes representing 3rd column

# Column-Oriented Storage — Suitable for OLAP systems!


Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 billion rows      30 columns      100 bytes per entry      = 100 GB

Data transfer speed: 200 MB/s

Total transfer time?  $\frac{100 \text{ GB or } 100,000 \text{ MB}}{200 \text{ MB/s}} = 8.33 \text{ minutes}$

```
SELECT SUM(price)
FROM my_table
```



## Row-oriented Storage

Transfer 1 billion rows  
from disk to memory

**4 hours**

# Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

Terrible for  
transactional  
workloads!

Stores data  
Column by column

**Physical Storage**

1	2	3	4	40	23	45	50	45865	90234	12558	45682	...
---	---	---	---	----	----	----	----	-------	-------	-------	-------	-----

Deserialize the column, modify it,  
then write it back to storage




DeepLearning.AI

# Storage Systems

---

## **Graph Databases**

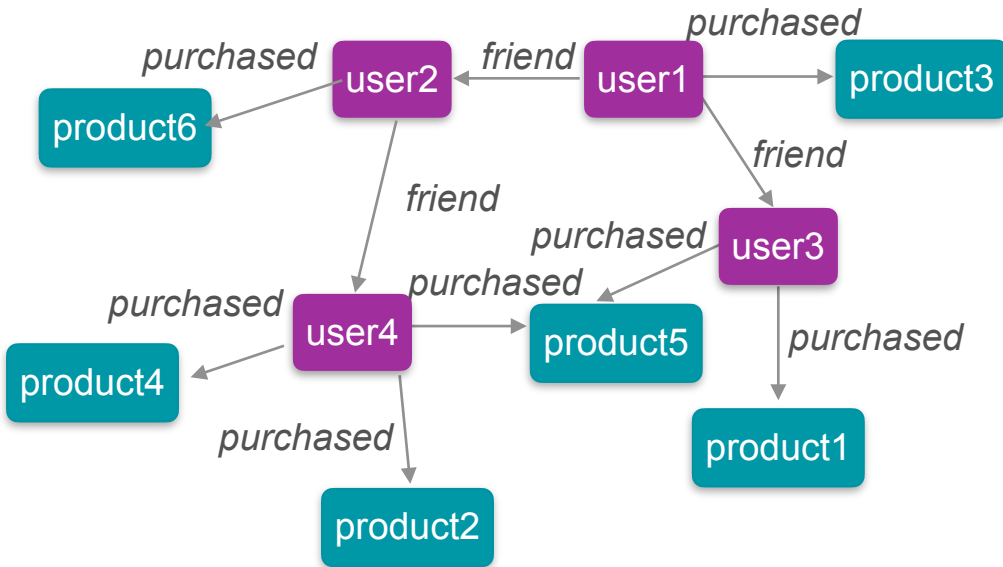
# Graph Database



- Nodes represent data items
- Edges represent connection between the data items
- Graph databases model complex connections between data entities

# Graph Database

## Relationships are first-class citizens



## Relational database

*purchase*

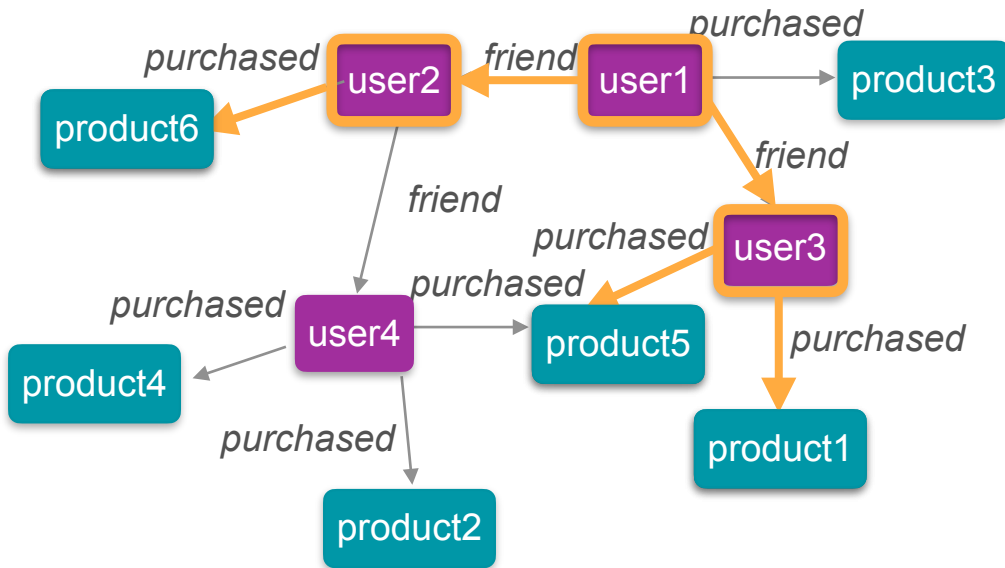
user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

*friendship*

user	friend
user1	user3
user1	user2
user2	user4

# Querying Data

Traverse the graph to query relationships



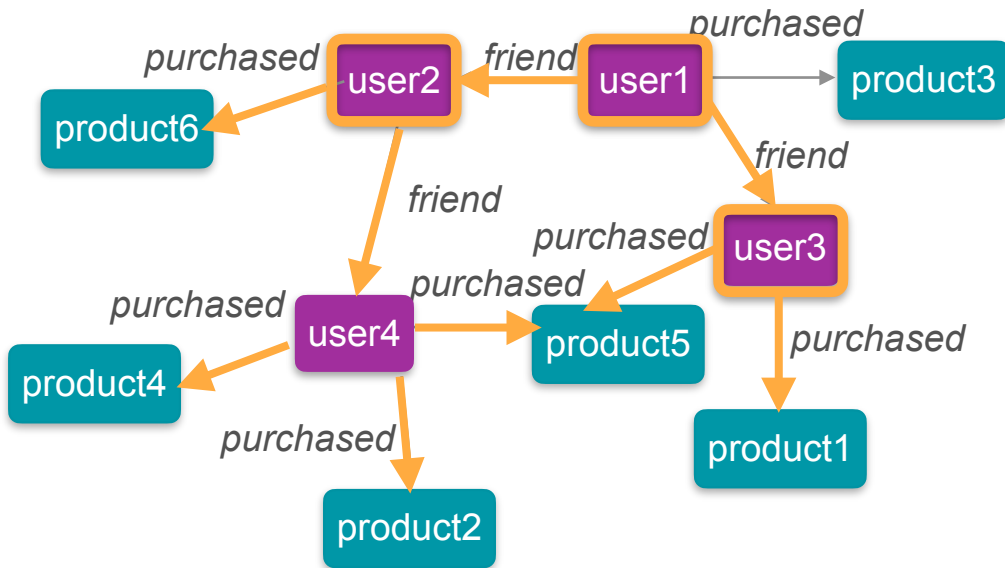
user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

user	friend
user1	user3
user1	user2
user2	user4

Recommendations  
for user1

# Querying Data

## Traverse the graph to query relationships



## Relational database

*purchase*

user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

*friendship*

user	friend
user1	user3
user1	user2
user2	user4

Less efficient in querying complex relationships!

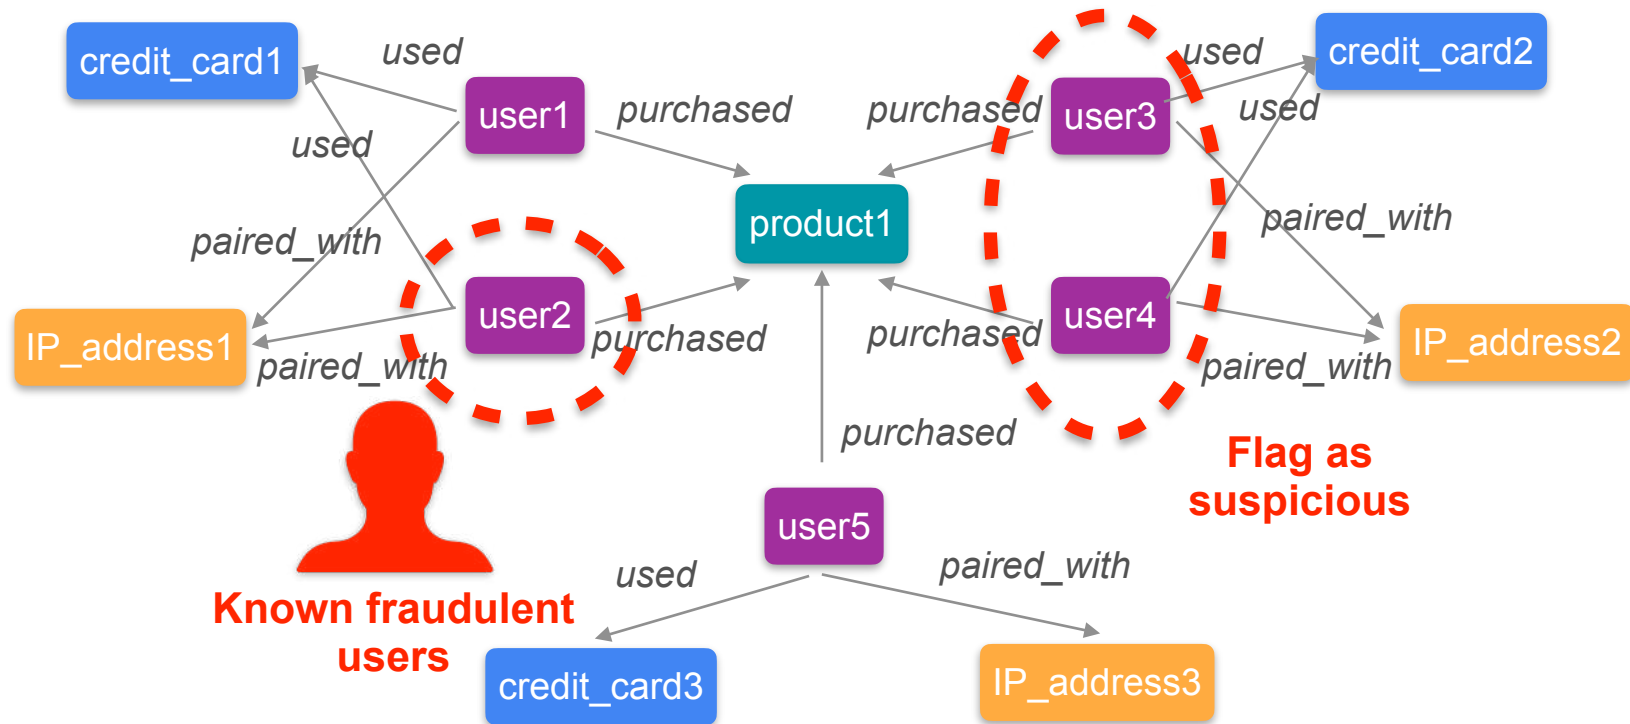
```
SELECT DISTINCT purchase.product
FROM friendship
JOIN purchase ON friendship.friend = purchase.user
WHERE friendship.user = 'user1'
```



# Graph Database - Use Cases

- Recommending products
- Modeling social networks
- Representing network and IT operations
- Simulating supply chains logistics
- Tracing data lineage

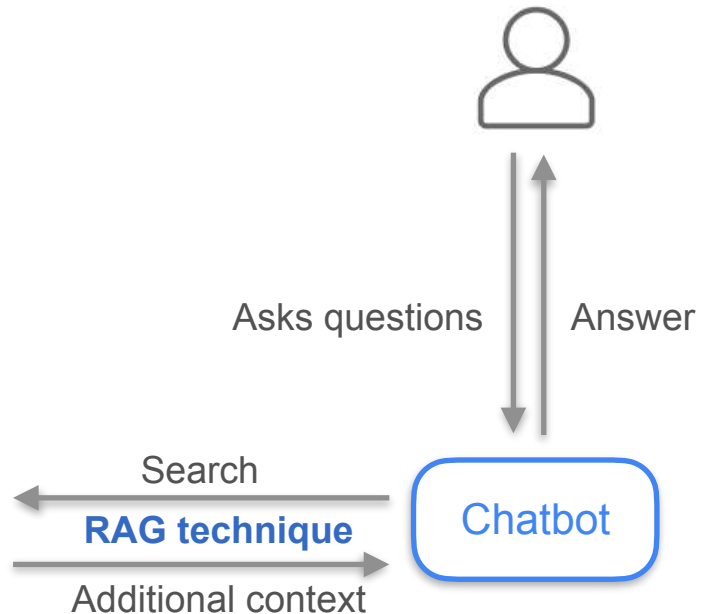
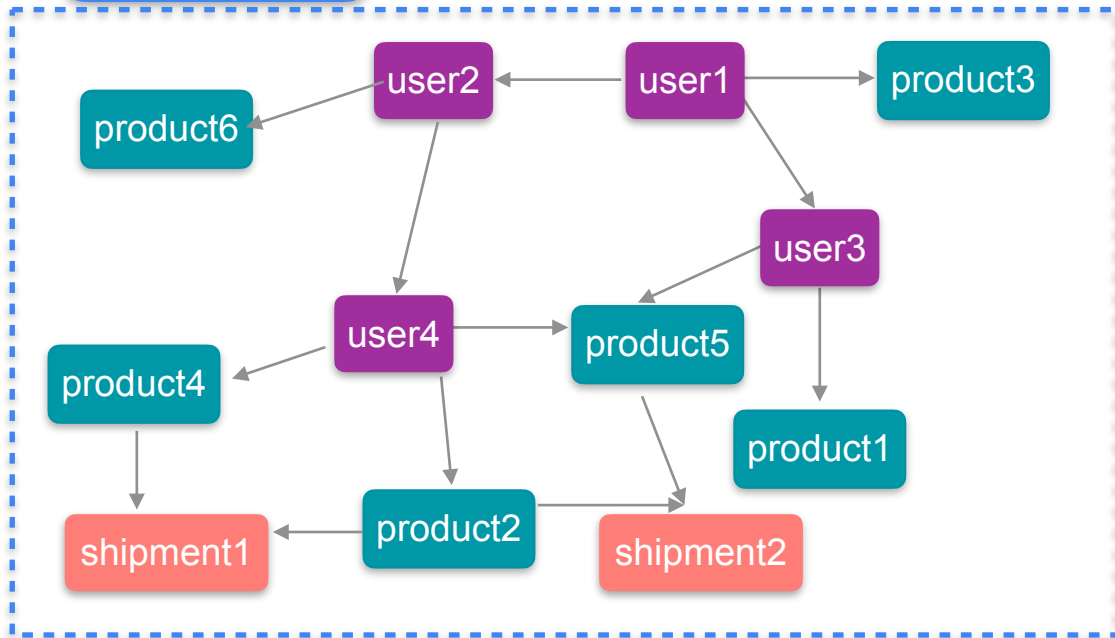
# Use Case - Fraud Detection



# Use Case - Knowledge Graph

## Knowledge Graph

Connect diverse data from disparate sources



# Graph Databases

## Examples of Graph Databases



Amazon Neptune

## Examples of Graph Query Language

Cypher

Gremlin

SparQL



DeepLearning.AI

# Storage Systems

---

## **Vector Databases**

## Vector data

Consists of numerical values arranged in an array



Image



0.5	0.5	0.7	0.5	0.3	0.1	0.7	0.2
0.4	0.8	0.9	0.1	0.3	0.1	0.4	0.2
0.3	0.5	0.7	0.8	0.3	0.1	0.6	0.2

## Vector embeddings


Capture semantic meaning of an item, like a text document or image




- Can convert an entire database of docs or text into embeddings
- Embeddings help you more efficiently find and retrieve similar items
- Example: Finding similar text
  - Compute embeddings for the query item
  - Database returns similar vectors (based on closeness)

# Distance Metric


**Vector database uses a distance metric to find similar vectors**



Euclidean distance




Cosine distance



Manhattan distance

# Similarity Search - Popular Algorithm


## K-nearest neighbors (KNN)



- Calculates distance to all vector embeddings
- Becomes inefficient when the data size increases
- Suffers from the curse of dimensionality

## ANN (Approximate Nearest Neighbors)

- Find a good guess for the nearest neighbors
- More efficient than K-NN
- Vector databases are built to support ANN algorithms





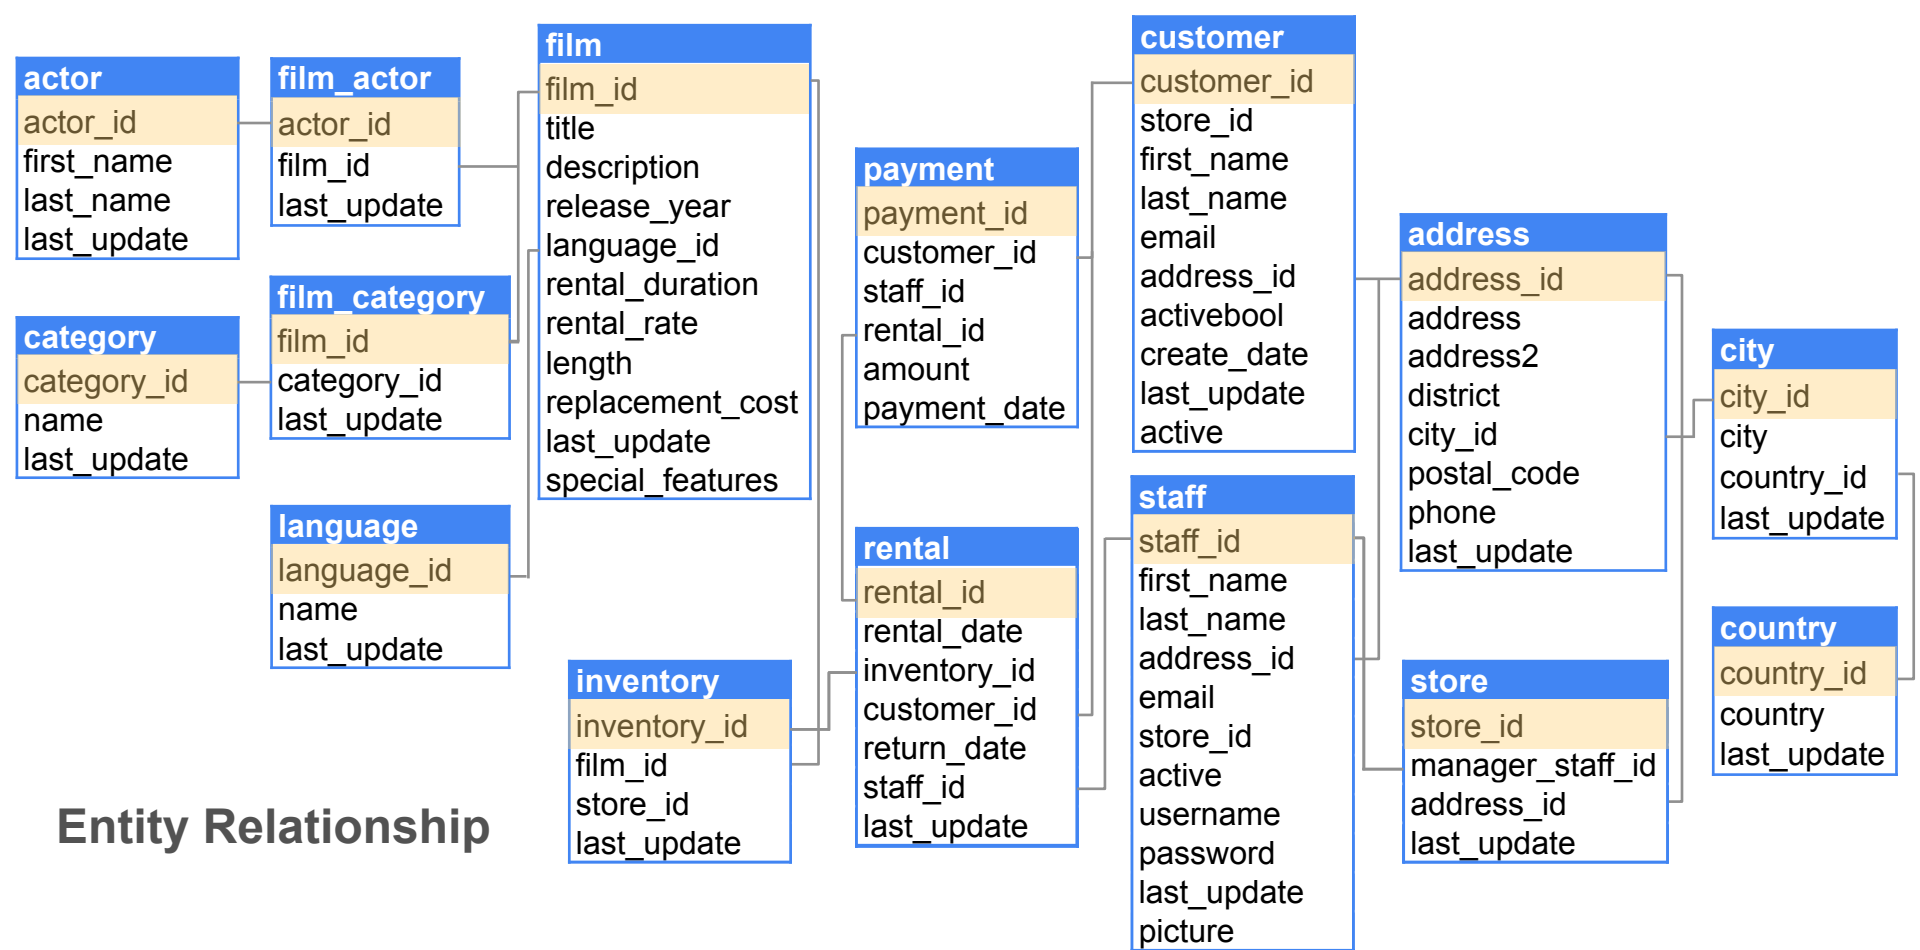


DeepLearning.AI

# Storage Systems

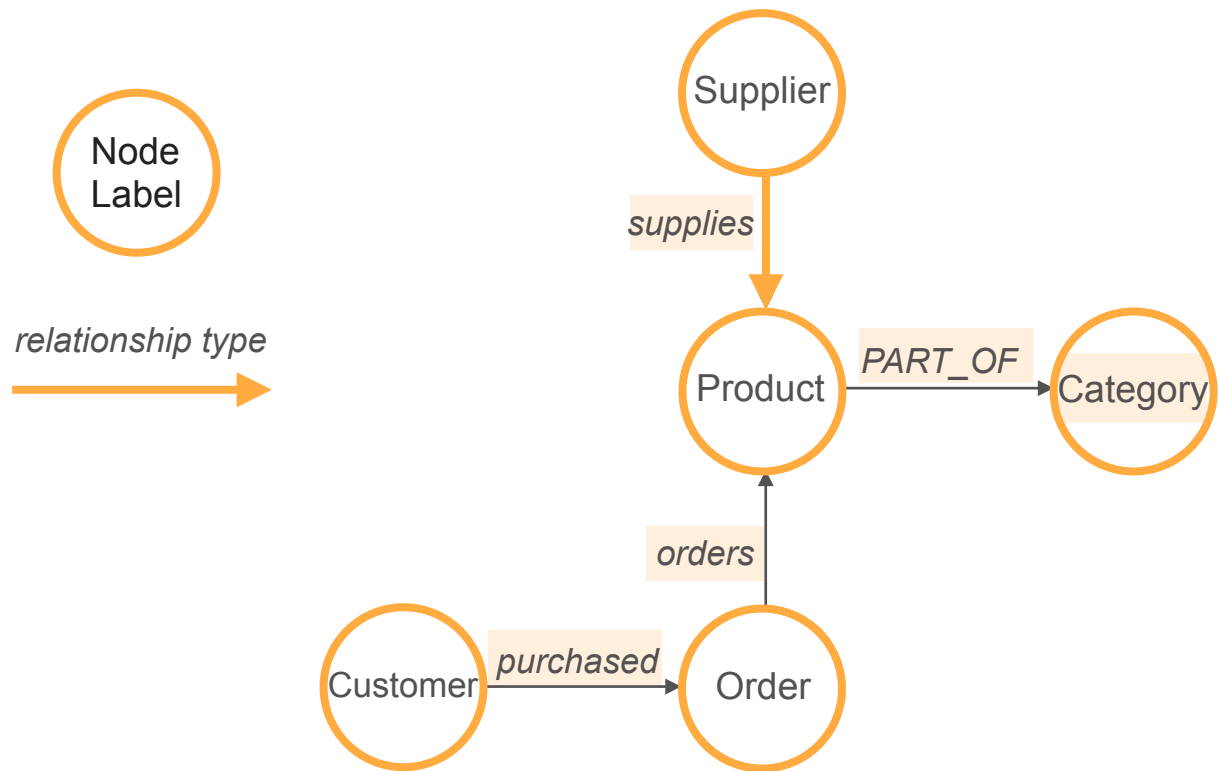
---

## **Neo4j Graph Database & Cypher Query Language (Part 1)**




# Entity Relationship

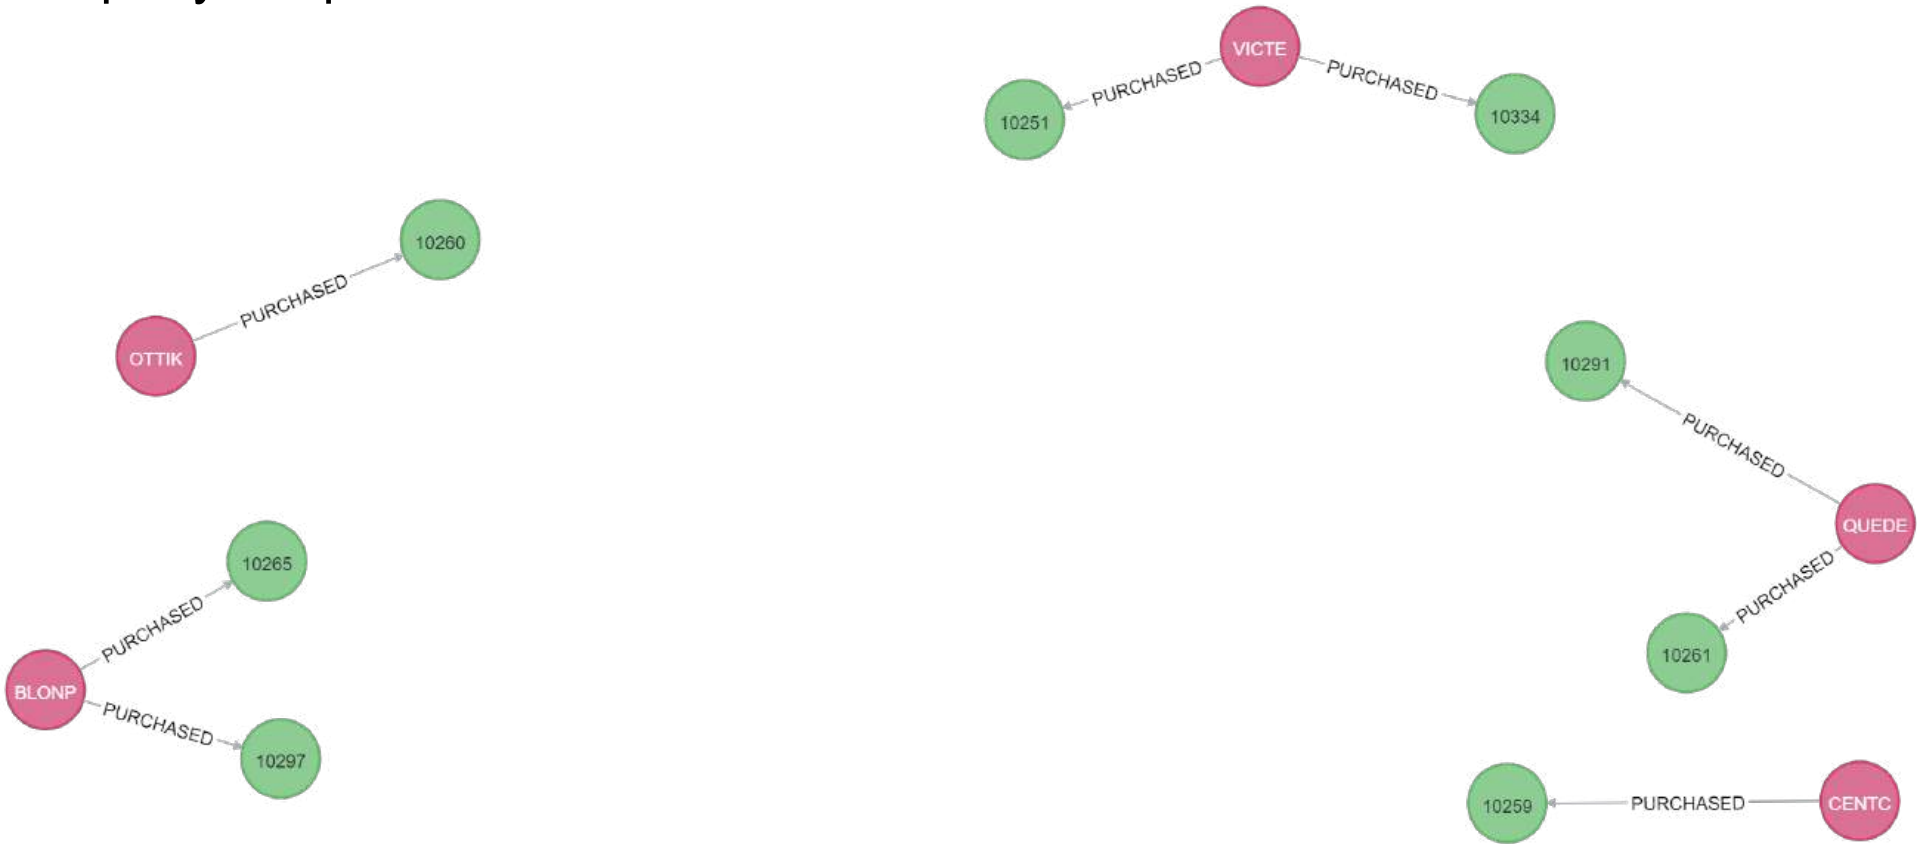
# Property Graph Model



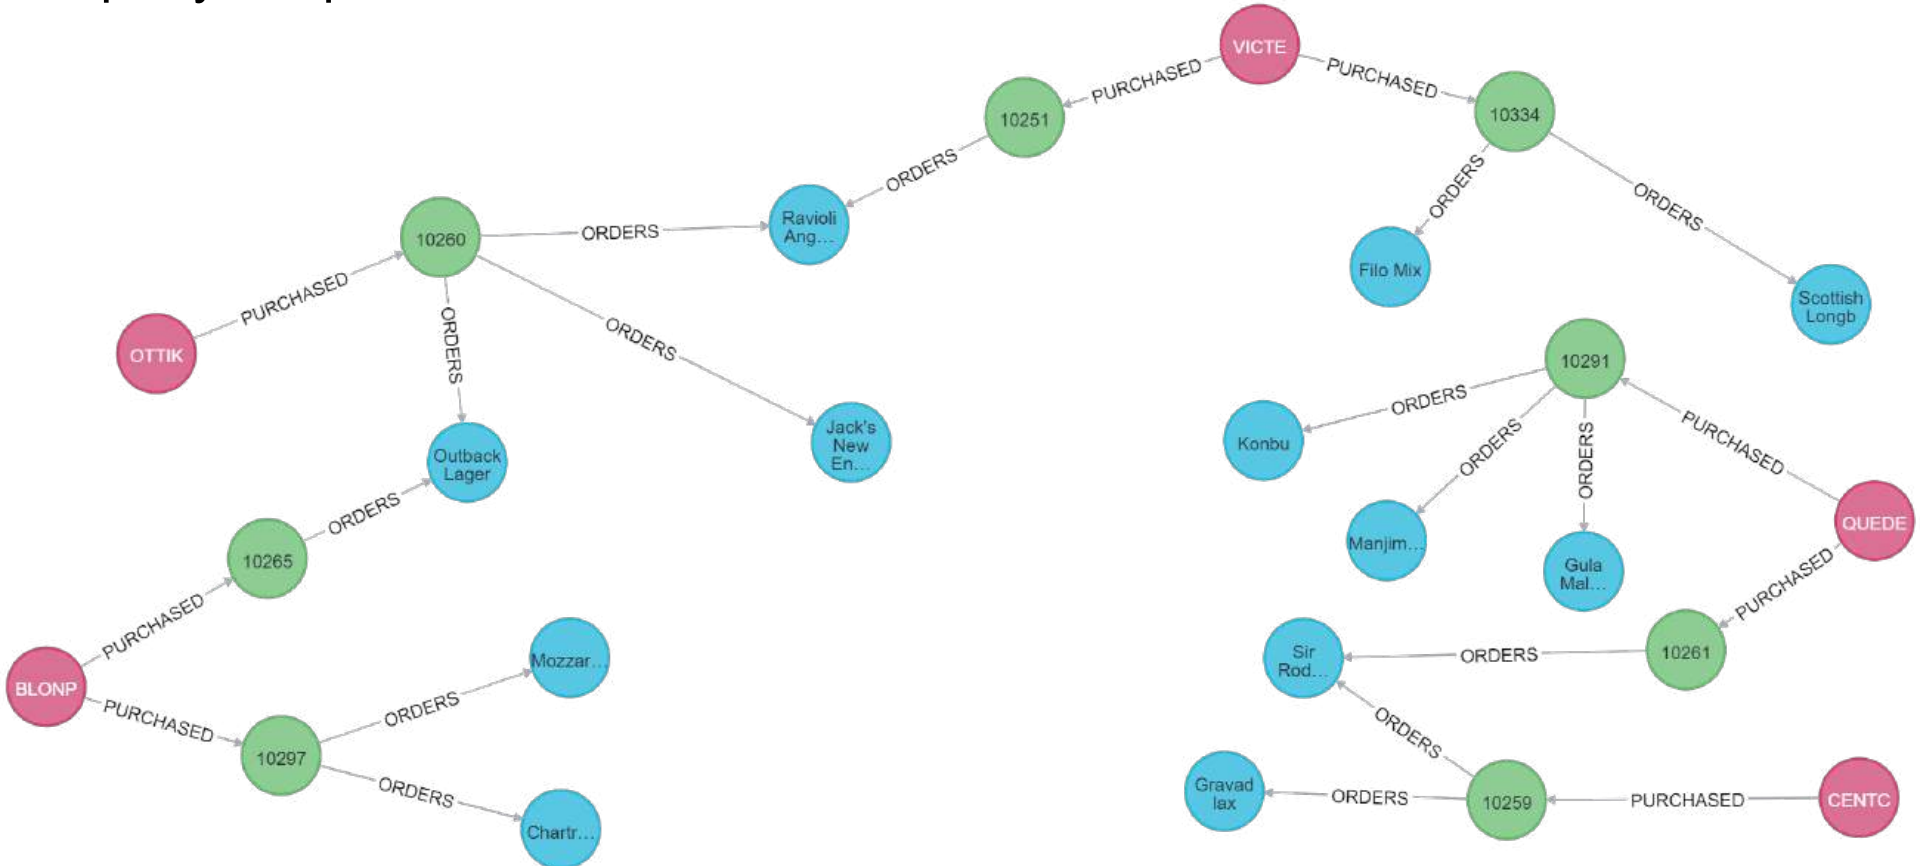
# Property Graph Model



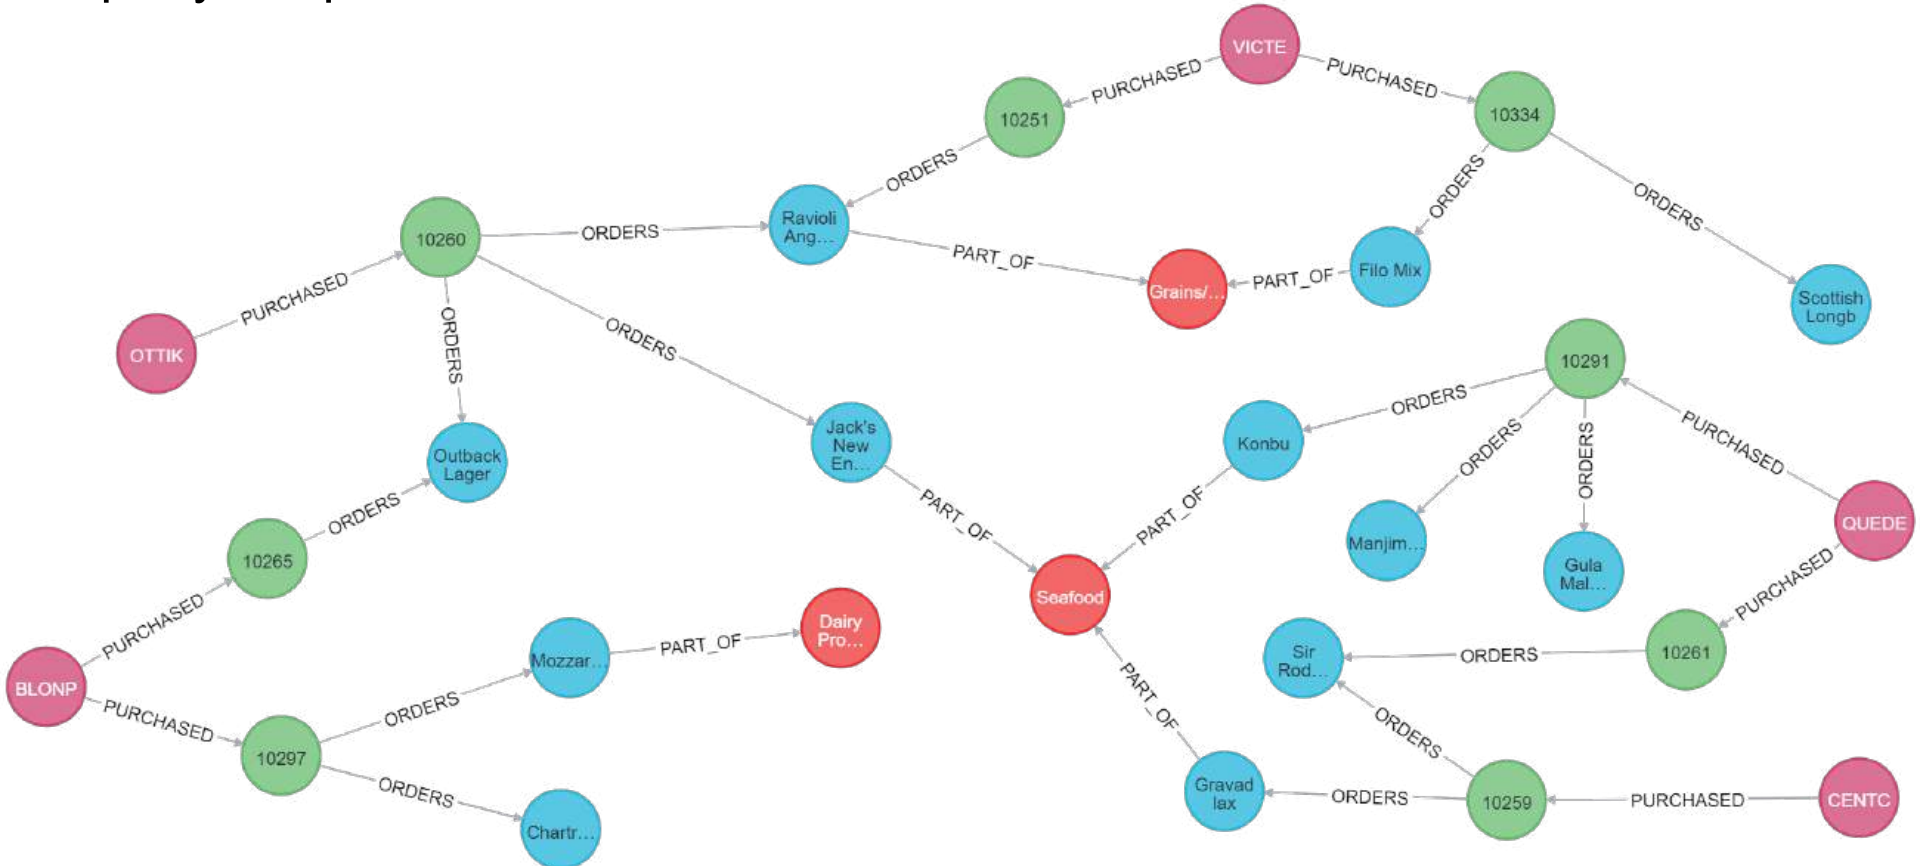
# Property Graph Model



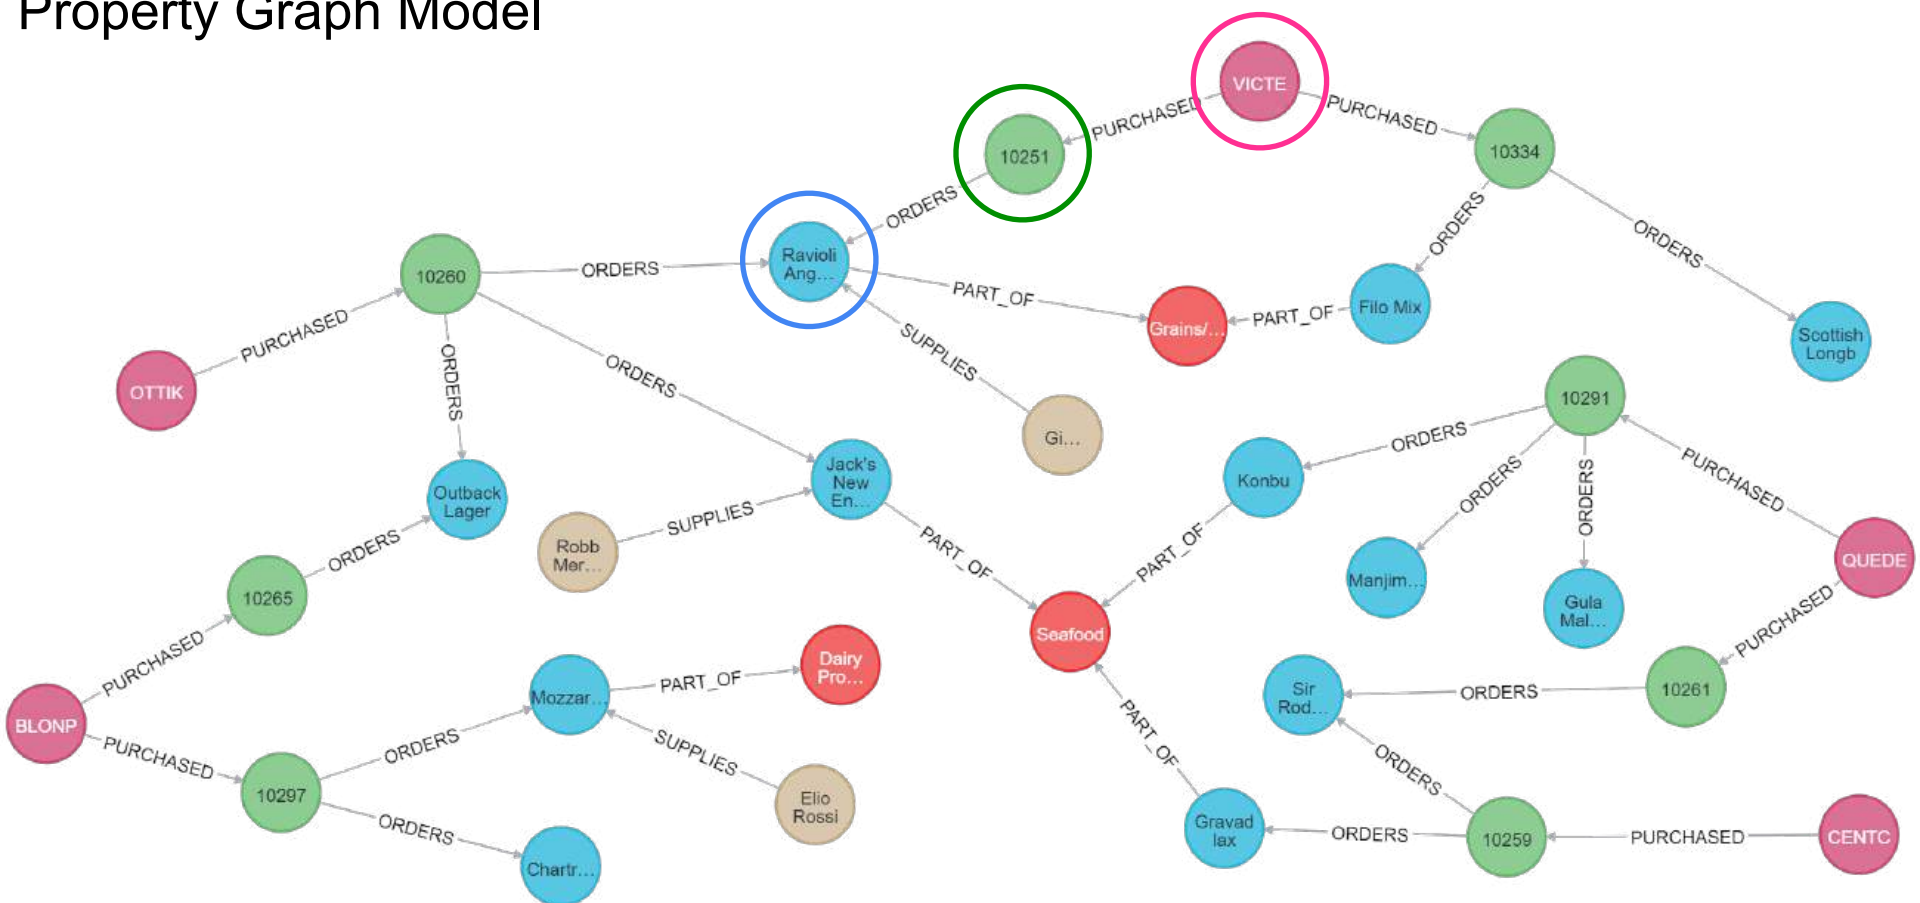
# Property Graph Model



# Property Graph Model

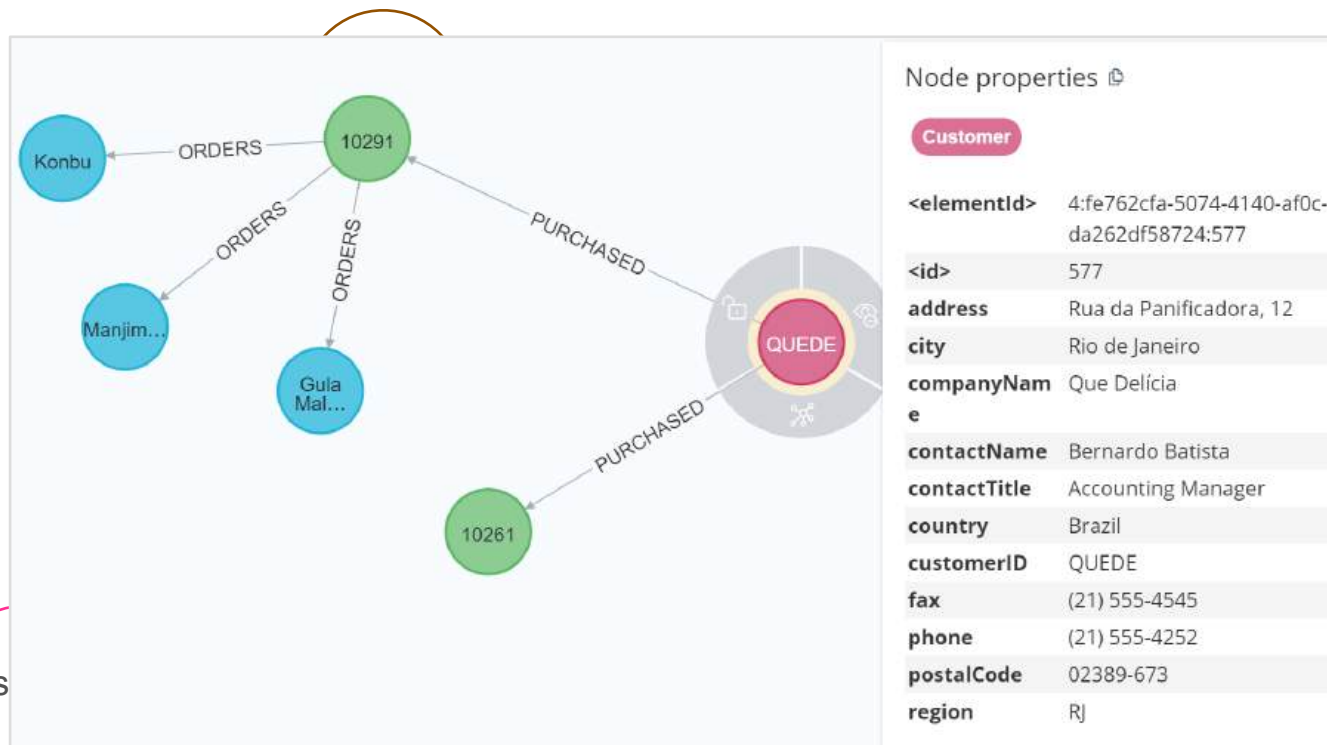


# Property Graph Model





# Property Graph Model



## Customer Properties

address  
city  
companyName  
contactName  
contactTitle  
country  
customerID

Cus

.....

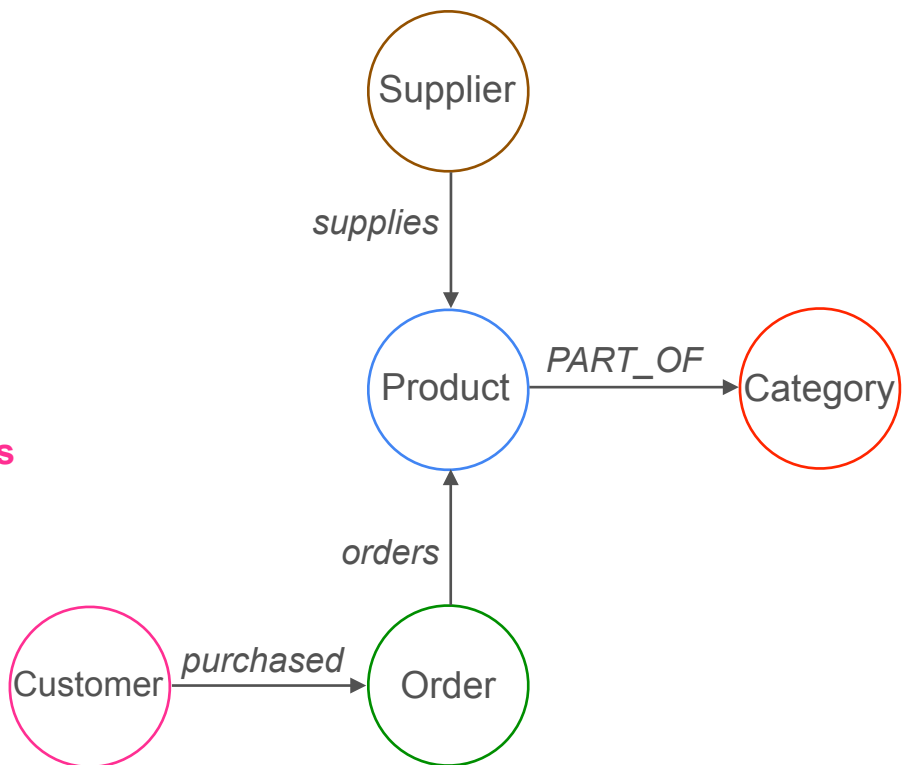
# Property Graph Model

## Product Properties

productID  
productName  
unitPrice  
unitsInStock  
unitsOnOrder

## Customer Properties

address  
city  
companyName  
contactName  
contactTitle  
country  
customerID  
.....



## Supplier Properties

address  
city  
contactName  
fax  
region  
supplierID  
postalCode

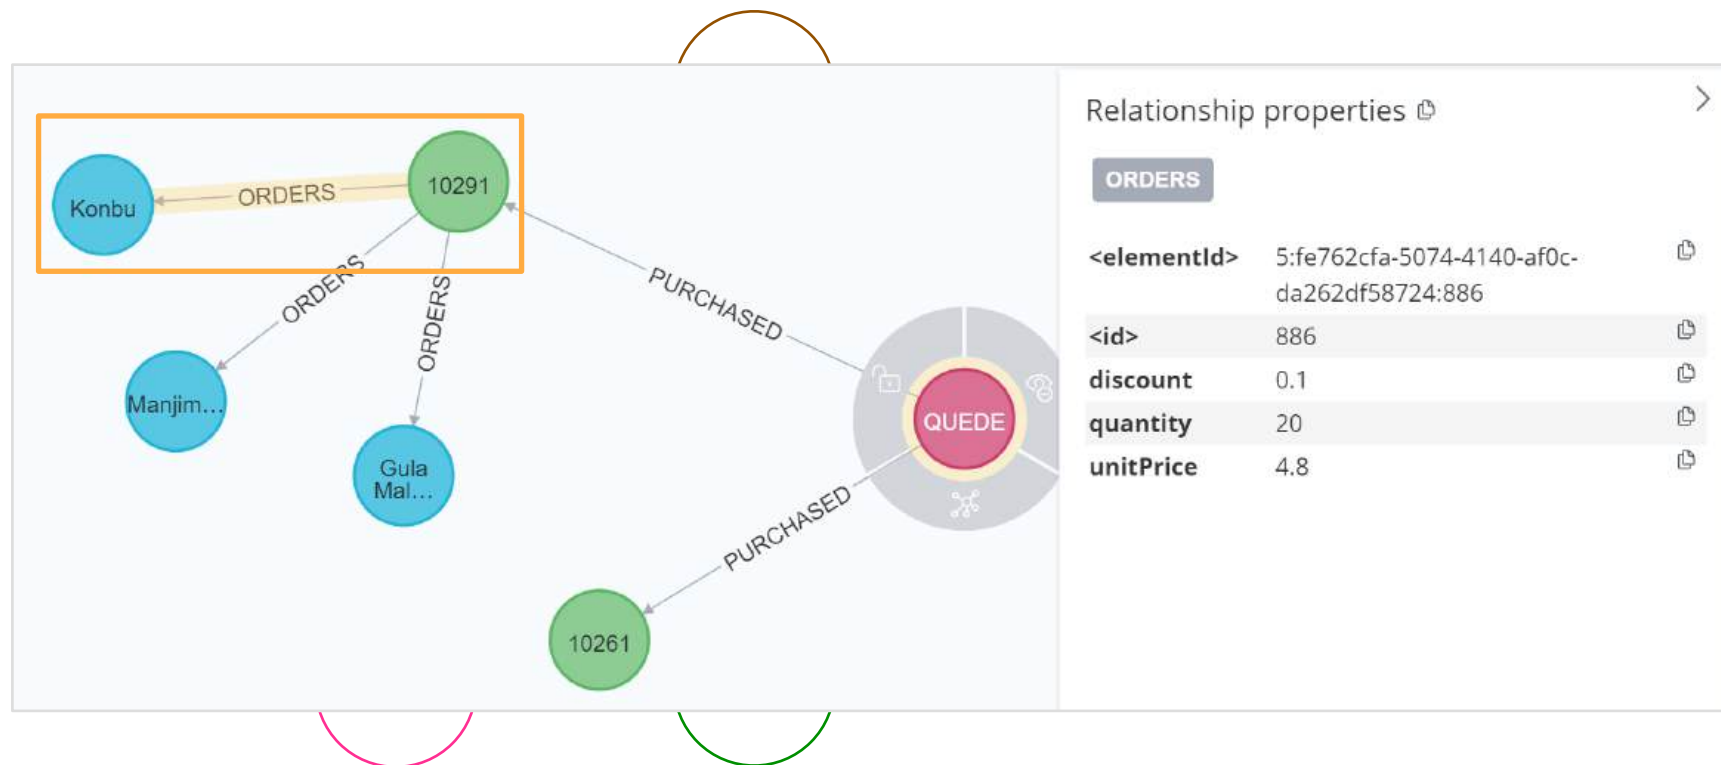
## Category Properties

categoryName

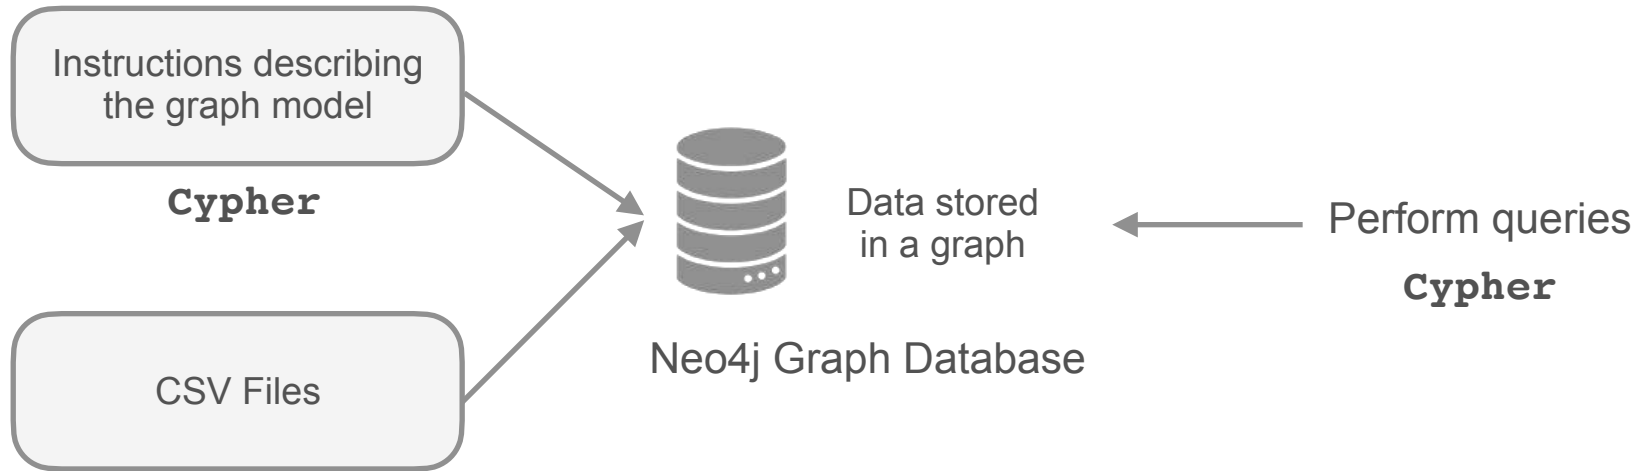
## Order Properties

freight  
orderDate  
orderID  
requiredDate  
shipAddress  
.....

# Property Graph Model



# Creating a Graph Database



- In the next video, we'll go through some queries examples.
- In the lab, you'll also practice CRUD operations.



DeepLearning.AI


# Storage Systems

---

## **Neo4j Graph Database & Cypher Query Language (Part 2)**

MATCH *pattern* RETURN *result*

node ( )



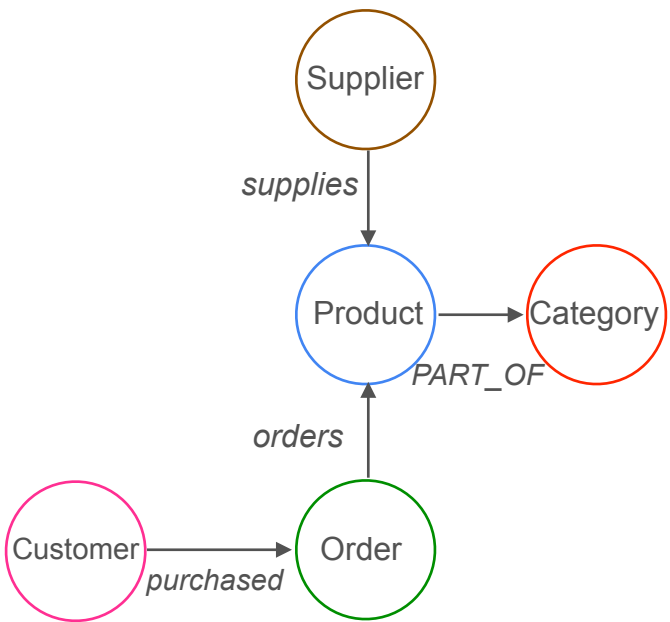
# MATCH Statement

Retrieve all nodes



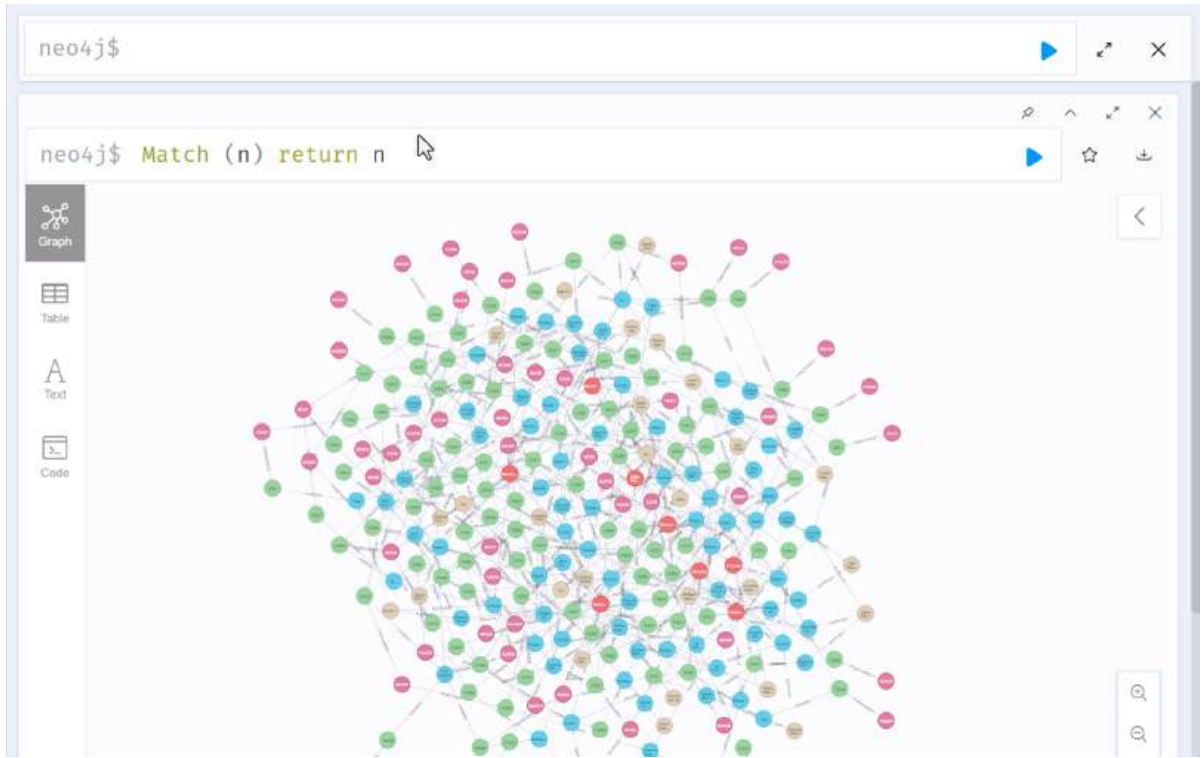
MATCH *pattern* RETURN *result*

node ( )




# MATCH Statement

Get the total number of nodes



MATCH *pattern* RETURN *result*

node ( )



# MATCH Statement

Explore the node labels using the `labels` function

```
neo4j$
```

```
neo4j$ Match (n) return count(n)
```

	count(n)
1	265


Started streaming 1 records after 1 ms and completed after 1 ms.

```
neo4j$ Match (n) return n
```



MATCH *pattern* RETURN *result*

node ( )



# MATCH Statement

Specify the label of the node

neo4j\$


neo4j\$ Match (n) return distinct labels(n)

	labels(n)
1	["Supplier"]
2	["Category"]
3	["Product"]
4	["Order"]
5	["Customer"]

Started streaming 5 records in less than 1 ms and completed after 23 ms.

MATCH *pattern* RETURN *result*

node ( )



# MATCH Statement

Explore the properties of each order node using the `Properties` function

```
neo4j$
```

```
neo4j$ Match (n:Order) return count(n)
```

	count(n)
1	99


Started streaming 1 records after 2 ms and completed after 10 ms.

```
neo4j$ Match (n) return distinct labels(n)
```

	labels(n)
--	-----------

MATCH *pattern* RETURN *result*

node ( )



# MATCH Statement

Explore the properties of each order node using the `Properties` function

neo4j\$

neo4j\$ `Match (n:Order) return Properties(n)`

Table

Text

Code

Properties(n)


1

```
{
  "shipCity": "Reims",
  "orderId": "10248",
  "shippedDate": "00:00.0",
  "orderDate": "00:00.0",
  "shipRegion": "NULL",
  "freight": "32.38",
  "shipName": "Vins et alcools Chevalier",
  "shipCountry": "France",
  "shipAddress": "59 rue de l'Abbaye",
  "requiredDate": "00:00.0",
  "shipPostalCode": "51100"
}
```

2

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



# MATCH Statement

Count all the directed paths

neo4j\$

neo4j\$ Match (n:Order) return Properties(n) limit 1

Table

Properties(n)

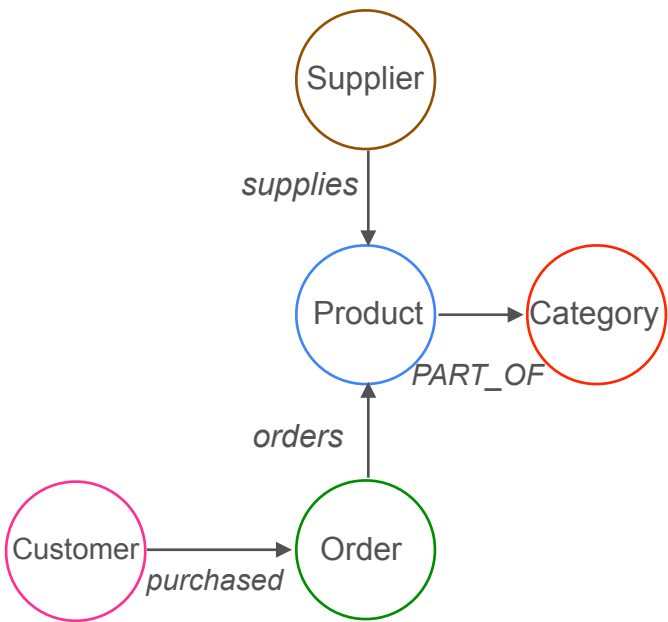
1

```
{
  "shipCity": "Reims",
  "orderId": "10248",
  "shippedDate": "00:00.0",
  "orderDate": "00:00.0",
  "shipRegion": "NULL",
  "freight": "32.38",
  "shipName": "Vins et alcools Chevalier",
  "shipCountry": "France",
  "shipAddress": "59 rue de l'Abbaye",
  "requiredDate": "00:00.0",
  "shipPostalCode": "51100"
}
```

Started streaming 1 records after 2 ms and completed after 8 ms.

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



# MATCH Statement

Return the types of relationships

```
neo4j$
```

```
neo4j$ Match ()-[r]->() return count(r)
```

count(r)	
1	518


Started streaming 1 records in less than 1 ms and completed after 5 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

Properties(n)	
---------------	--

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



# MATCH Statement

Specify the type of the relationship

```
neo4j$
```

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

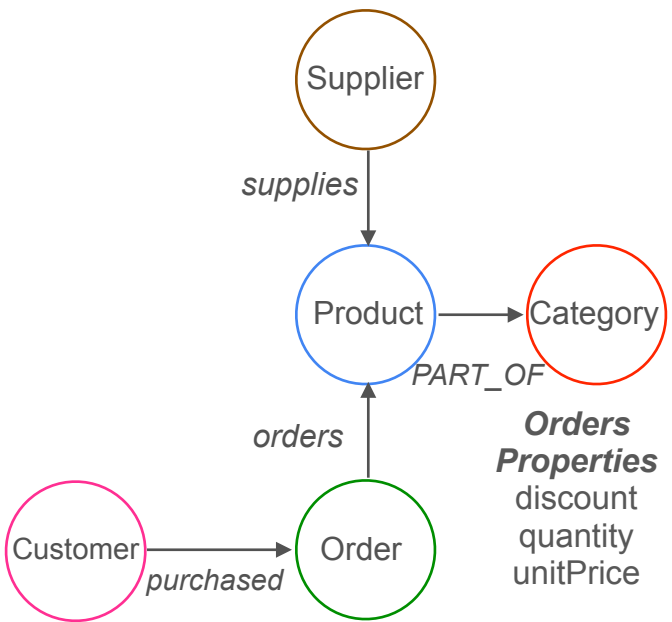
	type(r)
1	"SUPPLIES"
2	"PART_OF"
3	"PURCHASED"
4	"ORDERS"

Started streaming 4 records after 1 ms and completed after 21 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



# MATCH Statement

Return the properties of a relationship

```
neo4j$ Match ()-[r:ORDERS]->()
```

```
neo4j$ Match ()-[r]->() return distinct type(r)
```


	type(r)
1	"SUPPLIES"
2	"PART_OF"
3	"PURCHASED"
4	"ORDERS"

Started streaming 4 records after 1 ms and completed after 21 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Return the properties of a relationship

```
neo4j$
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice)
```

AVG(r.quantity*r.unitPrice)	
1	502.4645283018867

Started streaming 1 records after 1 ms and completed after 25 ms.


```
neo4j$ Match ()-[r]->() return distinct type(r)
```

type(r)	
---------	--



MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Get the average price for all orders grouped by product category

```
neo4j$
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice) as a...
```

	average_price
1	502.4645283018867


Started streaming 1 records in less than 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

	type(r)
--	---------

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Get the average price for all orders grouped by product category

```
neo4j$ MATCH ()-[r:ORDERS]->()-[part:PART_OF]->(c:Category) return  
AVG(r.quantity*r.unitPrice) as average_price
```

```
neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice) as a...
```

average_price
502.4645283018867


Started streaming 1 records in less than 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

type(r)
---------

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”


neo4j\$

neo4j\$ Match ()-[r:ORDERS]->()-[part:PART\_OF]->(c:Category) return c.c...

	c.categoryName	average_price
2	"Confections"	531.36666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”


```
neo4j$ Match (p:Product)-[:PART_OF]->(c:Category)
```

```
neo4j$ Match ()-[r:ORDERS]->()-[part:PART_OF]->(c:Category) return c.c...
```

	c.categoryName	average_price
2	"Confections"	531.3666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”


```
1 Match (p:Product)-[:PART_OF]→(c:Category)
2 where c.categoryName="Meat/Poultry"
```

```
neo4j$ Match ()-[r:ORDERS]→()-[part:PART_OF]→(c:Category) return c.c...
```

	c.categoryName	average_price
3	"Confections"	531.36666666666664
4	"Dairy Products"	504.31836734693877
5	"Grains/Cereals"	278.67
6	"Meat/Poultry"	766.9142857142859
7	"Produce"	645.6571428571427
8	"Seafood"	434.8282051282051

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

neo4j\$


```
1 Match (p:Product)-[:PART_OF]-(c:Category)
2 where c.categoryName="Meat/Poultry"
3 return p.productName, p.unitPrice
```

Table

	p.productName	p.unitPrice
1	"Perth Pasties"	32.8
2	"Tourtière"	7.45
3	"Alice Mutton"	39.0
4	"Pâté chinois"	24.0
5	"Thüringer Rostbratwurst"	123.79
6	"Mishi Kobe Niku"	97.0

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the product name of all products ordered by the customer “QUEDE”


neo4j\$

neo4j\$ Match (p:Product)-[:PART\_OF]-(c:Category {categoryName:"Meat/P...}

	p.productName	p.unitPrice
1	"Perth Pasties"	32.8
2	"Tourtière"	7.45
3	"Alice Mutton"	39.0
4	"Pâté chinois"	24.0
5	"Thüringer Rostbratwurst"	123.79
6	"Mishi Kobe Niku"	97.0

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Get the ID of other customers who ordered the same products as “QUEDE”

```
neo4j$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:ORDERS]->(p:Product)
```

```
neo4j$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:OR...
```


	p.productName
1	"Steeleye Stout"
2	"Sir Rodney's Scones"
3	"Gula Malacca"
4	"Konbu"
5	"Manjimup Dried Apples"

Started streaming 5 records after 1 ms and completed after 1 ms.



MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the orders that contain at most two products

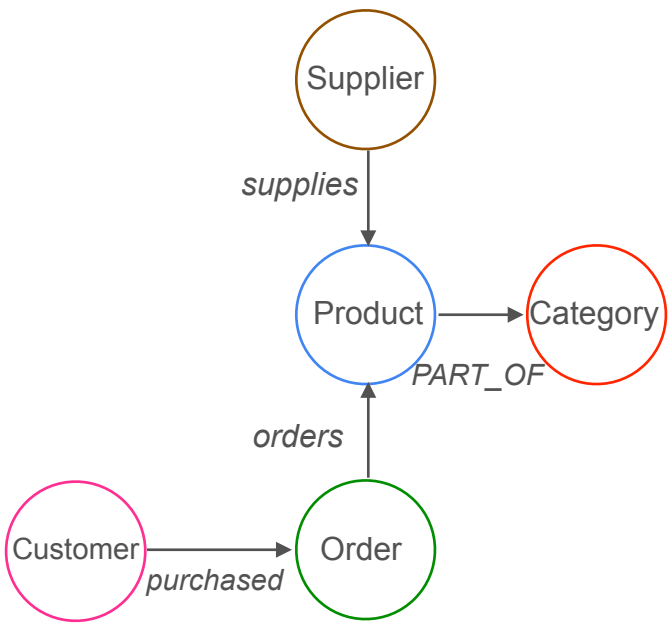
neo4j\$

neo4j\$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]->()-[:OR...

	c2.customerID
1	"QUICK"
2	"SAVEA"
3	"ROMEY"
4	"ISLAT"
5	"WARTH"
6	"CENTC"

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



# MATCH Statement

Retrieve the orders that contain at most two products

neo4j\$

```
1 Match (o:Order)-[:ORDERS]-(p:Product)
2 return o.orderID as ID, count(p) as countProd
```

ID	countProd
"10294"	5
"10317"	1
"10285"	3
"10342"	4
"10255"	4
"10327"	4



DeepLearning.AI

# Storage Systems

---

## Summary

# Raw Storage Ingredients

## Persistent Storage Medium

Magnetic disk




Solid-state storage



## Volatile Memory

RAM




CPU cache




# Cloud Storage Options

**File Storage**



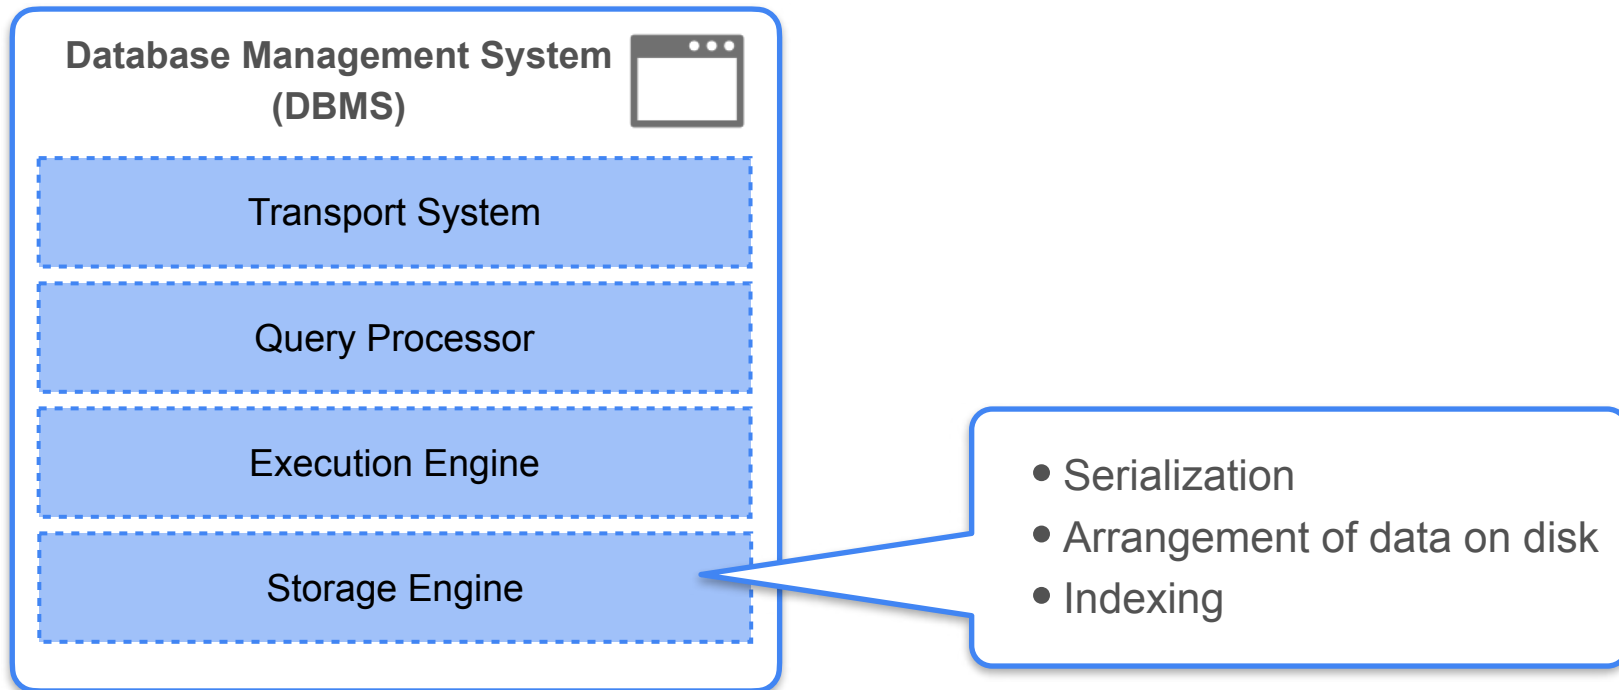
**Block Storage**



**Object Storage**



# Storage in Databases



# Row and Columnar Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	458650	10	67t
2	23	902348	14	56t
3	45	1255893	12	87q
4	50	456829	13	98q

## Row-oriented storage


bytes representing the 1st row	bytes representing the 2nd row	...	bytes representing the last row
--------------------------------	--------------------------------	-----	---------------------------------

## Column-oriented storage

bytes representing the 1st column	bytes representing the 2nd column	...	bytes representing the last column
-----------------------------------	-----------------------------------	-----	------------------------------------

# Databases

## Graph Databases



Cypher

## Vector Databases

